# The XSB System
Version 3.2
Volume 2: Libraries, Interfaces and Packages

March 29, 2009

Credits

Interfaces have become an increasingly important part of XSB. The interface from C to Prolog was implemented by David Warren as was the DLL interface; the interface from Prolog to C (foreign language interface) was developed by Jiyang Xu, Kostis Sagonas and Steve Dawson. The Oracle interface was written by Hassan Davulcu and Ernie Johnson. The ODBC took as its starting point the Oracle interface, and was written by Lily Dong and Baoqiu Cui, and maintained by David Warren. The interface to POSIX regular expression and wildcard matching as well as the Libwww-based Web access package was written by Michael Kifer. The interface to Perl pattern matching routines was written by Michael Kifer and Jin Yu. The SModels interface was written by Luis F. Castro.

The SLX preprocessor was written by José Júlio Alferes and Luís Moniz Pereira. Unix-style scripting libraries were written by Terrance Swift, and the ordset library was written by Richard O'Keefe.

Contents

Chapter 1

Library Utilities

In this chapter we introduce libraries of some useful predicates that are supplied with XSB. Interfaces and more elaborate packages are documented in later chapters. These predicates are available only when imported them from (or explicitly consult) the corresponding modules.

1.1   List Processing

The XSB library contains various list utilities, some of which are listed below. These predicates should be explicitly imported from the module specified after the skeletal specification of each predicate. There are a lot more useful list processing predicates in various modules of the XSB system, and the interested user can find them by looking at the sources.

`append(?List1, ?List2, ?List3)`   module: basics
>   Succeeds if list `List3` is the concatenation of lists `List1` and `List2`.

`member(?Element, ?List)`   module: basics
>   Checks whether `Element` unifies with any element of list `List`, succeeding more than once if there are multiple such elements.

`memberchk(?Element, ?List)`   module: basics
>   Similar to `member`/2, except that `memberchk`/2 is deterministic, i.e. does not succeed more than once for any call.

`ith(?Index, ?List, ?Element)`   module: basics
>   Succeeds if the $Index^{th}$ element of the list `List` unifies with `Element`. Fails if `Index` is not a positive integer or greater than the length of `List`. Either `Index` and `List`, or `List` and `Element`, should be instantiated (but not necessarily ground) at the time of the call.

`delete_ith(+Index, +List, ?Element, ?RestList)`   module: listutil
>   Succeeds if the $Index^{th}$ element of the list `List` unifies with `Element`, and `RestList` is `List` with `Element` removed. Fails if `Index` is not a positive integer or greater than the length of `List`.

1

`log_ith(?Index, ?Tree, ?Element)`   module: basics

Succeeds if the Index[th] element of the Tree `Tree` unifies with `Element`. Fails if `Index` is not a positive integer or greater than the number of elements that can be in `Tree`. Either `Index` and `Tree`, or `Tree` and `Element`, should be instantiated (but not necessarily ground) at the time of the call. Tree is a list of full binary trees, the first being of depth 0, and each one being of depth one greater than its predecessor. So `log_ith/3` is very similar to `ith/3` except it uses a tree instead of a list to obtain log-time access to its elements.

`log_ith_bound(?Index, ?Tree, ?Element)`   module: basics

is like `log_ith/3`, but only if the Index[th] element of `Tree` is non-variable and equal to Element. This predicate can be used in both directions, and is most useful with Index unbound, since it will then bind `Index` and `Element` for each non-variable element in `Tree` (in time proportional to $N * logN$, for N the number of non-variable entries in `Tree`.)

`length(?List, ?Length)`   module: basics

Succeeds if the length of the list `List` is `Length`. This predicate is deterministic if `List` is instantiated to a list of definite length, but is nondeterministic if `List` is a variable or has a variable tail. If `List` is uninstantiated, it is unified with a list of length `Length` that contains variables.

`same_length(?List1, ?List2)`   module: basics

Succeeds if list `List1` and `List2` are both lists of the same number of elements. No relation between the types or values of their elements is implied. This predicate may be used to generate either list (containing variables as elements) given the other, or to generate two lists of the same length, in which case the arguments will be bound to lists of length $0, 1, 2, \ldots$.

`select(?Element, ?L1, ?L2)`   module: basics

`List2` derives from `List1` by selecting (removing) an `Element` non-deterministically.

`reverse(+List, ?ReversedList)`   module: basics

Succeeds if `ReversedList` is the reverse of list `List`. If `List` is not a proper list, `reverse/2` can succeed arbitrarily many times. It works only one way.

`perm(+List, ?Perm)`   module: basics

Succeeds when `List` and `Perm` are permutations of each other. The main use of `perm/2` is to generate permutations of a given list. `List` must be a proper list. `Perm` may be partly instantiated.

`subseq(?Sequence, ?SubSequence, ?Complement)`   module: basics

Succeeds when `SubSequence` and `Complement` are both subsequences of the list `Sequence` (the order of corresponding elements being preserved) and every element of `Sequence` which is not in `SubSequence` is in the `Complement` and vice versa. That is,

$$length(\texttt{Sequence}) = length(\texttt{SubSequence}) + length(\texttt{Complement})$$

for example, `subseq([1,2,3,4], [1,3], [2,4])`. The main use of `subseq/3` is to generate subsets and their complements together, but can also be used to interleave two lists in all possible ways.

`merge(+List1, +List2, ?List3)`   module: listutil

 Succeeds if `List3` is the list resulting from "merging" lists `List1` and `List2`, i.e. the elements of `List1` together with any element of `List2` not occurring in `List1`. If `List1` or `List2` contain duplicates, `List3` may also contain duplicates.

`absmerge(+List1, +List2, ?List3)`   module: listutil

 Predicate `absmerge/3` is similar to `merge/3`, except that it uses predicate `absmember/2` described below rather than `member/2`.

`absmember(+Element, +List)`   module: listutil

 Similar to `member/2`, except that it checks for identity (through the use of predicate `'=='/2`) rather than unifiability (through `'='/2`) of `Element` with elements of `List`.

`member2(?Element, ?List)`   module: listutil

 Checks whether `Element` unifies with any of the actual elements of `List`. The only difference between this predicate and predicate `member/2` is on lists having a variable tail, e.g. `[a, b, c | _ ]`: while `member/2` would insert `Element` at the end of such a list if it did not find it, Predicate `member2/2` only checks for membership but does not insert the `Element` into the list if it is not there.

`closetail(?List)`   module: listutil

 Predicate `closetail/1` closes the tail of an open-ended list. It succeeds only once.

1.1.1 Processing Comma Lists

It is often useful to process comma lists when meta-interpreting or preprocessing. XSB libraries include the following simple utilities.

`comma_to_list(+CommaList,-List)`   module: basics

 Transforms `CommaList` to `List`.

`comma_append(?CL1,?CL2,?CL3)`   basics
`comma_length(?CommaList,?Length)`   basics
`comma_member(?Element,?CommaList)`   basics
`comma_member(?Element,?CommaList)`   module: basics

 Analogues for comma lists of `append/3`, `length/3`, `member/2` and `memberchk/2`, respectively.

1.2 Attributed Variables

Attributed variables are a special data type that associates variables with arbitrary attributes as well as supports extensible unification. Attributed variables have proven to be a flexible and powerful mechanism to extend a classic logic programming system with the ability of constraint solving. Our low-level API for constraints closely resembles that of hProlog [8] and SWI [24].

1.2.1 Lowlevel Interface

Attributes of variables are pairs of attribute module names and values. An attribute module name can be any atom. A value can be any XSB value (term, variable, atom, . . .). Any variable has at most one attribute for a particular attribute module. Attribute modules are distinct from XSB modules: although it is most efficient to keep each handlers for each attribute module in their own XSB module. c Attributes can be manipulated with the following three predicates (`get_attr/3`, `put_attr/3` and `del_attr/2`) defined in the module `machine`.

`get_attr(-Var,+Mod, ?Val)`   module: machine

> Gets the value of the attribute of `Var` in attribute module `Mod`. Non-variable terms in `Var` cause a type error. `Val` will be unified with the value of the attribute, if it exists. Otherwise the predicate fails.

`put_attr(-Var,+Mod, ?Val)`   module: machine

> Sets the value of the attribute of `Var` in attribute module `Mod`. Non-variable terms in `Var` cause a type error. The previous value of the attribute is overwritten, if it exists.

`del_attr(-Var, +Mod)`   module: machine

> Removes the attribute of `Var` in attribute module `Mod`. Non-variable terms in `Var` cause a type error. The previous value of the attribute is removed, if it exists.

One has to extend the default unification algorithm for used attributes by installing a handler in the following way:

$$\text{:- install\_verify\_attribute\_handler}(+Mod, -AttrValue, -Target, +Handler)$$

The `install_verify_attribute_handler/4` predicate is defined in module `machine`. *Mod* is the attribute `Module` and *Handler* is a term with arguments *AttrValue* and *Target*. The *Handler* term has to correspond to a handler predicate that takes the value of the attribute (*AttrValue*) and the term that the attributed value is bound to (*Target*) as arguments. To get good efficiency, it is usually best to keep the handlers for each attribute module in separate XSB modules.

The handler is called after the unification of an attributed variable with a term or other attributed variable, if the attributed variable has an attribute in the corresponding module. The two arguments of the unification are already bound at the time the handler is called, i.e. the handler is a post-unify handler.

Here, by giving the implementation of a simple finite domain constraint solver (see the file `fd.P` below), we show how these lowlevel predicates for attributed variables can be used. In this example, an attribute in the module `fd` is used and the value of this attribute is a list of terms.

```
%% File: fd.P
%%
%% A simple finite domain constrait solver implemented using the lowlevel
%% attributes variables interface.

:- import put_attr/3, get_attr/3, del_attr/3,
```

```
    install_verify_attribute_handler/4 from machine.
:- import member/2 from basics.

:- install_verify_attribute_handler(fd,AttrValue,Target,fd_handler(AttrValue,Target)).

fd_handler(Da, Target) :-
        (var(Target),                     % Target is an attributed variable
         get_attr(Target, fd, Db) ->      % has a domain
           intersection(Da, Db, [E|Es]),  % intersection not empty
           (Es = [] ->                    % exactly one element
              Target = E                  % bind Var (and Value) to E
           ;  put_attr(Target, fd, [E|Es]) % update Var's (and Value's)
           )
        ;  member(Target, Da)             % is Target a member of Da?
        ).

intersection([], _, []).
intersection([H|T], L2, [H|L3]) :-
        member(H, L2), !,
        intersection(T, L2, L3).
intersection([_|T], L2, L3) :-
        intersection(T, L2, L3).

domain(X, Dom) :-
        var(Dom), !,
        get_attr(X, fd, Dom).
domain(X, List) :-
        List = [El|Els],                  % at least one element
        (Els = []                         % exactly one element
         -> X = El                        % implied binding
        ;  put_attr(Fresh, fd, List),     % create a new attributed variable
           X = Fresh                      % may call verify_attributes/2
        ).

show_domain(X) :-                         % print out the domain of X
        var(X),                           % X must be a variable
        get_attr(X, fd, D),
        write('Domain of '), write(X),
        write(' is '), writeln(D).
```

When writing or porting a constraint package, it is usually useful to adjust the way that correct answer substitutions are shown in the command line. This can be controlled using the following two predicates:

install_attribute_portray_hook(Module,Attribute,Handler)   module: machine
> This hook is called by the command-line interpreter when printing out the value of each variable in a top-level query. When a printing out an attributed variable, any appropriate handlers are called to portray the constraints represented by the attribute. As an example, the bounds package (Section 11.2) uses a hook to print out the bounds of variables:

```
| ?- X in 1..10,Y in 1..10,X + 4 #< Y -3.

X = _h629 { bounds : 1 .. 2 }
Y = _h673 { bounds : 9 .. 10 }
```

Writing a handler can be as simple as possible or as elaborate as desired. In the case of `bounds` the handler is simple:

```
bounds_attr_portray_hook(bounds(L,U,_)) :- write(L..U).
```

The hook is installed when the constraint package is loaded by placing in the package loader directive such as:

```
:- install_attribute_portray_hook(bounds,Attr,bounds_attr_portray_hook(Attr)).
```

Note that the hook will be indexed on the module associated with the attribute (in this case `bounds`). XSB's command-line interpreter will unify the second argument of the portray hook with the attribute, and then call `Handler`.

`install_attribute_constraint_hook(Module,Vars,Names,Handler)`   module: machine
For some constraint packages, it may not be particularly useful to associate constraints with variables: instead, the projection of global constraints onto the variables of the top-level query may be more useful. This is the case in the CLP(R) package (Section 11.1), where the command-line interaction may look as follows:

```
| ?- {X = 2*Y,Y >= 7},inf(X,F).
 { X >= 14.0000 }
 { Y = 0.5000 * X }

X = _h8841
Y = _h9506
F = 14.0000
```

In XSB, the (projection of the) global constraints in CLP(R) are displayed by the following routines:

```
clpr_portray_varlist(Vars,Names):-
filter_varlist(Vars,Names,V1,N1),
dump(V1,N1,Constraints),
member(C,Constraints),
console_write(' { '), console_write(C),console_writeln(' } '),
fail.
clpr_portray_varlist(_V,_N).

filter_varlist([],[],[],[]).
filter_varlist([V1|R1],[N1|R2],[V1|R3],[N1|R4]):-
var(V1),!,
filter_varlist(R1,R2,R3,R4).
filter_varlist([_V1|R1],[_N1|R2],R3,R4):-
filter_varlist(R1,R2,R3,R4).
```

This predicate sets up a call to the CLP(R) library predicate `dump/3`, whose constraints it then writes out to the console. Analogous to the portray hook, the console hook is installed using the directive:

```
:- install_constraint_portray_hook(clpr,Vars,Names,clpr_portray_varlist(Vars,Names)).
```

If the `clpr` module is loaded, the command line interpreter checks any constraint portray hooks upon the first success of a top-level goal. It then unifies the second argument `Vars` with the variables of the goal, and `Names` with the names of the variables of the goal which are then passed on to `Handler`

1.3   `constraintLib`: a library for CLP

XSB supports constraint logic programming through its engine-level support of attributed variables (Section 1.2), and its support for constraint handling rules (CHR) (Chapter 12). The `constraintLib` library includes routines for delaying and examining bindings that are commonly used to implement CHR and other constraint libraries.

When processing constraints, it is often useful to delay a goal based on the instantiation level of a term or set of terms. For instance a `3 > X + Y` should be delayed until both `X` and `Y` are instantiated. However the goal should be reinvoked as soon as possible after both are instantiated in order to prune search paths that may not be useful to pursue. The predicate `when/2` provides a useful mechanism to delay goals based on instantiation patterns [1].

`when(+Condition,Goal)`   module: `constraintLib`

> Delays the execution of `Goal` until `Condition` is satisfied, whereupon `Goal` will be executed. Condition can have the form
>
> - `?=(Term1,Term2)`
> - `nonvar(Term)`
> - `ground(Term)` [2]
> - `(Condition,Condition)`
> - `(Condition ; Condition)`
>
> **Example:** The following session illustrates the use of when/2 to delay a goal.
>
> ```
> |?- when(nonvar(X),writeln(test(1-2,nonvar))),writeln(test(1,nonvar)),X = f(_Y).
>
> test(1,nonvar)
> test(1 - 2,nonvar)
>
> X = f(_h245)
> ```

`unifiable(X, Y, -Unifier)`   module: `constraintLib`

> If X and Y can unify, succeeds unifying `Unifier` with a list of terms of the form `Var =`

---

[1]Despite the similar name, this method of delaying is conceptually different from SLG DELAYING discussed in Volume 1 of this manual, which is used for resolving cycles of dependencies in computing the well-founded semantics, and is not based on the state of instantiation of a term.

[2]To use `ground/1` in the condition, it must be imported into the file where it is used.

`Value` representing a most general unifier of `X` and `Y`. `unifiable/3` can handle cyclic terms. Attributed variables are handles as normal variables. Associated hooks are not executed [3].

`setarg(+Index,+Term,+Value)`   module: `constraintLib`

The predicate `setarg/3` provides an efficient but non-logical way to update argument `Index` of a Prolog term `Term` to `Value` via destructive assignment and without the necessity of copying `Term`. `setarg/3` should be used sparingly, to ensure both clarity and portability of code.

Example

```
|?- X = p(f(1),g(2),r([a])),
    writeln(zero(X)),
    (  set_arg(X,2,g([b])),
       writeln(one(X)),
       fail
     ; writeln(two(X))).
zero(p(f(1),g(2),r([a])))
one(p(f(1),g([b]),r([a])))
two(p(f(1),g(2),r([a])))

X = p(f(1),g(2),r([a]))
```

Error Cases

- `Index` is a variable
    - `instantiation_error`
- `Index` neither a variable nor an integer
    - `type_error(integer,Index)`
- `Index` is less than 0
    - `domain_error(not_less_than_zero,Index)`
- `Term` is a variable
    - `instantiation_error`
- `Term` neither a variable nor a compound term
    - `type_error(compound,Term)`

`term_variables(+Term,-Variables`   module: `constraintLib`

Given any Prolog term `Term` as input, returns a sorted list of variables in the term.

1.4   Formatted Output

`format(+String,+Control)`   module: `format`

---

[3]In Version 3.2, `unifiable/3` is written as a Prolog predicate and so is slower than many of the predicates in this section.

`format(+Stream,+String,+Control)`   module: format

>   `format/2` and `format/3` act as a Prolog analog to the C `stdio` function `printf()`, allowing formatted output [4].
>
>   Output is formatted according to `String` which can contain either a format control sequence, or any other character which will appear verbatim in the output. Control sequences act as place-holders for the actual terms that will be output. Thus
>
>>   `?- format("Hello ~q!",world).`
>
>   will print `Hello world!`.
>
>   If there is only one control sequence, the corresponding element may be supplied alone in `Control`. If there are more, `Control` must be a list of these elements. If there are none then `Control` must be an empty list. There have to be as many elements in `Control` as control sequences in `String`.
>
>   The character `~` introduces a control sequence. To print a `~` just repeat it:
>
>>   `?- format("Hello ~~world!", []).`
>
>   will output `Hello ~world!`.
>
>   The general format of a control sequence is `~NC`. The character `C` determines the type of the control sequence. `N` is an optional numeric argument. An alternative form of `N` is `*`. `*` implies that the next argument in `Arguments` should be used as a numeric argument in the control sequence. For example:
>
>   `?- format("Hello~4cworld!", [0'x]).`
>
>   and
>
>   `?- format("Hello~*cworld!", [4,0'x]).`
>
>   both produce
>
>   `Helloxxxxworld!`
>
>   The following control sequences are available in XSB.
>
>   - `~a` The argument is an atom. The atom is printed without quoting.
>   - `~Nc` (Print character.) The argument is a number that will be interpreted as an ASCII code. `N` defaults to one and is interpreted as the number of times to print the character.

---

[4]The `format` family of predicates is due to Quintus Prolog, by way of Ciao.

- ~f (Print float). The argument is a float. The float will be printed out by XSB.

- ~d (Print integer). The argument is an integer, and will be printed out by XSB.

- ~Ns (Print string.) The argument is a list of ASCII codes. Exactly N characters will be printed. N defaults to the length of the string. Example:

  ```
  ?- format("Hello ~4s ~4s!", ["new","world"]).
  ?- format("Hello ~s world!", ["new"]).
  ```

  will print as

  ```
  Hello new worl!
  Hello new world!
  ```

  respectively.

- ~i (Ignore argument.) The argument may be of any type. The argument will be ignored. Example:

  ```
  ?- format("Hello ~i~s world!", ["old","new"]).
  ```

  will print as

  ```
  Hello new world!
  ```

- ~k (Print canonical.) The argument may be of any type. The argument will be passed to write_canonical/2). Example:

  ```
  ?- format("Hello ~k world!", a+b+c).
  ```

  will print as

  ```
  Hello +(+(a,b),c) world!
  ```

- ~q (Print quoted.) The argument may be of any type. The argument will be passed to writeq/2. Example:

  ```
  ?- format("Hello ~q world!", [['A','B']]).
  ```

  will print as

  ```
  Hello ['A','B'] world!
  ```

- ~w (write.) The argument may be of any type. The argument will be passed to write/2. Example:

  ```
  ?- format("Hello ~w world!", [['A','B']]).
  ```

  will print as

  ```
  Hello [A,B] world!
  ```

- ~Nn (Print newline.) Print N newlines. N defaults to 1. Example:

```
?- format("Hello ~n world!", []).
```

will print as

```
Hello
 world!
```

1.5   String Manipulation

XSB has a number of powerful predicates that simplify the job of string manipulation. These predicates are especially powerful when they are combined with pattern-matching facilities provided by the `regmatch` package described in Chapter 7 [5].

str_sub(+Sub, +Str, ?Pos)   module: string

    Succeeds if `Sub` is a substring of `Str`. In that case, `Pos` unifies with the position where the match occurred. Positions start from 0. `str_sub/2` is also available, which is equivalent to having _ in the third argument of `str_sub/3`.

str_match(+Sub, +Str, +Direction, ?Beg, ?End)   module: string

    This is an enhanced version of the previous predicate. `Direction` can be `forward` or `reverse` (or any abbreviation of these). If `forward`, the predicate finds the first match of `Sub` from the beginning of `Str`. If `reverse`, it finds the first match from the end of the string (*i.e.*, the last match of `Sub` from the beginning of `Str`). `Beg` and `End` must be integers or unbound variables. (It is possible that one is bound and another is not.) `Beg` unifies with the offset of the first character where `Sub` matched, and `End` unifies with the offset of the next character to the right of `Sub` (such a character might not exist, but the offset is still defined). Offsets start from 0.

    Both `Beg` and `End` can be bound to negative integers. In this case, the value represents the offset from the *second* character past the end of `Str`. Thus -1 represents the character next to the end of `Str` and can be used to check where the end of `Sub` matches in `Str`. In the following examples

```
?- string_match(Sub,Str,forw,X,-1).
?- string_match(Sub,Str,rev,X,-1).
?- string_match(Sub,Str,forw,0,X).
```

    the first checks if the *first* match of `Sub` from the beginning of `Str` is a suffix of `Str` (because `End` represents the character next to the last character in `Sub`, so `End=-1` means that the last characters of `Sub` and of `Str` occupy the same position). If so, `X` is bound to the offset (from the end of `Str`) of the first character of `Sub`. The second example checks if the *last* match of

---

[5]Not all string manipulation predicates have been made thread-safe in Version 3.2.

Sub in Str is a suffix of Str and binds X to the offset of the beginning of that match (counted from the beginning of Str). The last example checks if the first match of Sub is a prefix of Str. If so, X is bound to the offset (from the beginning of Str) of the last character of Sub.

substring(+String, +BeginOffset, +EndOffset, -Result)   module: string
String can be an atom or a list of characters, and the offsets must be integers. If EndOffset is negative, endof(String)+EndOffset+1 is assumed. Thus, -1 means end of string. If BeginOffset is less than 0, then 0 is assumed; if it is greater than the length of the string, then string end is assumed. If EndOffset is non-negative, but is less than BeginOffset, then empty string is returned.

Offsets start from 0.

The result returned in the fourth argument is a string, if String is an atom, or a list of characters, if so is String.

The substring/4 predicate always succeeds (unless there is an error, such as wrong argument type).

Here are some examples:

```
| ?- substring('abcdefg', 3, 5, L).

L = de

| ?- substring("abcdefg", 4, -1, L).

L = [101,102]
```

(*i.e.*, L = ef represented using ASCII codes).

string_substitute(+InpStr, +SubstrList, +SubstitutionList, -OutStr) module: string

InputStr can an atom or a list of characters. SubstrList must be a list of terms of the form s(BegOffset, EndOffset), where the name of the functor is immaterial. The meaning of the offsets is the same as for substring/4. (In particular, negative offsets represent offsets from the first character past the end of String.) Each such term specifies a substring (between BegOffset and EndOffset; negative EndOffset stands for the end of string) to be replaced. SubstitutionList must be a list of atoms or character lists.

Offsets start from 0, as in C/Java.

This predicate replaces the substrings specified in SubstrList with the corresponding strings from SubstitutionList. The result is returned in OutStr. OutStr is a list of characters, if so is InputStr; otherwise, it is an atom.

If SubstitutionList is shorter than SubstrList then the last string in SubstitutionList is used for substituting the extra substrings specified in SubstitutionList. As a special case, this makes it possible to replace all specified substrings with a single string.

As in the case of re_substring/4, if OutStr is an atom, it is not interned. The user should either intern this string or convert it into a list, as explained previously.

The `string_substitute/4` predicate always succeeds.

Here are some examples:

```
| ?- string_substitute('qaddf', [s(2,4)], ['123'] ,L).

L = qa123f

| ?- string_substitute('qaddf', [s(2,-1)], ['123'] ,L).

L = qa123

| ?- string_substitute("abcdefg", [s(4,-1)], ["123"],L).

L = [97,98,99,100,49,50,51]

| ?- string_substitute('1234567890123', [f(1,5),f(5,7),f(9,-2)], ["pppp", lll],X).

X = 1ppppll89lll

| ?- string_substitute('1234567890123', [f(1,5),f(6,7),f(9,-2)], ['---'],X).

X = 1---6---89---
```

concat_atom(+AtomList,?Atom)   module: string

AtomList must be a list structure containing atoms, integers and/or floats. This predicate flattens `AtomList` and concatenates the atoms and integers into a single atom, returned in `Atom`. Integers and floats are converted to character strings using `number_codes/2`.

concat_atom(+AtomList,+Sep,?Atom)   module: string

AtomList must be a list containing atoms, integers and/or floats, and `Sep` must be an atom. This predicate concatenates the atoms and integers into a single atom, separating each by `Sep`, return the resulting atom in `Atom`. Integers and floats are converted to character strings using `number_codes/2`.

1.6   Script Writing Utilities

Prolog, (in particular XSB!) can be useful for writing scripts. Prolog's simple syntax and declarative semantics make it especially suitable for scripts that involve text processing. There are several ways to access script-writing commands from XSB. The first is to execute the command via the predicates `shell/1` or `shell/2`. These predicates can execute any command but they do not provide streamability across UNIX and Windows commands, and they do not return any output of

commands to Prolog. Special predicates are provided to handle cross-platform compatibility and to bring output into XSB.

Effort has been made to make the these thread-safe; however in Version 3.2, calls to the XSB script writing utilities go through a single mutex, and may cause contention if many threads seek to concurrently use sockets.

`expand_filename(+FileName,-ExpandedName)`   module: machine

Expands the file name passed as the first argument and binds the variable in the second argument to the expanded name. This includes (1) expanding Unix tildes, (2) prepending `FileName` to the current directory, and (3) "rectifying" the expanded file name. In rectification, the expanded file name is "rectified" so that multiple repeated slashes are replaced with a single slash, the intervening "./" are removed, and "../" are applied so that the preceding item in the path name is deleted. For instance, if the current directory is `/home`, then `abc//cde/..///ff/./b` will be converted into `/home/abc/ff/b`.

Under Windows, this predicates does rectification as described above, (using backslashes when appropriate), but it does not expand the tildes.

`expand_filename_no_prepend(+FileName,-ExpandedName)`   module: shell

This predicate behaves as `expand_filename/2`, but only expands tildes and does rectification. It does not prepend the current working directory to relative file names.

`parse_filename(+FileName,-Dir,-Base,-Extension)`   module: machine

This predicate parses file names by separating the directory part, the base name part, and file extension. If file extension is found, it is removed from the base name. Also, directory names are rectified and if a directory name starts with a tilde (in Unix), then it is expanded. Directory names always end with a slash or a backslash, as appropriate for the OS at hand.

For instance, ∼`john///doe/dir1//../foo.bar` will be parsed into: `/home/john/doe/`, `foo`, and `bar` (where we assume that `/home/john` is what ∼`john` expands into).

`sleep(+Seconds)`   module: shell

Put XSB to sleep for a given number of seconds.

Error Cases

- `Seconds` is a variable
    - `instantiation_error`.
- `Seconds` is not an integer
    - `type_error(integer, Seconds)`.

`sys_pid(-Pid)`   module: shell

Get Id of the current process.

`getenv(+VarName,-VarVal`   module: machine

Unifies `VarVal` with the value of `VarName` in the current shell. If `VarName` is not an environment varible, the predicate fails.

Example:

```
:- import getenv/2 from machine.

yes
| ?- getenv('HOSTTYPE',F).

F = intel-pc
```

putenv(+String)   module: machine

If `String` is of the form `VarName=Value`, inserts or resets the environment variable `VarName`. If `VarName` does not exist, it is inserted with `VarVal`. If the `VarName` does exist, it is reset to `VarVal`. `putenv/2` always succeeds.

Exceptions:

`instantiation_error` `String` is not instantiated at the time of call.

`type_error` `VarName` or `VarVal` is not an atom or a list of atoms.

1.6.1   Communication with Subprocesses

In the previous section, we have seen several predicates that allow XSB to create other processes. However, these predicates offer only a very limited way to communicate with these processes. The predicate `spawn_process/5` and friends come to the rescue. It allows a user to spawn any process (including multiple copies of XSB) and redirect its standard input and output to XSB streams. XSB can then write to the process and read from it. The section of socket I/O describes yet another mode of interprocess communication.

In addition, the predicate `pipe_open/2` described in this section lets one create any number of pipes (that do not need to be connected to the standard I/O stream) and talk to child processes through these pipes. All predicates in this section, except `pipe_open/2` and `fd2stream/2`, must be imported from module `shell`. The predicates `pipe_open/2` and `fd2stream/2` must be imported from `file_io`.

spawn_process(+CmdSpec,-StreamToProc,-StreamFromProc,-ProcStderrStream,-ProcId)

Spawn a new process specified by `CmdSpec`. `CmdSpec` must be either a single atom or a *list* of atoms. If it is an atom, then it must represent a shell command. If it is a list, the first member of the list must be the name of the program to run and the other elements must be arguments to the program. Program name must be specified in such a way as to make sure the OS can find it using the contents of the environment variable `PATH`. Also note that pipes, I/O redirection and such are not allowed in command specification. That is, `CmdSpec` must represent a single command. (But read about process plumbing below and about the related predicate `shell/5`.)

The next three parameters of `spawn_process` are XSB I/O stream identifiers for the process (leading to the subprocess standard input), from the process (from its standard output), and a stream capturing the subprocess standard error output. The last parameter is the system process id.

Here is a simple example of how it works.

```
| ?- import file_flush/2, file_read_line_atom/2 from file_io.
| ?- import file_nl/1 , file_write/2 from xsb_writ.

| ?- spawn_process([cat, '-'], To, From, Stderr, Pid),
     writeln(To,'Hello cat!'), flush_output(To,_), file_read_line_atom(From,Y).

To = 3
From = 4
Stderr = 5
Pid = 14328
Y = Hello cat!

yes
```

Here we created a new process, which runs the "`cat`" program with argument "–". This forces
`cat` to read from standard input and write to standard output. The next line writes an atom and
newline to the XSB stream `To`, which is bound to the standard input of the `cat` process (proc id
`14328`). The `cat` process then copies the input to its standard output. Since standard output of the
cat process is redirected to the XSB stream `From` in the parent process, the last line in our program
is able to read it and return in the variable `Y`. Note that in the second line we used `flush_output/2`.
Flushing the output is extremely important here, because XSB I/O pipe (file) streams are buffered.
Thus, `cat` might not see its input until the buffer is filled up, so the above clause might hang.
`flush_output/2` makes sure that the input is immediately available to the subprocess.

In addition to the above general schema, the user can tell `spawn_process/5` not to open one
of the communication streams or to use one of the existing communication streams. This is useful
when you do not expect to write or read to/from the subprocess or when one process wants to write
to another (see the process plumbing example below). To tell that a certain stream is not needed,
it suffices to bind that stream to an atom. For instance,

```
| ?- spawn_process([cat, '-'], To, none, none, _),
     nl(To), writeln(To,'Hello cat!'), flush_output(To).


To = 3,
Hello cat!
```

reads from XSB and copies the result to standard output. Likewise,

```
| ?- spawn_process('cat library.tex', none, From, none, _),
     file_read_line_atom(From, S).

From = 4
S = \chapter{Library Utilities} \label{library_utilities}
```

In each case, only one of the streams is open. (Note that the shell command is specified as an atom rather than a list.) Finally, if both streams are suppressed, then `spawn_process` reduces to the usual `shell/1` call (in fact, this is how `shell/1` is implemented):

```
| ?- spawn_process([pwd], none, none).
```

```
/usr/local/foo/bar
```

On the other hand, if any one of the three stream variables in `spawn_process` is bound to an already existing file stream, then the subprocess will use that stream (see the process plumbing example below).

One of the uses of XSB subprocesses is to create XSB servers that spawn subprocesses and control them. A spawned subprocess can be another XSB process. The following example shows one XSB process spawning another, sending it a goal to evaluate and obtaining the result:

```
| ?- spawn_process([xsb], To, From,Err,_),
     write(To,'assert(p(1)).'), flush_output(To,_),
     write(To,'p(X), writeln(X).'), flush_output(To,_),
     file_read_line_atom(From,XX).
```

```
XX = 1
```

```
yes
| ?-
```

Here the parent XSB process sends "`assert(p(1)).`" and then "`p(X), writeln(X).`" to the spawned XSB subprocess. The latter evaluates the goal and prints (via "`writeln(X)`") to its standard output. The main process reads it through the `From` stream and binds the variable `XX` to that output.

Finally, we should note that the stream variables in the `spawn_process` predicate can be used to do process plumbing, *i.e.*, redirect output of one subprocess into the input of another. Here is an example:

```
| ?- open(test,write,Stream),
     spawn_process([cat, 'data'], none, FromCat1, none, _),
     spawn_process([sort], FromCat1,Stream, none, _).
```

Here, we first open file `test`. Then `cat data` is spawned. This process has the input and standard error stream blocked (as indicated by the atom `none`), and its output goes into stream `FromCat1`. Then we spawn another process, `sort`, which picks the output from the first process (since it uses the stream `FromCat1` as its input) and sends its own output (the sorted version of `data`) to its output stream `Stream`. However, `Stream` has already been open for output into the file `test`. Thus, the overall result of the above clause is tantamount to the following shell command:

```
        cat data | sort > test
```

***Important notes about spawned processes*:**

1. Asynchronous processes spawned by XSB do not disappear (at least on Unix) when they terminate, *unless* the XSB program executes a *wait* on them (see `process_control` below). Instead, such processes become defunct *zombies* (in Unix terminology); they do not do anything, but consume resources (such as file descriptors). So, when a subprocess is known to terminate, it must be waited on.

2. The XSB parent process must know how to terminate the asynchronous subprocesses it spawns. The drastic way is to kill it (see `process_control` below). Sometimes a subprocess might terminate by itself (*e.g.*, having finished reading a file). In other cases, the parent and the child programs must agree on a protocol by which the parent can tell the child to exit. The programs in the XSB subdirectory `examples/subprocess` illustrate this idea. If the child subprocess is another XSB process, then it can be terminated by sending the atom `end_of_file` or `halt` to the standard input of the child. (For this to work, the child XSB must waiting at the prompt).

3. It is very important to not forget to close the streams that the parent uses to communicate with the child. These are the streams that are provided in arguments 2,3,4 of `spawn_process`. The reason is that the child might terminate, but these streams to the standard input of the child will remain open, since they belong to the parent process. As a result, the parent will own defunct I/O streams and might eventually run out of file descriptors or streams.

`process_status(+Pid,-Status)`
> This predicate always succeeds. Given a process id, it binds the second argument (which must be an unbound variable) to one of the following atoms: `running`, `stopped`, `exited_normally`, `exited_abnormally`, `aborted`, `invalid`, and `unknown`. The `invalid` status is given to processes that never existed or that are not children of the parent XSB process. The `unknown` status is assigned when none of the other statuses can be assigned.

> Note: process status (other than `running`) is system dependent. Windows does not seem to support `stopped` and `aborted`. Also, processes killed using the `process_control` predicate (described next) are often marked as `invalid` rather than `exited`, because Windows seems to lose all information about such processes. Process status might be inaccurate in some Unix systems as well, if the process has terminated and `wait()` has been executed on that process.

`process_control(+Pid,+Operation)`
> Perform a process control `operation` on the process with the given `Pid`. Currently, the only supported operations are `kill` (an atom) and `wait(Code)` (a term). The former causes the process to exit unconditionally, and the latter waits for process completion. When the process exits, `Code` is bound to the process exit code. The code for normal termination is 0.

> This predicate succeeds, if the operation was performed successfully. Otherwise, it fails. The `wait` operation fails if the process specified in `Pid` does not exist or is not a child of the parent XSB process.

> The `kill` operation might fail, if the process to be killed does not exist or if the parent XSB process does not have the permission to terminate that process. Unix and Windows have

different ideas as to what these permissions are. See *kill(2)* for Unix and *TerminateProcess* for Windows.

Note: under Windows, the programmer's manual warns of dire consequences if one kills a process that has DLLs attached to it.

get_process_table(-ProcessList)   module: shell

This predicate is imported from module `shell`. It binds `ProcessList` to the list of terms, each describing one of the active XSB subprocesses (created via `spawn_process/5`). Each term has the form:

process(Pid,ToStream,FromStream,StderrStream,CommandLine).

The first argument in the term is the process id of the corresponding process, the next three arguments describe the three standard streams of the process, and the last is an atom that shows the command line used to invoke the process. This predicate always succeeds.

shell(+CmdSpec,-StreamToProc, -StreamFromProc, -ProcStderr, -ErrorCode)

The arguments of this predicate are similar to those of `spawn_process`, except for the following: (1) The first argument is an atom or a list of atoms, like in `spawn_process`. However, if it is a list of atoms, then the resulting shell command is obtained by string concatenation. This is different from `spawn_process` where each member of the list must represent an argument to the program being invoked (and which must be the first member of that list). (2) The last argument is the error code returned by the shell command and not a process id. The code -1 and 127 mean that the shell command failed.

The `shell/5` predicate is similar to `spawn_process` in that it spawns another process and can capture that process' input and output streams. The important difference, however, is that XSB will ways until the process spawned by `shell/5` terminates. In contrast, the process spawned by `spawn_process` will run concurrently with XSB. In this latter case, XSB must explicitly synchronize with the spawned subprocess using the predicate `process_control/2` (using the `wait` operation), as described earlier.

The fact that XSB must wait until `shell/5` finishes has a very important implication: the amount of data the can be sent to and from the shell command is limited (1K is probably safe). This is because the shell command communicates with XSB via pipes, which have limited capacity. So, if the pipe is filled, XSB will hang waiting for `shell/5` to finish and `shell/5` will wait for XSB to consume data from the pipe. Thus, use `spawn_process/5` for any kind of significant data exchange between external processes and XSB.

Another difference between these two forms of spawning subprocesses is that `CmdSpec` in `shell/5` can represent *any* shell statement, including those that have pipes and I/O redirection. In contrast, `spawn_process` only allows command of the form "program args". For instance,

```
| ?- open(test,write,Stream),
     shell('cat | sort > data', Stream, none, none, ErrCode)
```

As seen from this example, the same rules for blocking I/O streams apply to `shell/5`. Finally, we should note that the already familiar standard predicates `shell/1` and `shell/2` (documented in Volume 1) are implemented using `shell/5`, and `shell/5` shares their error cases.

Notes:

1. With `shell/5`, you do not have to worry about terminating child processes: XSB waits until the child exits automatically. However, since communication pipes have limited capacity, this method can be used only for exchanging small amounts of information between parent and child.

2. The earlier remark about the need to close I/O streams to the child *does* apply.

`pipe_open(-ReadPipe, -WritePipe)`
Open a new pipe and return the read end and the write end of that pipe. If the operation fails, both `ReadPipe` and `WritePipe` are bound to negative numbers. The pipes returned by the `pipe_open/2` predicate are small integers that represent file descriptors used by the underlying OS. They are **not XSB I/O streams**, and they cannot be used for I/O directly. To use them, one must convert them to streams using `open/3` or `open/4`. [6]

The best way to illustrate how one can create a new pipe to a child (even if the child has been created earlier) is to show an example. Consider two programs, `parent.P` and `child.P`. The parent copy of XSB consults `parent.P`, which does the following: First, it creates a pipe and spawns a copy of XSB. Then it tells the child copy of XSB to assert the fact `pipe(RP)`, where `RP` is a number representing the read part of the pipe. Next, the parent XSB tells the child XSB to consult the program `child.P`. Finally, it sends the message `Hello!`.

The `child.P` program gets the pipe from predicate `pipe/1` (note that the parent tells the child XSB to first assert `pipe(RP)` and only then to consult the `child.P` file). After that, the child reads a message from the pipe and prints it to its standard output. Both programs are shown below:

```
%% parent.P
:- import pipe_open/2 from file_io.
%% Create the pipe and pass it to the child process
?- pipe_open(RP,WP),
   %% WF is now the XSB I/O stream bound to the write part of the pipe
   open(pipe(WP),write,WF),
   %% ProcInput becomes the XSB stream leading directly to the child's stdin
   spawn_process(nxsb1, ProcInput, block, block, Process),
   %% Tell the child where the reading part of the pipe is
   fmt_write(ProcInput, "assert(pipe(%d)).\n", arg(RP)),
   fmt_write(ProcInput, "[child].\n", _),
```

---

[6] XSB does not convert pipe file descriptors into I/O streams automatically. Because of the way XSB I/O streams are represented, they are not inherited by the child process and they do not make sense to the child process (especially if the child is not another XSB process). Therefore, we must pass the child processes an OS file descriptor instead. The child then converts these descriptor into XSB I/O streams.

```
        flush_output(ProcInput, _),
        %% Pass a message through the pipe
        fmt_write(WF, "Hello!\n", _),
        flush_output(WF, _),
        fmt_write(ProcInput, "end_of_file.\n",_), % send end_of_file atom to child
        flush_output(ProcInput, _),
        %% wait for child (so as to not leave zombies around;
        %% zombies quit when the parent finishes, but they consume resources)
        process_control(Process, wait),
        %% Close the ports used to commuicate with the process
        %% Otherwise, the parent might run out of file descriptors
        %% (if many processes were spawned)
        close(ProcInput), close(WF).



%% child.P
:- import file_read_line_atom/2 from file_io.
:- dynamic pipe/1.
?- pipe(P), open(pipe(P),read,F),
        %% Acknowledge receipt of the pipe
        fmt_write("\nPipe %d received\n", arg(P)),
        %% Get a message from the parent and print it to stdout
        file_read_line_atom(F, Line), write('Message was: '), writeln(Line).
```

This produces the following output:

```
| ?- [parent].                      <- parent XSB consults parent.P
[parent loaded]
yes
| ?- [xsb_configuration loaded]   <- parent.P spawns a child copy of XSB
[sysinitrc loaded]                      Here we see the startup messages of
[packaging loaded]                       the child copy
XSB Version 2.0 (Gouden Carolus) of June 27, 1999
[i686-pc-linux-gnu; mode: optimal; engine: slg-wam; scheduling: batched]
| ?-
yes
| ?- [Compiling ./child]          <- The child copy of received the pipe from
[child compiled, cpu time used: 0.1300 seconds]      the parent and then the
[child loaded]                                        request to consult child.P
Pipe 15 received                     <- child.P acknowledges receipt of the pipe
Message was: Hello!                  <- child.P gets the message and prints it
yes
```

Observe that the parent process is very careful about making sure that the child terminates and also about closing the I/O streams after they are no longer needed.

Finally, we should note that this mechanism can be used to communicate through pipes with non-XSB processes as well. Indeed, an XSB process can create a pipe using `pipe_open` (*before* spawning a child process), pass one end of the pipe to a child process (which can be a C program), and use `open/3` to convert the other end of the pipe to an XSB stream. The C program, of course, does not need `open/3`, since it can use the pipe file handle directly. Likewise, a C program can spawn off an XSB process and pass it one end of a pipe. The XSB child-process can then convert this pipe fd to a file using `fd2iostream` and then talk to the parent C program.

`fd2iostream(+Pipe, -IOstream)`
> Take a file descriptor and convert it to an XSB I/O stream. This predicate should be used only for user-defined I/O. Otherwise, use `open/{3,4}` when possible.

1.7 Socket I/O

The XSB socket library defines a number of predicates for communication over BSD-style sockets. Most are modeled after and are interfaces to the socket functions with the same name. For detailed information on sockets, the reader is referred to the Unix man pages (another good source is *Unix Network Programming*, by W. Richard Stevens). Several examples of the use of the XSB sockets interface can be found in the `XSB/examples/` directory in the XSB distribution.

XSB supports two modes of communication via sockets: *stream-oriented* and *message-oriented*. In turn, stream-oriented communication can be *buffered* or *character-at-a-time*.

To use *buffered* stream-oriented communication, system socket handles must be converted to XSB I/O streams using `fd2iostream/2`. In these stream-oriented communication, messages have no boundaries, and communication appears to the processes as reading and writing to a file. At present, buffered stream-oriented communication works under Unix only.

Character-at-a-time stream communication is accomplished using the primitives `socket_put/3` and `socket_get0/3`. These correspond to the usual Prolog `put/1` and `get0/1` I/O primitives.

In message-oriented communication, processes exchange messages that have well-defined boundaries. The communicating processes use `socket_send/3` and `socket_recv/3` to talk to each other. XSB messages are represented as strings where the first four bytes (`sizeof(int)`) is an integer (represented in the binary network format — see the functions `htonl` and `ntohl` in socket documentation) and the rest is the body of the message. The integer in the header represents the length of the message body.

Effort has been made to make the socket interface thread-safe; however in Version 3.2, calls to the XSB socket interface go through a single mutex, and may cause contention if many threads seek to concurrently use sockets.

We now describe the XSB socket interface. All predicates below must be imported from the module `socket`. Note that almost all predicates have the last argument that unifies with the error

code returned from the corresponding socket operation. This argument is explained separately.

General socket calls. These are used to open/close sockets, to establish connections, and set special socket options.

`socket(-Sockfd, ?ErrorCode)`
> A socket `Sockfd` in the AF_INET domain is created. (The AF_UNIX domain is not yet implemented). `Sockfd` is bound to a small integer, called socket descriptor or socket handle.

`socket_set_option(+Sockfd,+OptionName,+Value)`
> Set socket option. At present, only the `linger` option is supported. "Lingering" is a situation when a socket continues to live after it was shut down by the owner. This is used in order to let the client program that uses the socket to finish reading or writing from/to the socket. `Value` represents the number of seconds to linger. The value -1 means do not linger at all.

`socket_close(+Sockfd, ?ErrorCode)`
> `Sockfd` is closed. Sockets used in `socket_connect/2` should not be closed by `socket_close/1` as they will be closed when the corresponding stream is closed.

`socket_bind(+Sockfd,+Port, ?ErrorCode)`
> The socket `Sockfd` is bound to the specified local port number.

`socket_connect(+Sockfd,+Port,+Hostname,?ErrorCode)`
> The socket Sockfd is connected to the address (`Hostname` and `Port`). If `socket_connect/4` terminates abnormally for any reason (connection refused, timeout, etc.), then XSb closes the socket `Sockfd` automatically, because such a socket cannot be used according to the BSD semantics. Therefore, it is always a good idea to check to the return code and reopen the socket, if the error code is not `SOCK_OK`.

`socket_listen(+Socket, +Length, ?ErrorCode)`
> The socket `Sockfd` is defined to have a maximum backlog queue of `Length` pending connections.

`socket_accept(+Sockfd,-SockOut, ?ErrorCode)`
> Block the caller until a connection attempt arrives. If the incoming queue is not empty, the first connection request is accepted, the call succeeds and returns a new socket, `SockOut`, which can be used for this new connection.

Buffered, message-based communication. These calls are similar to the `recv` and `send` calls in C, except that XSB wraps a higher-level message protocol around these low-level functions. More precisely, `socket_send/3` prepends a 4-byte field to each message, which indicates the length of the message body. When `socket_recv/3` reads a message, it first reads the 4-byte field to determine the length of the message and then reads the remainder of the message.

All this is transparent to the XSB user, but you should know these details if you want to use these details to communicate with external processes written in C and such. All this means that these external programs must implement the same protocol. The subtle point here is that different

machines represent integers differently, so an integer must first be converted into the machine-independent network format using the functions `htonl` and `ntohl` provided by the socket library. For instance, to send a message to XSB, one must do something like this:

```
char *message, *msg_body;
unsigned int msg_body_len, network_encoded_len;

  msg_body_len = strlen(msg_body);
  network_encoded_len = (unsigned int) htonl((unsigned long int) msg_body_len);
  memcpy((void *) message, (void *) &network_encoded_len, 4);
  strcpy(message+4, msg_body);
```

To read a message sent by XSB, one can do as follows:

```
int actual_len;
char lenbuf[4], msg_buff;
unsigned int msglen, net_encoded_len;

  actual_len = (long)recvfrom(sock_handle, lenbuf, 4, 0, NULL, 0);
  memcpy((void *) &net_encoded_len, (void *) lenbuf, 4);
  msglen = ntohl(net_encoded_len);

  msg_buff = calloc(msglen+1, sizeof(char))); // check if this suceeded!!!
  recvfrom(sock_handle, msg_buff, msglen, 0, NULL, 0);
```

If making the external processes follow the XSB protocol is not practical (because you did not write these programs), then you should use the character-at-a-time interface or, better, the buffered stream-based interface both of which are described in this section. At present, however, the buffered stream-based interface does not work on Windows.

socket_recv(+Sockfd,-Message, ?ErrorCode)
> Receives a message from the connection identified by the socket descriptor `Sockfd`. Binds `Message` to the message. `socket_recv/3` provides a message-oriented interface. It understands message boundaries set by `socket_send/3`.

socket_send(+Sockfd,+Message, ?ErrorCode)
> Takes a message (which must be an atom) and sends it through the connection specified by `Sockfd`. `socket_send/3` provides message-oriented communication. It prepends a 4-byte header to the message, which tells `socket_recv/3` the length of the message body.

Stream-oriented, character-at-a-time interface. Internally, this interface uses the same `sendto` and `recvfrom` socket calls, but they are executed for each character separately. This interface is appropriate when the message format is not known or when message boundaries are determined using special delimiters.

`socket_get0/3` creates the end-of-file condition when it receives the end-of-file character `CH_EOF_P` (a.k.a. 255) defined in `char_defs.h` (which must be included in the XSB program). C programs that need to send an end-of-file character should send `(char)-1`.

`socket_get0(+Sockfd, -Char, ?ErrorCode)`
> The equivalent of `get0` for sockets.

`socket_put(+Sockfd, +Char, ?ErrorCode)`
> Similar to put/1, but works on sockets.

Socket-probing. With the help of the predicate `socket_select/6` one can establish a group of asynchronous or synchronous socket connections. In the synchronous mode, this call is blocked until one of the sockets in the group becomes available for reading or writing, as described below. In the asynchronous mode, this call is used to probe the sockets periodically, to find out which sockets have data available for reading or which sockets have room in the buffer to write to.

The directory `XSB/examples/socket/select/` has a number of examples of the use of the socket-probing calls.

`socket_select(+SymConName,+Timeout,-ReadSockL,-WriteSockL,-ErrSockL,?ErrorCode)`
> `SymConName` must be an atom that denotes an existing connection group, which must be previously created with `socket_set_select/4` (described below). `ReadSockL`, `WriteSockL`, `ErrSockL` are lists of socket handles (as returned by `socket/2`) that specify the available sockets that are available for reading, writing, or on which exception conditions occurred. `Timeout` must be an integer that specifies the timeout in seconds (0 means probe and exit immediately). If `Timeout` is a variable, then wait indefinitely until one of the sockets becomes available.

`socket_set_select(+SymConName,+ReadSockFdLst,+WriteSockFdLst,+ErrorSockFdLst)`
> Creates a connection group with the symbolic name `SymConName` (an atom) for subsequent use by `socket_select/6`. `ReadSockFdLst`, `WriteSockFdLst`, and `ErrorSockFdLst` are lists of sockets for which `socket_select/6` will be used to monitor read, write, or exception conditions.

`socket_select_destroy(+SymConName)`
> Destroys the specified connection group.

Error codes. The error code argument unifies with the error code returned by the corresponding socket commands. The error code -2 signifies *timeout* for timeout-enabled primitives (see below). The error code of zero signifies normal termination. Positive error codes denote specific failures, as defined in BSD sockets. When such a failure occurs, an error message is printed, but the predicate succeeds anyway. The specific error codes are part of the socket documentation. Unfortunately, the symbolic names and error numbers of these failures are different between Unix compilers and Visual C++. Thus, there is no portable, reliable way to refer to these error codes. The only reliably portable error codes that can be used in XSB programs defined through these symbolic constants:

```
#include "socket_defs_xsb.h"

#define SOCK_OK        0      /* indicates sucessful return from socket    */
#define SOCK_EOF      -1      /* end of file in socket_recv, socket_get0   */


#include "timer_defs_xsb.h"

#define TIMEOUT_ERR -2                    /* Timeout error code */
```

Timeouts.   XSB socket interface allows the programer to specify timeouts for certain operations. If the operations does not finish within the specified period of time, the operation is aborted and the corresponding predicate succeeds with the `TIMEOUT_ERR` error code. The following primitives are timeout-enabled: `socket_connect/4`, `socket_accept/3`, `socket_recv/3`, `socket_send/3`, `socket_get0/3`, and `socket_put/3`. To set a timeout value for any of the above primitives, the user should execute `set_timer/1` right before the subgoal to be timed. Note that timeouts are disabled after the corresponding timeout-enabled call completes or times out. Therefore, one must use `set_timer/1` before each call that needs to be controlled by a timeout mechanism.

The most common use of timeouts is to either abort or retry the operation that times out. For the latter, XSB provides the `sleep/1` primitive, which allows the program to wait for a few seconds before retrying.

The `set_timer/1` and `sleep/1` primitives are described below. They are standard predicates and do not need to be explicitly imported.

`set_timer(+Seconds)`
> Set timeout value. If a timer-enabled goal executes after this value is set, the clock begins ticking. If the goal does not finish in time, it succeeds with the error code set to `TIMEOUT_ERR`. The timer is turned off after the goal executes (whether timed out or not and whether it succeeds or fails). This goal always succeeds.
>
> Note that if the timer is not set, the timer-enabled goals execute "normally," without timeouts. In particular, they might block (say, on `socket_recv`, if data is not available).

`sleep(+Seconds)`
> Put XSB to sleep for the specified number of seconds. Execution resumes after the `Seconds` number of seconds. This goal always succeeds.

Here is an example of the use of the timer:

```
:- compiler_options([xpp_on]).
#include "timer_defs_xsb.h"

?- set_timer(3),  % wait for 3 secs
   socket_recv(Sockfd, Msg, ErrorCode),
   (ErrorCode == TIMEOUT_ERR
   -> writeln('Socket read timed out, retrying'),
      try_again(Sockfd)
   ;  write('Data received: '), writeln(Msg)
   ).
```

Apart from the above timer-enabled primitives, a timeout value can be given to `socket_select/6` directly, as an argument.

Buffered, stream-oriented communication. In Unix, socket descriptors can be "promoted" to file streams and the regular read/write commands can be used with such streams. In XSB, such promotion can be done using the following predicate:

`fd2ioport(+Pipe, -IOport)`   module: shell
> Take a socket descriptor and convert it to an XSB I/O port that can be used for regular file I/O.

Once `IOport` is obtained, all normal I/O primitives can be used by specifying the `IOport` as their first argument. This is, perhaps, the easiest and the most convenient way to use sockets in XSB. (This feature has not been implemented for Windows.)

Here is an example of the use of this feature:

```
:- compiler_options([xpp_on]).
#include "socket_defs_xsb.h"

?- (socket(Sockfd, SOCK_OK)
   ->   socket_connect(Sockfd1, 6020, localhost, Ecode),
        (Ecode == SOCK_OK
        -> fd2ioport(Sockfd, SockIOport),
           file_write(SockIOport, 'Hello Server!')
        ;  writeln('Can''t connect to server')
        ),
   ;   writeln('Can''t open socket'), fail
   ).
```

1.8   Arrays

The module `array1` provides a simple backtrackable array implementation that requires no copying. In Version 3.2, this package was changed to make use of the backtrackable destructive assignment made possible by `setarg/3`. We note that as of Version 3.2 this library provides simple syntactic sugar for `functor/3`, `arg/3` and `setarg/3` and relies on error messages for these predicates.

`array_new(-Array,+Size)`   module: array
> Creates a one dimensional empty array of size `Size`. All the elements of this array are variables.

`array_elt(+Array, +Index, ?Element)`   module: array
> Succeeds iff `Element` unifies with the `Index`-th element of array `Array`.

`array_update(+Array, +Index, +Elem)`   module: array
> Updates the array `Array` such that the `Index`-th element of the new array is `Elem` using destructive assignement. The implementation is quite efficient in that it avoids the copying of the entire array.

The following example shows the use of these predicates:

```
| ?- import array_new/2, array_elt/3, array_update/4 from array.

yes
| ?- array_new(A,3), array_update(A,1,1), array_update(A,2,2),
     ( array_update(A,3,3), writeln(first(A))
     ; array_update(A,3,6), writeln(second(A))
     ; array_update(A,3,7), writeln(third(A))),fail.

first(array(1,2,3))
second(array(1,2,6))
third(array(1,2,7))

no
```

1.9   The Profiling Library

XSB can provide Prolog-level profiling for Prolog programs, which allows the Prolog programmer to determine what proportion of time is spent executing code for each predicate. To enable profiling, XSB must be started with the command line parameter of -p. The module `xsb_profiling` contains the predicate `profile_call/1` that invokes profiling [7].

`profile_call(+Goal)`   module: xsb_profiling

---

[7]The profiling library should only be used with the single-threaded engine in Version 3.2.

Calls the goal `Goal`, and when it first returns, prints out (on userout) a table of predicate names indicating for each, the percentage of time spent executing that predicate's code. XSB must have been started with the -p command line parameter to obtain this table. The percentages are accumulated by module. `Goal` may backtrack, but profiling is done only for the time to the first success. So it is most appropriate to profile succeeding deterministic goals.

Profiling works by starting another thread that interrupts every 100th of a second and sets a flag so that the XSB emulator will determine the predicate of the currently executing code. The printout also includes the total number of interrupts and for each predicate, the raw number of times its code was determined to be executing. A predicate is printed only if its code was interrupted at least once. The numbers will be meaningful only for relatively long-running predicates, taking more than a couple of seconds.

When an interrupt occurs, the **next** WAM call, execute, proceed or trust instruction to be executed will log its associated predicate. The system does not keep track of code addresses for tries (used to represent the results of completed tables, and trie-indexed asserted code), so for some interrupts the associated executing predicate cannot be determined. In these cases the interrupt is charged against an "unknown/?" pseudo-predicate, and this count is included in the output.

Profiling does not give the context from which the predicate is called, so you may want to make renamed copies of basic predicates to use in particular circumstances to determine their times.

Predicates compiled with the "optimize" option may provide misleading results under profiling. Note that all system predicates (including those in `basics`) are compiled with the "optimize" option, by default. That option causes tail-recursive predicates to use a "jump" instruction rather than an "execute" instruction to make the recursive call, and so an interrupt in such a loop will not be charged until the next **logging** instruction is executed. If much time is spent in the recursion, this might not be for a long time, and the interrupt might ultimately be charged to another predicate. (If an interrupt has not been charged by the time of the next interrupt, it is lost.)

Profiling currently works under Windows, and somewhat under Linux. For the profiling algorithm to work, the thread that wakes and sets the interrupt flag must be of high priority and given the CPU when it wants it. I haven't figured out how to get this scheduling on some machines, so if you want profiling to work somewhere it doesn't, maybe you can help figure out how to get appropriate scheduling.

The profiling module also provides some support for determining the modes in which a predicate is called. This can be used to determine whether indexing is appropriately declared for important dynamic predicates. This is a primitive facility, but has been found useful.

`profile_mode_call(+Goal)`   module: `xsb_profiling`

Calls the goal `Goal` and constructs a table of the modes in which the predicate is called and the number of times it is called in that mode. Modes are simply "b" for ground and "f" for variable. Counts are kept in a table with entries of the form `Pred(Md1,Md2,..,Mdn)` where Pred is the name of the called predicate and the Mdi are either 'f' or 'b', indicating free or

bound for the corresponding argument. The table can be printed using `profile_mode_dump/0` and can be cleared using `profile_mode_init/0`.

`profile_mode_dump`   module: `xsb_profiling`

   Prints out the counts of calls in particular modes as accumulated using `profile_mode_call(+Goal)`.

`profile_mode_init`   module: `xsb_profiling`

   Clears the table that accumulates counts of calls in particular modes (done by `profile_mode_call(+Goal)`.

1.10   Gensym

The Gensym library provides a convenient way to generate unique integers or constants.

`prepare(+Index)`   module: `gensym`
   Sets the initial integer to be used for generation to `Index`. Thus, the command `?- prepare(0)` would cause the first call to `gennum/1` to return `1`. `Index` must be a non-negative integer.

`gennum(-Var)`   module: `gensym`
   Unifies `Var` with a new integer.

`gensym(+Atom,-Var)`   module: `gensym`
   Generates a new integer, and concatenates this integer with `Atom`, unifying the result with `Var`. For instance a call `?- gensym(foo,Var)` might unify `Var` with `foo32`.

1.11   Random Number Generator

The following predicates are provided in module `random` to generate random numbers (both integers and floating numbers), based on the Wichmann-Hill Algorithm [23, 17]. The random number generator is entirely portable, and does not require any calls to the operating system. As noted below, it does require 3 seeds, each of which must be an integer in a given range. These seeds are thread-specific: thus different threads may generate independent sequences of random numbers.

`random(-Number)`   module: `random`
   Binds `Number` to a random float in the interval [0.0, 1.0). Note that 1.0 will never be generated.

`random(+Lower,+Upper,-Number)`   module: `random`
   Binds `Number` to a random integer in the interval [`Lower`,`Upper`) if `Lower` and `Upper` are integers. Otherwise `Number` is bound to a random float between `Lower` and `Upper`. `Upper` will never be generated.

`getrand(?State)`   module: random
>   Tries to unify `State` with the term `rand(X,Y,Z)` where `X`,`Y`,and `Z` are integers describing the state of the random generator.

`setrand(rand(+X,+Y,+Z))`   module: random
>   Sets the state of the random generator. `X`,`Y`, and `Z` must be integers in the ranges [1,30269), [1,30307), [1,30323), respectively.

`datime_setrand`   module: random
>   This simple initialization utility sets the random seed triple based on a function of the current day, hour, minute and second.

`randseq(+K, +N, -RandomSeq)`   module: random
>   Generates a sequence of `K` unique integers chosen randomly in the range from 1 to `N`. `RandomSeq` is not returned in any particular order.

`randset(+K, +N, -RandomSet)`   module: random
>   Generates an ordered set of `K` unique integers chosen randomly in the range from 1 to `N`. The set is returned in reversed order, with the largest element first and the smallest last.

`gauss(-G1,-G2)`   module: ramdom
>   Generates two random numbers that are normally distributed with mean 0 and standard deviation 1. It uses the polar form of the Box-Muller transformation [5] of uniform random variables as generated by `random/1`.

`weibull(K,Lambda,X)`   module: ramdom
>   Generates a random number for the Weibull distribution:

$$f(x; k, \lambda) = \frac{k}{\lambda}(\frac{x}{\lambda})^{k-1} e^{-(x/\lambda)^h}$$

>   based on the transformation

$$x = \lambda(-ln(U))^{1/k}$$

>   of a uniformly distributed ranom variable produce by `random/1`

`exponential(K,X)`   module: ramdom
>   Generates a random number for the exponential distribution:

$$f(x; k, \lambda) = \frac{e^{-(x/\lambda)^h}}{\lambda}$$

>   based on the transformation

$$x = \lambda(-ln(U))$$

>   of a uniformly distributed ranom variable produce by `random/1`. This is the same as the Weibull distribution with $k = 1$.

1.12 Scanning in Prolog

Scanners, (sometimes called tokenizers) take an input string, usually in ASCII or similar format, and produce a scanned sequence of tokens. The requirements that various applications have for scanning differ in small but important ways – a character that is special to one application may be part of the token of another; or some applications may want lower case text converted to uppercase test. The `stdscan.P` library provides a simple scanner written in XSB that can be configured in several ways. While useful, this scanner is not intended to be as powerful as general-purpose scanners such as *lex* or *flex*.

`scan(+List,-Tokens)`   module: stdscan

> Given as input a `List` of character codes, `scan/2` scans this list producing a list of atoms constituting the lexical tokens. Its parameters are set via `set_scan_pars/1`.

> Tokens produced are either a sequence of *letters* and/or *numbers* or consist of a single *special character* (e.g. `(` or `)`). Whitespaces may occur between tokens.

`scan(+List,+FieldSeparator,-Tokens)`   module: stdscan

> Given as input a `List` of character codes, along with a character code for a field separator, `scan/3` scans this list producing a list of list of atoms constituting the lexical tokens in each field. *scan/3* thus can be used to scan tabular information. Its parameters are set via `set_scan_pars/1`.

`set_scan_pars(+List)`   module: stdscan

> `set_scan_pars(+List)` is used to configure the tokenizer to a particular need. `List` is a list of parameters including the following:

> - `whitespace`. The default action of the scanner is to return a list of tokens, with any whitespace removed. If `whitespace` is a parameter, then the scanner returns the token `''` when it finds whitespace separating two tokens (unless the two tokens are letter sequences; since two letter sequences can be two tokens ONLY if they are separated by whitespace, such an indication of whitespace would be redundant.) Including the parameter `no_whitespace` undoes the effect of previously including `whitespace`.
> - `upper_case` The default action of the parser is to treat lowercase letter differently from uppercase letters. This parameter should be set if conversion to uppercase should be done when producing a token that does *not* consist entirely of letters (e.g. one with mixed letters and digits). Including the parameter `no_case` undoes the effect of previously including `upper_case`.
> - `upper_case_in_lit` The default action of the parser is to treat lowercase letter differently from uppercase letters. This parameter should be set if conversion to uppercase should be done when producing a token that consists entirely of letters. Including the parameter `no_case_in_lit` undoes the effect of previously including `upper_case`.

- `whitespace(Code)` adds `Code` as a whitespace code. By default, all ASCII codes less than or equal to 32 are regarded as whitespace.

- `letter(Code)` adds `Code` as a letter constituting a token. By default, ASCII codes for characters `a--z` and `A--Z` are regarded as letters.

- `special_char(Code)` adds `Code` as a special character. By default, ASCII codes for the following characters are regarded as special characters:

  `| { } [ ] " $ % & ' ( ) * + , - . / : ; < = > ? @ \ ^ _ ~ '`

`get_scan_pars(-List)`   module: stdscan

> `get_scan_pars/1` returns a list of the currently active parameters.

1.13   XSB Lint

The `xsb_lint_impexp.P` file contains a simple tool to analyze import/exports and definitions and uses of predicates. It tries to find possible inconsistencies, producing warnings when it finds them and generating `document_import`/`document_export` declarations that might be useful. It can be used after a large multi-file, multi-module XSB program has been written to find possible inconsistencies in (or interesting aspects of) how predicates are defined and used.

XSB source files that contain an `export` compiler directive are considered as modules. Predicates defined in modules, but not exported, are local to that module. When compiling a module, the XSB compiler generates useful warnings when predicates are used but not defined or defined but not used. All predicates that are defined in source files that do not contain an `export` directive are compiled to be defined in a global module, called `usermod`, and no warning messages are generated. The user may add `document_export` and `document_import` compiler directives (exactly analogous to the `export` and `import` directives) to non-module source files. These directives are ignored by the compiler for its compilation, but cause the define-use analysis to be done and any warning messages to be issued, if appropriate. This allows a user to get the benefit of the define-use analysis without using modules. (See Volume 1, Chapter 3 for more details.)

The `xsb_lint_impexp` utility processes both modules and regular XSB source files that contain `document_export` statements. `xsb_lint_impexp` is not itself a module. To use it, [`xsb_lint_impexp`] must be consulted, which will define the `checkImpExps/{1,2}` and `add_libraries/1` predicates in `usermod`.

`add_libraries(+DirectoryNameList)`

> `add_libraries/1` takes a list of directory names and adds them to the `library_directory/1` predicate. This causes the XSB system to look for XSB source code files in these directories. To use `checkImpExps/{1,2}`, all the directories that contain files (or modules) referenced (recursively) in the files to be processed must be in the `library_directory/1` predicate. This predicate can be used to add a number of directories at once.

`checkImpExps(+Options,+FileNameList)`

> `checkImpExps/1` reads all the XSB source files named in the list `FileNameList`, and all files

they reference (recursively), and produces a listing that describes properties of how they reference predicates.

`Options` is a list of atoms (from the following list) indicating details of how `checkImpExps` should work.

1. `used_elsewhere`: Print a warning message in the case of a predicate defined in a file, not used there, but used elsewhere (in a file in `FileNameList`). This can be useful to see whether it might be better to move the predicate definition to another file, but it produces many warnings for predicates in multi-use libraries.

2. `unused`: Print a warning message in the case of a predicate that is exported but never used. This can be useful to see if predicate is not used anywhere, and thus could be deleted. Again this produces many warnings for predicates in multi-use libraries.

3. `all_files`: By default, only predicates in files that contain a `:- document_export` or `:- export` declaration are processed. This option causes predicates of *all* files (and modules) to be processed.

4. `all_symbol_uses`: Treat *all* non-predicate uses of symbols (even constants) as predicate uses for the purpose of generating imports.

5. `no_symbol_uses`: Don't treat any non-predicate uses of symbols as predicate uses for the purpose of generating imports.

We further explain the final two options, which allow the user to determine more precisely what uses of a symbol are considered as uses of it as the predicate symbol. All uses of symbols that appear in a "predicate context", i.e., in the body of a rule or in a meta-predicate argument position of a use of a meta-predicate, are considered uses of that predicate symbol. The default is also to allow nonconstant symbols appearing in any other context to also count as uses of that symbol as that predicate symbol. This is useful for programs that define their own meta-predicates.

`checkImpExps(+FileNameList)`
> `checkImpExps/1` is currently equivalent to `checkImpExps([],FileNameList)`.

1.14 Miscellaneous Predicates

`term_hash(+Term,+HashSize,-HashVal)`   module: machine
> Given an arbitrary Prolog term, `Term`, that is to be hashed into a table of `HashSize` buckets, this predicate returns a hash value for `Term` that is between `0` and `HashSize -1`.

`pretty_print(+ClausePairs)`   module: pretty_print
> The input to `pretty_print/1`, `ClausePairs`, can be either a list of clause pairs or a single clause pair. A clause pair is either a Prolog clause (or declaration) or a pair:

$$(\texttt{Clause,Dict})$$

Where `Dict` is a list of the form `A = V` where `V` is a variable in `Clause` and `A` is the string to be used to denote the variable [8].

By default, `pretty_print/1` outputs atomic terms using `writeq/1`, but specialized output can be configured via asserting in `usermod` a term of the form

$$\texttt{user\_replacement\_hook(Term,Call)}$$

which will use Call to output an atomic literal A whenever A unifies with `Term`. For example, pretty printing weight constraints in XSB's `XASP` package is done via the hook

$$\texttt{user\_replacement\_hook(weight\_constr(Term),output\_weight\_constr(Term))}$$

which outputs a weight constraint in a (non-Prolog) syntax that is used by several ASP systems.

1.15 Other Libraries

Not all XSB libraries are fully documented. We provide brief summaries of some of these other libraries.

1.15.1 Justification

By Hai-Feng Guo

Most Prolog debuggers, including XSB's, are based on a mechanism that allows a user to trace the evaluation of a goal by interrupting the evaluation at call, success, retry, or failure of various subgoals. While this has proved an excellent mechanism for evaluating SLD(NF) executions, it is difficult at best to use such a mechanism during a tabled evaluation. This is because, unlike with SLD(NF), SLG requires answers to be returned to tabled subgoals at various times (depending on whether batched or local evaluation is used), negative subgoals to be sometimes be delayed and/or simplified, etc.

One approach to understanding tabled evaluation better is to abstract away the procedural aspects of debugging and to use the tables produced by an evaluation to construct a *justification* after the evaluation has finished. The justification library does just this using algorithms described in [12].

1.15.2 AVL Trees

By Mats Carlsson

AVL trees provide a mechanism to maintain key value pairs so that loop up, insertion, and deletion all have complexity $\mathcal{O}(\log n)$. This library contains predicates to transform a sorted list to an AVL tree and back, along with predicates to manipulate the AVL trees.

---

[8]Thus the list of variable names returned by `read_term/{2,3}` can be used directly in `Dict`.

1.15.3 Ordered Sets: `ordsets.P`

By Richard O'Keefe

Summary from code documentation) `ordset.P` provides an XSB port of the widely used ordset library, whose summary we paraphrase here. In the ordset library, sets are represented by ordered lists with no duplicates. Thus $\{c,r,a,f,t\}$ is represented as `[a,c,f,r,t]`. The ordering is defined by the `@<` family of term comparison predicates, which is the ordering used by `sort/2` and `setof/3`. The benefit of the ordered representation is that the elementary set operations can be done in time proportional to the sum of the argument sizes rather than their product. Some of the unordered set routines, such as `member/2`, `length/2`, or `select/3` can be used unchanged.

1.15.4 Unweighted Graphs: `ugraphs.P`

By Mats Carlsson

XSB also includes a library for unweighted graphs. This library allows for the representation and manipulation of directed and non-directed unlabelled graphs, including predicates to find the transitive closure of a graph, maximal paths, minimal paths, and other features. This library represents graphs as an ordered set of their edges and does not use tabling. As a result, it may be slower for large graphs than similar predicates based on a datalog representatoin of edges.

1.15.5 Heaps: `heaps.P`

By Richard O'Keefe

(Summary from code documentation). A heap is a labelled binary tree where the key of each node is less than or equal to the keys of its sons. The point of a heap is that we can keep on adding new elements to the heap and we can keep on taking out the minimum element. If there are N elements total, the total time is $\mathcal{O}(Nlg(N))$. If you know all the elements in advance, you are better off doing a merge-sort, but this file is for when you want to do say a best-first search, and have no idea when you start how many elements there will be, let alone what they are.

A heap is represented as a triple `t(N, Free, Tree)` where `N` is the number of elements in the tree, `Free` is a list of integers which specifies unused positions in the tree, and `Tree` is a tree made of t terms for empty subtrees and `t(Key,Datum,Lson,Rson)` terms for the rest The nodes of the tree are notionally numbered like this:

```
                      1
          2                           3
      4           6           5           7
    8    12    10    14     9    13    11    15
  .. .. .. .. .. .. .. .. .. .. .. .. .. .. .. ..
```

The idea is that if the maximum number of elements that have been in the heap so far is M, and the tree currently has K elements, the tree is some subtreee of the tree of this form having

exactly M elements, and the Free list is a list of $K - M$ integers saying which of the positions in the M-element tree are currently unoccupied. This free list is needed to ensure that the cost of passing N elements through the heap is $\mathcal{O}(Nlg(M))$ instead of $\mathcal{O}(NlgN)$. For M say 100 and N say 10^4 this means a factor of two.

Chapter 2

Foreign Language Interface

When XSB is used to build real-world systems, a foreign-language interface may be necessary to:

- combine XSB with existing programs and libraries, thereby forming composite systems;

- interface XSB with the operating system, graphical user interfaces or other system level programs;

- speed up certain critical operations.

XSB has both a high-level and the low-level interface to C. The low-level interface is much more flexible, but it requires greater attention to details of how the data is passed between XSB and C. To connect XSB to a C program using the high-level interface requires very little work, but the program must be used "as is" and it must take the input and produce the output supported by this high-level interface. Before describing the interfaces themselves, we first describe aspects common to both the lower- and higher-level foreign language interfaces.

The foreign language interface can also support `C++` programs. Since XSB is written in `C`, the interface functions in the foreign `C++` module must have the declaration `extern ''C''`, and a separate compiler option (e.g. specifying `g++` rather than `gcc`) may need to be given to ensure proper linkage, inclusion of `C++` libraries, etc. In addition, on certain platforms compilation may need to be done externally to XSB – see the `xasp` 1package for a example of using the foreign language interface with `C++` files. For the rest of this chapter, we restrict our attention to foreign predicates written in `C`.

2.1 Foreign Language Modules

Foreign predicates must always appear in modules, and these modules can contain only foreign predicates. A foreign module differs from a Prolog module in that the foreign module's source file must appear in a `*.c` file rather than a `*.P` file (or `.pl` file). This `*.c` file cannot contain a `main()` function. Furthermore, a `*.P` file with the same name *must not* be present or else the `*.c` file is ignored and the module is compiled as a regular Prolog module. The interface part of a foreign

module, which has the same syntax as that of a normal module, is written in Prolog and must appear in a *.H file. If the lower-level interface is used, this *.H file contains explicit `export/1` declarations for the the foreign predicates that are to be used by other modules; if the higher-level interface is used, the declarations have the form `foreign_pred/1`.

The Prolog predicates attached to foreign functions are deterministic, in the sense that they succeed at most once for a given call and are not re-entered on backtracking. Note that this requirement imposes no serious limitation, since it is always possible to divide a foreign predicate into the part to be done on the first call and the part to be redone on backtracking. Backtracking can then take place at the Prolog level where it is more naturally expressed.

A foreign module can be compiled or consulted just like a normal Prolog module. Currently, predicates `consult/[1,2]` recompile both the *.c and the *.H files of a foreign module when at least one of them has been changed from the time the corresponding object files have been created (see the section *Compiling and Consulting* in Volume 1) [1]. The C compiler used to compile the *.c files can be set as a defaults to that used for the configuration of XSB (refer to the section *Getting Started with XSB* in Volume 1). This default behavior includes the C compilation options used to compile XSB when it was configured, along with a default set of include files so that header files in XSB directories can be obtained. Alternately, the user can add options to be passed to the C compiler. To give an example, the following command will compile file `file.c` using the default C Compiler with optimization and by including `/usr/local/X11/R6/include` to the directories that will be searched for header files.

```
:- consult(file, [cc_opts('-O2 -I/usr/local/X11/R6/include')]).
```

Note in particular, that if XSB were compiled with the `-g` debugging option, then the C file will be also [2]. Any Prolog compiler options are ignored when compiling a foreign module.

Prolog-specific directives such as `index`, `hilog`, `table`, `auto_table` or even `import` make no sense in the case of a foreign module and thus are ignored by the compiler. However, another directive, namely `ldoption`, is recognized in a foreign module and is used to instruct the dynamic loading and linking of the module. The syntax of the `ldoption` directive is simply:

```
:- ldoption(Option).
```

where `Option` should either be an atom or a list of atoms. Multiple `ldoption` directives may appear in the same .H file of a foreign module [3]. In Unix-derived systems, the foreign language interface of XSB uses `ld` command that combines object programs to create an executable file or

---

[1]In addition, if a C module compiled by the single-threaded XSB engine is loaded by the multi-threaded engine, it will be recompiled, and vice-versa.

[2] In a 64-bit platform, users may override the default compilation of XSB by the configuration options `-with-bits32` or `-with-bits64`. If either of these options is used, the default compilation options will pass along the appropriate memory options. If XSB is compiled with a memory option that is not the default of the platform, and if an externally compiled C file is to be loaded into XSB, it must be ensured that the C file has been compiled with the appropriate memory options: `-m32` or `-m64` if `gcc` is used.

[3]Mac OSX users using 10.3 or above should have the environment variable `MACOSX_DEPLOYMENT_TARGET` set to `10.3` so that the compiler generates code that can be dynamically linked by XSB. This should be done automatically by XSB on initialization, but it is useful to check if encountering problems.

another object program suitable for further `ld` processing. Version 3.2 of XSB assumes that the `ld` command resides in the file `/usr/bin/ld`.

2.2   Lower-Level Foreign Language Interface

Creating a foreign predicate using the lower-level foreign language interface is almost entirely a matter of writing C code. Consider the foreign module `$XSBDIR/examples/XSB_calling_c/simple_foreign.[cH]`. The .H file has the form:

```
:- export minus_one/2, my_sqrt/2, change_char/4.

:- ldoption('-lm').     % link together with the math library
```

When the lower level foreign language interface is used, C functions that implement foreign predicates must return values of type `int`. The return value is not used by a Prolog argument; rather if a non-zero is returned, the foreign predicate succeeds; a zero return value means failure.

At the C level, the function that implements the Prolog predicate must have the same name as the Prolog predicate (that is declared in the `*.H` file), and must have a special *context parameter* macro. The context parameter macro allows C functions to be used with both the single-threaded and multi-threaded engines, and are described in detail in Section 2.2.1. The Prolog level arguments are converted to C data structures through several predefined functions rather than through direct parameter passing [4]. The C file `simple_foreign.c` corresponding to the above .H file is as follows.

```
/*----------------------------------------------------------------------*/

#include <math.h>
#include <stdio.h>
#include <string.h>
#include <alloca.h>

/*----- Make sure your C compiler finds the following header file.    -----
  ----- One way to do this is to include the directory XSB/emu on the -----
  ----- compiler's command line with the -I (/I in Windows) option    -----*/

#include "cinterf.h"

/*---------------------------------------*/

int minus_one(CTXTdecl)
{
    int  i = ptoc_int(CTXTc 1);
```

---

[4]The inclusion of context parameters changes the lower-level interface for Version 3.0. C files written for previous versions of XSB continue to work properly for the single-threaded engine in, but will not work properly for the multi-threaded engine.

```
   ctop_int(CTXTc 2, i-1);
   return TRUE;
}

/*----------------------------------------*/

int my_sqrt(CTXTdecl)
{
   int i = ptoc_int(CTXTc 1);

   ctop_float(CTXTc 2, (float) pow((double)i, 0.5));
   return TRUE;
}

/*----------------------------------------*/

int change_char(CTXTdecl)
{
   char *str_in;
   int  pos;
   int c;
   char *str_out;

   str_in = (char *) ptoc_string(CTXTc 1);
   str_out = (char *) alloca(strlen(str_in)+1);
   strcpy(str_out, str_in);
   pos = ptoc_int(CTXTc (2);
   c = ptoc_int(CTXTc (3);
   if (c < 0 || c > 255) /* not a character */
     return FALSE; /* this predicate will fail on the Prolog side */

   str_out[pos-1] = c;

   extern_ctop_string(CTXTc 4, str_out);
   return TRUE;
}

/*------------------------------------------------------------------*/
```

Before describing the C program used, here is a sample session illustrating the behavior of the predicates in `simple_foreign`.

```
XSB Version 2.0 (Gouden Carolus) of June 26, 1999
[i686-pc-linux-gnu; mode: optimal; engine: slg-wam; scheduling: batched]
| ?- [simple_foreign].
[Compiling C file ./simple_foreign.c using gcc]
[Compiling Foreign Module ./simple_foreign]
[simple_foreign compiled, cpu time used: 0.0099993 seconds]
[simple_foreign loaded]
```

```
yes
| ?- change_char('Kostis', 2, w, TempStr),
      change_char(TempStr, 5, h, GrkName).

TempStr = Kwstis
GrkName = Kwsths;

no
| ?- minus_one(43, X).

X = 42;

no
| ?- minus_one(43, 42).                   % No output unification is allowed
Wrong arg in ctop_int 2a2 (Reg = 2)

yes
| ?- my_sqrt(4,X).

X = 2

yes
| ?- my_sqrt(23,X).

X = 4.7958;

no
```

Consider the function `minus_one()` above. As discussed, it takes a context parameter (explained below), and returns an integer, and as can be seen the return values can be specified by the macros `TRUE` and `FALSE`. From the Prolog perspective the first argument to `minus_one/2` is an (integer) input argument, while the second is an (integer) output argument. Input arguments for basic C types are translated from their Prolog representation to a C representation by functions of the form `ptoc\_<type>()` – here `ctop_int()`. The single parameter of such a function is the number of the Prolog argument that is to be transformed and the function returns the C representation. Output arguments are converted from C to Prolog by corresponding functions of the form `ctop_<type>()` – here `ctop_int()`. For converting C back to Prolog, the first parameter of `ctop_int()` is the number of the Prolog argument to be transformed and the second is the C value to be transformed. In the session output above, if an improper argument is given to `minus_one/2` it will emit a warning, and succeed. Also note that the call `my_sqrt(23,X)` succeeds once, but fails on backtracking since it is deterministic, as are all other foreign language functions.

The above example illustrates the exchange of *basic* types through the lower-level interface – e.g. atoms, integers, and floating-point numbers. The lower-level interface also allows a user to pass lists and terms between XSB and C as will be discussed in Section 2.2.3.

2.2.1   Context Parameters

When using the lower-level interface, *context parameters* must be added to many C functions in order for the functions to be used with XSB's multi-threaded engine. In the multi-threaded engine, variables for Prolog's virtual machine, as well as for thread-private data structures are stored in a *context structure*. This context structure must be passed to any functions that need to access elements of a thread's virtual machine – including many of the functions that are used to exchange data between Prolog and C. We note in passing that when using the multi-threaded engine, a user must ensure that foreign-language functions are thread-safe, by using standard multi-threaded programming techniques, including XSB's mutex predicates (see the Section *Predicates for Thread Synchronization* in Volume 1 of this manual). On the other hand, in the single-threaded engine virtual machine elements are kept in static variables, so that context parameters are not required.

The lower-level C interface makes use of a set of macros to address the requirements of the different engines. The data exchange functions discussed in this chapter, `ptoc_xxx`, `ctop_xxx`, `c2p_xxx`, `p2c_xxx`, and `p2p_xxx` usually, but not always, require information about a threads virtual machine state. If a C function directly or indirectly calls a data interchange function that requires a context parameter, the function must have a context parameter in its declaration, calls, and prototypes in order to be used by the multi-threaded engine. These context parameters have the following forms:

- In function *declarations*, use the macro `CTXTdecl` in the code for a function that would otherwise be `void`, and `CTXTdeclc` as the first argument in the code for a function with parameters (`CTXTdeclc` and `CTXTdecl` are similar, except that macro expansion of `CTXTdeclc` for the multi-threaded engine includes a comma). The example for `minus_one(CTXTdecl)` shows use of this macro.

- In function *calls* use the macro `CTXT` in the code for a function that would otherwise be `void`, and `CTXTc` as the first argument in the code for a function with parameters. As an example, a call to `minus_one` would have the form `minus_one(CTXT)`.

- In function *prototypes* use the macro `CTXTdecltype` in the code for a function that would otherwise be `void`, and `CTXTdecltypec` as the first argument in the code for a function with parameters. As an example, a prototype for `minus_one` would have the form `minus_one(CTXTdecltype)`.

Fortunately, when compiling with the multi-threaded engine, it is easy to determine at compile time whether context parameters are correct. If compilation of a function `foo` gives an error along the lines of:

```
foofile.c: In function 'foo':
foofile.c:109: error: 'th' undeclared (first use in this function)
```

Then the declaration of `foo` omitted a context parameter. If compilation gives an error along the lines of

```
foofile.c: In function 'foo_caller':
```

```
    :
  foofile.c:149: error: too few arguments to function 'foo'
```

Then the call to `foo` may have omitted a context parameter.

Note that context parameters are *only* necessary if the lower-level interface is used. The higher-level interface automatically generates any context parameters it needs.

2.2.2   Exchanging Basic Data Types

The basic interface assumes that correct modes (*i.e.*, input or output parameters) and types are being passed between the C and Prolog levels. As a result, output unification should be explicitly performed in the Prolog level. The prototypes for the conversion functions between Prolog and C should be declared before the corresponding functions are used. This is done by including the `"cinterf.h"` header file. Under Unix, the XSB foreign C interface automatically finds this file in the `XSB/emu` directory. Under Windows (including Cygwin), the user must compile and create the DLL out of the C file manually, so the compiler option '`/I...\XSB\emu`' is necessary.

The following C functions are used to convert basic types between Prolog and C.

`int ptoc_int(CTXTdeclc int N)`
> `CTXTdeclc` is a context parameter; `N` is assumed to hold a Prolog integer corresponding to the Nth argument of a Prolog predicate. This function returns the value of that argument in as a C `int`.

`double ptoc_float(CTXTdeclc int N)`
> `CTXTdeclc` is a context parameter; `N` is assumed to hold a Prolog integer corresponding to the Nth argument of a Prolog predicate. This function returns the value of that argument as a C `double`. By default, XSB provides double precision, but if XSB was configured with `--enable-fast-floats` less than single precision can be provided [5].

`char *ptoc_string(CTXTdeclc int N)`
> `CTXTdeclc` is a context parameter; `N` is assumed to hold a Prolog integer corresponding to the Nth argument of a Prolog predicate. This function returns the value the C string (of type `char *`) that corresponds to this interned Prolog atom. *WARNING: the string should be copied before being manipulated in any way: otherwise unexpected results may arise whenever the interned Prolog atom is unified.*

`void ctop_int(CTXTdeclc int N, int V)`
> `CTXTdeclc` is a context parameter; argument `N` is assumed to hold a Prolog free variable, and this function binds that variable to an integer of value `V`.

`void ctop_float(CTXTdeclc int N, float V)`
> `CTXTdeclc` is a context parameter; argument `N` is assumed to hold a Prolog free variable, and this function binds that variable to a floating point number of value `V`.

---

[5] The fast float configuration option does represents floating point values as directly tagged single precision values rather than as indirectly tagged double precision values. Speed increases in arithmetic can be gained from this optimization, in exchange for significant precision loss on floating point numbers.

```
void extern_ctop_string(CTXTdeclc int N, char * V)
```
> CTXTdeclc is a context parameter; argument N is assumed to hold a Prolog free variable. If needed, this function interns the string to which V points as a Prolog atom and then binds the variable in argument N to that atom.

2.2.3   Exchanging Complex Data Types

If the lower-level interface is used, exchanging basic data types is sufficient for most applications. Exchanging complex data types is also possible, although doing so is slightly more involved than exchanging basic types. To exchange complex data types, the lower-level interface uses only one C data type: `prolog_term`, which can point to any XSB term. On the C side, the type of the term can be checked and then processed accordingly. For instance, if the term turns out to be a structure, then it can be decomposed and the functor can be extracted along with the arguments. If the term happens to be a list, then it can be processed in a loop and each list member can be further decomposed into its atomic components. The advanced interface also provides functions to check the types of these atomic components and for converting them into C types.

We begin by presenting the functions used to exchange complex data types, before presenting a detailed example below. As when exchanging basic C types, the file `emu/cinterf.h` must be included in the C program in order to make the prototypes of the relevant functions known to the C compiler.

The first set of functions is typically used to check the type of Prolog terms passed into the C program.

```
xsbBool is_attv((prolog_term) T)
```
> `is_attv(T)` returns TRUE if T represents an XSB attributed variable, and FALSE otherwise.

```
xsbBool is_float((prolog_term) T)
```
> `is_float(T)` returns TRUE if T represents an XSB float value, and FALSE otherwise.

```
xsbBool is_functor((prolog_term) T)
```
> `is_functor(T)` returns TRUE if T represents an XSB structure value (not a list), and FALSE otherwise.

```
xsbBool is_int((prolog_term) T)
```
> `is_int(T)` returns TRUE if T represents an XSB integer value, and FALSE otherwise.

```
xsbBool is_list((prolog_term) T)
```
> `is_list(T)` returns TRUE if T represents an XSB list value (not nil), and FALSE otherwise.

```
xsbBool is_nil((prolog_term) T)
```
> `is_nil(T)` returns TRUE if T represents an XSB [] (nil) value, and FALSE otherwise.

```
xsbBool is_string((prolog_term) T)
```
> `is_string(T)` returns TRUE if T represents an XSB atom value, and FALSE otherwise.

```
xsbBool is_var((prolog_term) T)
```
> `is_var(T)` returns TRUE if T represents an XSB variable, and FALSE otherwise.

After checking the types of the arguments passed in from the Prolog side, the next task usually is to convert Prolog data into the types understood by C. This is done with the following functions. The first three convert between the basic types. The last two extract the functor name and the arity. Extraction of the components of a list and the arguments of a structured term is explained later.

int p2c_int((prolog_term) V)
> The `prolog_term` parameter must represent a Prolog integer, and `p2c_int` returns the C representation of that integer.

double p2c_float((prolog_term) V)
> The `prolog_term` parameter must represent a Prolog floating point number, and `p2c_float` returns the C representation of that floating point number.

char *p2c_string((prolog_term) V)
> The `prolog_term` parameter must represent a (Prolog) atom, and `p2c_string` returns that atom as a C string. The pointer returned points to the actual atom name in XSB 's atom table, and thus it must NOT be modified by the calling program.

char *p2c_functor((prolog_term) V)
> The `prolog_term` parameter must represent a structured term (not a list). `p2c_functor` returns the name of the main functor symbol of that term as a string. The pointer returned points to the actual functor name in XSB 's space, and thus it must NOT be modified by the calling program.

int p2c_arity((prolog_term) V)
> The `prolog_term` parameter must represent a structured term (not a list). `p2c_arity` returns the arity of the main functor symbol of that term as a C `int`.

The next batch of functions support conversion of data in the opposite direction: from basic C types to the type `prolog_term`. These `c2p_*` functions all return a boolean value `TRUE` if successful and `FALSE` if unsuccessful. The XSB term argument must always contain an XSB variable, which will be bound to the indicated value as a side effect of the function call.

xsbBool c2p_int(CTXTdeclc (int) N, (prolog_term) V)
> CTXTdeclc is a context parameter; `c2p_int` binds the `prolog_term` V (which must be a variable) to the integer value N, creating a Prolog integer.

xsbBool c2p_float(CTXTdeclc (double) F, (prolog_term) V)
> CTXTdeclc is a context parameter; `c2p_float` binds the `prolog_term` V (which must be a variable) to the (double) float value F, creating a double Prolog float.

xsbBool c2p_string(CTXTdeclc (char *) S, (prolog_term) V)
> CTXTdeclc is a context parameter; `c2p_string` binds the `prolog_term` V (which must be a variable) to the Prolog atom corresponding to the `char *S`. During this process the Prolog atom is interned into XSB's atom table.

   The following functions create Prolog data structures within a C program. This is usually done in order to pass these structures back to the Prolog side.

xsbBool c2p_functor(CTXTdeclc (char *) S, (int) N, (prolog_term) V)
   CTXTdeclc is a context parameter; c2p_functor binds the prolog_term V (which must be a variable) to an open term whose main functor symbol is given by S (of type char *) and whose arity is N. An open term is one with all arguments as new distinct variables.

xsbBool c2p_list(CTXTdeclc (prolog_term) V)
   CTXTdeclc is a context parameter; c2p_list binds the prolog_term V (which must be a variable) to an open list term, i.e., a list term with both car and cdr as new distinct variables. Note: to create an empty list use the function c2p_nil described below.

xsbBool c2p_nil(CTXTdeclc (prolog_term) V)
   CTXTdeclc is a context parameter; c2p_nil binds the prolog_term V (which must be a variable) to the atom [] (nil).

prolog_term p2p_new()
   Create a new Prolog variable. This is sometimes needed when you want to create a Prolog term on the C side and pass it to the Prolog side.

   To use the above functions, one must be able to get access to the components of the structured Prolog terms. This is done with the help of the following functions:

prolog_term p2p_arg((prolog_term) T, (int) A)
   Parameter T must be a prolog_term that is a structured term (but not a list). A is a positive integer (no larger than the arity of the term) that specifies an argument position of the term T. p2p_arg returns the A^{th} subfield of the term T.

prolog_term p2p_car((prolog_term) T)
   Parameter T must be a prolog_term that is a list (not nil). p2p_car returns the car (i.e., head of the list) of the term T.

prolog_term p2p_cdr((prolog_term) T)
   Parameter T must be a prolog_term that is a list (not nil). p2p_cdr returns the cdr (i.e., tail of the list) of the term T.

   It is important to realize that these functions return the actual Prolog term that is, say, the head of a list or the actual argument of a structured term. Thus, assigning a value to such a Prolog term also modifies the head of the corresponding list or the relevant argument of the structured term. It is precisely this feature that allows passing structured terms and lists from the C side to the Prolog side. For instance,

```
   prolog_term plist,      /* a Prolog list            */
              structure;   /* something like f(a,b,c) */
   prolog_term tail, arg;
```

```
..........
tail = p2p_cdr(plist);          /* get the list tail  */
arg  = p2p_arg(structure, 2);  /* get the second arg */

/* Assume that the list tail was supposed to be a prolog variable */
if (is_var(tail))
   c2p_nil(CTXTc tail);  /* terminate the list */
else {
   fprintf(stderr, "Something wrong with the list tail!");
   exit(1);
}
/* Assume that the argument was supposed to be a prolog variable */
c2p_string(CTXTc "abcdef", arg);
```

In the above program fragment, we assume that both the tail of the list and the second argument of the term were supposed to be bound to Prolog variables. In case of the tail, we check if this is, indeed, the case. In case of the argument, no checks are done; XSB will issue an error (which might be hard to track down) if the second argument is not currently bound to a variable.

The last batch of functions is useful for passing data in and out of the Prolog side of XSB. The first function is the only way to get a `prolog_term` out of the Prolog side; the second function is sometimes needed in order to pass complex structures from C into Prolog.

`prolog_term reg_term(CTXTdeclc (int) R)`
> `CTXTdeclc` is a context parameter. Parameter R is an argument number of the Prolog predicate implemented by this C function (range 1 to 255). The function `reg_term` returns the `prolog_term` in that predicate argument.

`xsbBool p2p_unify(CTXTdeclc prolog_term T1, prolog_term T2)`
> Unify the two Prolog terms. This is useful when an argument of the Prolog predicate (implemented in C) is a structured term or a list, which acts both as input and output parameter. `CTXTdeclc` is a context parameter.

For instance, consider the Prolog call `test(X, f(Z))`, which is implemented by a C function with the following fragment:

```
prolog_term newterm, newvar, z_var, arg2;
.....
/* process argument 1 */
c2p_functor(CTXTc "func",1,reg_term(CTXTc 1));
c2p_string(CTXTc "str",p2p_arg(reg_term(CTXTc 1),1));
/* process argument 2 */
arg2 = reg_term(CTXTc 2);
z_var = p2p_arg(arg2, 1);  /* get the var Z */
/* bind newterm to abc(V), where V is a new var */
c2p_functor(CTXTc "abc", 1, newterm);
```

```
    newvar = p2p_arg(newterm, 1);
    newvar = p2p_new();
    ....
    /* return TRUE (success), if unify; FALSE (failure) otherwise */
    return p2p_unify(CTXTc z_var, newterm);
```

On exit, the variable X will be bound to the term func(str). Processing argument 2 is more interesting. Here, argument 2 is used both for input and output. If test is called as above, then on exit Z will be bound to abc(_h123), where _h123 is some new Prolog variable. But if the call is test(X,f(1)) or test(X,f(Z,V)) then this call will *fail* (fail as in Prolog, *i.e.*, it is not an error), because the term passed back, abc(_h123), does not unify with f(1) or f(Z,V). This effect is achieved by the use of p2p_unify above.

We conclude this section with two real examples of functions that pass complex data in and out of the Prolog side of XSB. These functions are part of the POSIX regular expression matching package of XSB. The first function uses argument 2 to accept a list of complex Prolog terms from the Prolog side and does the processing on the C side. The second function does the opposite: it constructs a list of complex Prolog terms on the C side and passes it over to the Prolog side in argument 5.

(We should note that this second function could cause a heap overflow in XSB were it to build a large list of values. Instead of building a large list of values on the XSB heap, one would better design the functions to return smaller values, in which case XSB will be able to automatically expand the heap as necessary.)

```
/* XSB string substitution entry point: replace substrings specified in Arg2
   with strings in Arg3.
   In:
       Arg1: string
       Arg2: substring specification, a list [s(B1,E1),s(B2,E2),...]
       Arg3: list of replacement string
   Out:
       Arg4: new (output) string
   Always succeeds, unless error.
*/
int do_regsubstitute__(CTXTdecl)
{
  /* Prolog args are first assigned to these, so we could examine the types
     of these objects to determine if we got strings or atoms. */
  prolog_term input_term, output_term;
  prolog_term subst_reg_term, subst_spec_list_term, subst_spec_list_term1;
  prolog_term subst_str_term=(prolog_term)0,
    subst_str_list_term, subst_str_list_term1;
  char *input_string=NULL;    /* string where matches are to be found */
  char *subst_string=NULL;
  prolog_term beg_term, end_term;
  int beg_offset=0, end_offset=0, input_len;
  int last_pos = 0; /* last scanned pos in input string */
  /* the output buffer is made large enough to include the input string and the
```

```
    substitution string. */
char subst_buf[MAXBUFSIZE];
char *output_ptr;
int conversion_required=FALSE; /* from C string to Prolog char list */

input_term = reg_term(CTXTc 1);  /* Arg1: string to find matches in */
if (is_string(input_term)) /* check it */
  input_string = string_val(input_term);
else if (is_list(input_term)) {
  input_string =
    p_charlist_to_c_string(input_term, input_buffer, sizeof(input_buffer),
                           "RE_SUBSTITUTE", "input string");
  conversion_required = TRUE;
} else
  xsb_abort("RE_SUBSTITUTE: Arg 1 (the input string) must be an atom or a character list");

input_len = strlen(input_string);

/* arg 2: substring specification */
subst_spec_list_term = reg_term(CTXTc 2);
if (!is_list(subst_spec_list_term) && !is_nil(subst_spec_list_term))
  xsb_abort("RE_SUBSTITUTE: Arg 2 must be a list [s(B1,E1),s(B2,E2),...]");

/* handle substitution string */
subst_str_list_term = reg_term(CTXTc 3);
if (! is_list(subst_str_list_term))
  xsb_abort("RE_SUBSTITUTE: Arg 3 must be a list of strings");

output_term = reg_term(CTXTc 4);
if (! is_var(output_term))
  xsb_abort("RE_SUBSTITUTE: Arg 4 (the output) must be an unbound variable");

subst_spec_list_term1 = subst_spec_list_term;
subst_str_list_term1 = subst_str_list_term;

if (is_nil(subst_spec_list_term1)) {
  strncpy(output_buffer, input_string, sizeof(output_buffer));
  goto EXIT;
}
if (is_nil(subst_str_list_term1))
  xsb_abort("RE_SUBSTITUTE: Arg 3 must not be an empty list");

/* initialize output buf */
output_ptr = output_buffer;

do {
  subst_reg_term = p2p_car(subst_spec_list_term1);
  subst_spec_list_term1 = p2p_cdr(subst_spec_list_term1);

  if (!is_nil(subst_str_list_term1)) {
    subst_str_term = p2p_car(subst_str_list_term1);
```

```
      subst_str_list_term1 = p2p_cdr(subst_str_list_term1);

    if (is_string(subst_str_term)) {
      subst_string = string_val(subst_str_term);
    } else if (is_list(subst_str_term)) {
      subst_string =
        p_charlist_to_c_string(subst_str_term, subst_buf, sizeof(subst_buf),
                               "RE_SUBSTITUTE", "substitution string");
    } else
      xsb_abort("RE_SUBSTITUTE: Arg 3 must be a list of strings");
  }

  beg_term = p2p_arg(subst_reg_term,1);
  end_term = p2p_arg(subst_reg_term,2);

  if (!is_int(beg_term) || !is_int(end_term))
    xsb_abort("RE_SUBSTITUTE: Non-integer in Arg 2");
  else{
    beg_offset = int_val(beg_term);
    end_offset = int_val(end_term);
  }
  /* -1 means end of string */
  if (end_offset < 0)
    end_offset = input_len;
  if ((end_offset < beg_offset) || (beg_offset < last_pos))
    xsb_abort("RE_SUBSTITUTE: Substitution regions in Arg 2 not sorted");

  /* do the actual replacement */
  strncpy(output_ptr, input_string + last_pos, beg_offset - last_pos);
  output_ptr = output_ptr + beg_offset - last_pos;
  if (sizeof(output_buffer)
      > (output_ptr - output_buffer + strlen(subst_string)))
    strcpy(output_ptr, subst_string);
  else
    xsb_abort("RE_SUBSTITUTE: Substitution result size %d > maximum %d",
              beg_offset + strlen(subst_string),
              sizeof(output_buffer));

  last_pos = end_offset;
  output_ptr = output_ptr + strlen(subst_string);

} while (!is_nil(subst_spec_list_term1));

if (sizeof(output_buffer) > (output_ptr-output_buffer+input_len-end_offset))
  strcat(output_ptr, input_string+end_offset);

EXIT:
/* get result out */
if (conversion_required)
  c_string_to_p_charlist(output_buffer,output_term,"RE_SUBSTITUTE","Arg 4");
else
```

```
    /* DO NOT intern. When atom table garbage collection is in place, then
       replace the instruction with this:
                    c2p_string(CTXTc output_buffer, output_term);
       The reason for not interning is that in Web page
       manipulation it is often necessary to process the same string many
       times. This can cause atom table overflow. Not interning allows us to
       circumvent the problem.  */
    extern_ctop_string(CTXTc 4, output_buffer);

  return(TRUE);
}



/* XSB regular expression matcher entry point
   In:
       Arg1: regexp
       Arg2: string
       Arg3: offset
       Arg4: ignorecase
   Out:
       Arg5: list of the form [match(bo0,eo0), match(bo1,eo1),...]
             where bo*,eo* specify the beginning and ending offsets of the
             matched substrings.
             All matched substrings are returned. Parenthesized expressions are
             ignored.
*/
int do_bulkmatch__(CTXTdecl)
{
  prolog_term listHead, listTail;
  /* Prolog args are first assigned to these, so we could examine the types
     of these objects to determine if we got strings or atoms. */
  prolog_term regexp_term, input_term, offset_term;
  prolog_term output_term = p2p_new();
  char *regexp_ptr=NULL;       /* regular expression ptr                */
  char *input_string=NULL;    /* string where matches are to be found */
  int ignorecase=FALSE;
  int return_code, paren_number, offset;
  regmatch_t *match_array;
  int last_pos=0, input_len;
  char regexp_buffer[MAXBUFSIZE];

  if (first_call)
    initialize_regexp_tbl();

  regexp_term = reg_term(CTXTc 1);  /* Arg1: regexp */
  if (is_string(regexp_term)) /* check it */
    regexp_ptr = string_val(regexp_term);
  else if (is_list(regexp_term))
    regexp_ptr =
      p_charlist_to_c_string(regexp_term, regexp_buffer, sizeof(regexp_buffer),
                             "RE_MATCH", "regular expression");
```

```
    else
      xsb_abort("RE_MATCH: Arg 1 (the regular expression) must be an atom or a character list");

    input_term = reg_term(CTXTc 2);  /* Arg2: string to find matches in */
    if (is_string(input_term)) /* check it */
      input_string = string_val(input_term);
    else if (is_list(input_term)) {
      input_string =
        p_charlist_to_c_string(input_term, input_buffer, sizeof(input_buffer),
                               "RE_MATCH", "input string");
    } else
      xsb_abort("RE_MATCH: Arg 2 (the input string) must be an atom or a character list");

    input_len = strlen(input_string);

    offset_term = reg_term(CTXTc 3); /* arg3: offset within the string */
    if (! is_int(offset_term))
      xsb_abort("RE_MATCH: Arg 3 (the offset) must be an integer");
    offset = int_val(offset_term);
    if (offset < 0 || offset > input_len)
      xsb_abort("RE_MATCH: Arg 3 (=%d) must be between 0 and %d", input_len);

    /* If arg 4 is bound to anything, then consider this as ignore case flag */
    if (! is_var(reg_term(CTXTc 4)))
      ignorecase = TRUE;

    last_pos = offset;
    /* returned result */
    listTail = output_term;
    while (last_pos < input_len) {
      c2p_list(CTXTc listTail); /* make it into a list */
      listHead = p2p_car(listTail); /* get head of the list */

      return_code = xsb_re_match(regexp_ptr, input_string+last_pos, ignorecase,
                                 &match_array, &paren_number);
      /* exit on no match */
      if (! return_code) break;

      /* bind i-th match to listHead as match(beg,end) */
      c2p_functor(CTXTc "match", 2, listHead);
      c2p_int(CTXTc match_array[0].rm_so+last_pos, p2p_arg(listHead,1));
      c2p_int(CTXTc match_array[0].rm_eo+last_pos, p2p_arg(listHead,2));

      listTail = p2p_cdr(listTail);
      last_pos = match_array[0].rm_eo+last_pos;
    }
  c2p_nil(CTXTc listTail); /* bind tail to nil */
  return p2p_unify(CTXTc output_term, reg_term(CTXTc 5));
}
```

2.3   Foreign Modules That Call XSB Predicates

A C function that has been called from XSB through the lower-level foreign language interface
may want to call back into XSB to have XSB evaluate a predicate. This can be done by using the
interface described in Chapter 3 (Volume 2) on calling XSB from another language. The interface
described there allows a caller to initialize XSB and pass queries to it. However, since XSB has
already called a foreign module, XSB does not need to be initialized. However it does need to
manage the registers that are in use to support interaction with the foreign module currently
executing. So there are some minor differences with the interface described in Chapter 3.

First, XSB should not be initialized. I.e., a foreign module should **not** call `xsb_init` or
`xsb_init_string`. Second, the foreign module must protect the XSB registers it is currently using
when it calls XSB. To do this, after it has retrieved its arguments into local variables and before
it calls any XSB predicate, it must call `xsb_query_save(NumRegs)`, which saves the current XSB
registers and initializes them to be able to accept a new query. `NumRegs` is the number of registers
used to interact with the currently executing foreign routine (i.e., the arity of the predicate that
called this foreign code.) When the foreign routine has completed its work, it will set the appro-
priate registers with the appropriate return values and return to the caller. Before it does this,
it must call `xsb_query_restore()` to restore the saved registers and prepare XSB for the return.
Note that it must be called before any of the output registers are accessed to set return values. (It
must also be called even if no values are returned.)

In summary the extra functions needed to call XSB from a foreign module are:

`int xsb_query_save(CTXTc (byte) NumRegs)`
    This function is used in a foreign routine that is called from XSB. It is used to save the
    current contents of the XSB registers and to initialize them to be prepared to accept a query.
    It must be called after a foreign routine collects its input arguments from the XSB registers
    and before it invokes any XSB predicate.

`int xsb_query_restore(CTXT)`
    This function is used in a foreign routine that is called from XSB and in turn calls an XSB
    predicate. It is used to restore the previously saved contents of the XSB registers. It must be
    called after all XSB predicates have been called and returned, and before the current foreign
    routine sets its output parameters and returns to XSB.

An example where a foreign module and XSB call each other recursively can be found in the
directory `$XSB_DIR/examples/XSB_calling_c` and files `fibr.[cH]` and `fibp.P`.

2.4   Foreign Modules That Link Dynamically with Other Libraries

Sometimes a foreign module might have to link dynamically with other (non-XSB) libraries. Typ-
ically, this happens when the foreign module implements an interface to a large external library
of utilities. One example of this is the package `libwww` in the XSB distribution, which provides a

high-level interface to the W3C's Libwww library for accessing the Web. The library is compiled into a set of shared objects and the `libwww` module has to link with them as well as with XSB.

The problem here is that the loader must know at run time where to look for the shared objects to link with. On Unix systems, this is specified using the environment variable `LD_LIBRARY_PATH`; on Windows, the variable name is `LIBPATH`. For instance, under Bourne shell or its derivatives, the following will do:

```
LD_LIBRARY_PATH=dir1:dir2:dir3
export LD_LIBRARY_PATH
```

One problem with this approach is that this variable must be set before starting XSB. The other problem is that such a global setting might interact with other foreign modules.

To alleviate the problem, XSB dynamically sets `LD_LIBRARY_PATH` (`LIBPATH` on Windows) before loading foreign modules by adding the directories specified in the `-L` option in `ldoption`. Unfortunately, this works on some systems (Linux), but not on others (Solaris). One route around this difficulty is to build a runtime library search path directly into the object code of the foreign module. This can be specified using a loader flag in `ldoption`. The problem here is that different systems use a different flag! To circumvent this, XSB provides a predicate that tries to guess the right flag for your system:

```
runtime_loader_flag(+Hint,-Flag)
```

Currently it knows about a handful of the most popular systems, but this will be expanded. The argument `Hint` is not currently used. It might be used in the future to provide `runtime_loader_flag` with additional information that can improve the accuracy of finding the right runtime flags for various systems.

The above predicate can be used as follows:

```
...,
runtime_loader_flag(_,Flag),
fmt_write_string(LDoptions, '%sdir1:dir2:dir2 %s', args(Flag,OldLDoption)),
fmt_write(File, ':- ldoption(%s).', LDoptions),
file_nl(File).
```

2.5   Higher-Level Foreign Language Interface

The high-level foreign predicate interface was designed to release the programmer from the burden of having to write low-level code to transfer data from XSB to C and vice-versa. Instead, all the user needs to do is to describe each C function and its corresponding Prolog predicates in the `.H` files. The interface then automatically generates *wrappers* that translate Prolog terms and structures to proper C types, and vice-versa. These wrappers also check for type-correctness of arguments to the C function; in addition, in Unix-derived systems the wrappers are automatically compiled and loaded along with the foreign predicates in the `.c` file [6].

---

[6]for Windows, please see special instructions in Section 2.6.

As with the lower-level foreign interfaces, when predicates are defined in a foreign module `myfile.[cH]`, the predicates must be explicitly imported from the module to be used [7]. For an example of using the higher level interface, see `$XSBDIR/examples/XSB_calling_c/second_foreign.[cH]`.

2.5.1 Declaration of high level foreign predicates

The basic formats of a foreign predicate declaration are:

> `:- foreign_pred` *predname*`([+-]`*parg1*`, [+-]`*parg2*`,...)`
>   `from` *funcname*`(`*carg1:type1*`, `*carg2:type2*`, ...):`*functype*`.`

and

> `:- private_foreign_pred` *predname*`([+-]`*parg1*`, [+-]`*parg2*`,...)`
>   `from` *funcname*`(`*carg1:type1*`, `*carg2:type2*`, ...):`*functype*`.`

where:

`foreign_pred, private_foreign_pred`
    declares a new foreign predicate. For most cases, the declaration `foreign_pred` can be used in both the multi-threaded and the sequential engine. The declaration `private_foreign_pred` needs to be used only in the multi-threaded engine when the external foreign function, `funcname` contains a context parameter as its first argument because `funcname` needs to access thread-private data or other information from the context of the XSB thread (see Section 2.2.1). This case is uncommen, and mostly occurs for users who are creating XSB packages (e.g. the XASP interface to Smodels).

`predname`
    is the name of the foreign Prolog predicate.

`parg1, parg2, ...`
    are the predicate arguments. Each argument is preceded by either '+' or '-', indicating its mode as input or output respectively. The names of the arguments must be the same as those used in the declaration of the corresponding C function. If a C argument is used both for input and output, then the corresponding Prolog argument can appear twice: once with "+" and once with "-". In addition, a special argument `retval` is used to denote the argument that corresponds to the return value of the C function; it must always have the mode '-'.

`funcname`
    is the name of the function in the `.c` file. At compile-time a C function with name `predname` will be generated which will translate arguments from Prolog to C, call `funcname`, and then translate arguments back from C to Prolog.

---

[7]In Version 3.2, a foreign module that uses the higher-level C interface must be explicitly consulted before it can be used.

`carg1, carg2, ...`
      is the list of arguments of the C function. The names used for the arguments must match the
      names used in the Prolog declaration.

`type1, type2, ...`
      are the types associated to the arguments of the C function. This is not the set of C types,
      but rather a set of descriptive types, as defined in Table 2.5.1.

`functype`
      is the return type of the C function.

Using the higher-level interface, the same C code can be used for both the sequential and the
multi-threaded engines, and no context parameters are required in a user's C code unless thread
context information is explicitly needed. However, a foreign module compiled for the single-threaded
engine will need to be recompiled for the multi-threaded engine and vice-versa.

Table 2.5.1 provides the correspondence between the types allowed on the C side of a foreign
module declaration and the types allowed on the Prolog side of the declaration.

In all modes and types, checks are performed to ensure the types of the arguments. Also, all
arguments of type '-' are checked to be free variables at call time.

2.6   Compiling Foreign Modules on Windows and under Cygwin

Due to the complexity of creating makefiles for the different compilers under Windows, XSB doesn't
attempt to compile and build DLL's for the Windows foreign modules automatically. However, for
almost all typical cases the user should be able to easily adapt the sample makefile for Microsoft
VC++:

      XSB/examples/XSB_calling_c/MakefileForCreatingDLLs

It is important that the C program will have the following lines near the top of the file:

```
#include "xsb_config.h"
#ifdef WIN_NT
#define XSB_DLL
#endif
#include "cinterf.h"
```

Note that these same DLLs will work under Cygwin — XSB's C interface under Cygwin is like
that under Windows rather than Unix.

If the above makefile cannot be adapted, then the user has to create the DLL herself. The
process is, roughly, as follows: first, compile the module from within XSB. This will create the
XSB-specific object file, and (if using the higher-level C interface) the *wrappers*. The *wrappers* are
created in a file named `xsb_wrap_`*modulename*`.c`.

Descriptive Type	Mode Usage	Associated C Type	Comments
int	+	int	integer numbers
float	+	double	floating point numbers
atom	+	unsigned long	atom represented as an unsigned long
chars	+	char *	the textual representation of an atom is passed to C as a string
chars(*size*)	+	char *	the textual representation of an atom is passed to C as a string in a buffer of size *size*
string	+	char *	a prolog list of characters is passed to C as a string
string(*size*)	+	char *	a prolog list of characters is passed to C as a string
term	+	prolog_term	the unique representation of a term
intptr	+	int *	the location of a given integer
floatptr	+	double *	the location of a given floating point number
atomptr	+	unsigned long *	the location of the unique representation of a given atom
charsptr	+	char **	the location of the textual representation of an atom
stringptr	+	char **	the location of the textual representation of a list of characters
termptr	+	prolog_term *	the location of the unique representation of a term
intptr	-	int *	the integer value returned is passed to Prolog
floatptr	-	double *	the floating point number is passed back to Prolog
charsptr	-	char **	the string returned is passed to Prolog as an atom
stringptr	-	char **	the string returned is passed back as a list of characters
atomptr	-	unsigned long *	the number returned is passed back to Prolog as the unique representation of an atom
termptr	-	prolog_term *	the number returned is passed to Prolog as the unique representation of a term
chars(*size*)	+-	char *	the atom is copied from Prolog to a buffer, passed to C and converted back to Prolog afterwards
string(*size*)	+-	char *	the list of characters is copied from Prolog to a buffer, passed to C and back to Prolog afterwards
intptr	+-	int *	an integer is passed from Prolog to C and from C back to Prolog
floatptr	+-	double *	a float number is passed from Prolog to C, and back to Prolog
atomptr	+-	unsigned long *	the unique representation of an atom is passed to C, and back to Prolog
charsptr	+-	char **	the atom is passed to C as a string, and a string is passed to Prolog as an atom
stringptr	+-	char **	the list of characters is passed to C, and a string passed to Prolog as a list of characters
termptr	+-	prolog_term *	the unique representation of a term is passed to C, and back to Prolog

Table 2.1: Allowed combinations of types and modes, and their meanings

Then, create a project, using the compiler of choice, for a dynamically-linked library that exports symbols. In this project, the user must include the source code of the module along with the *wrapper* created by XSB. This DLL should be linked against the library

```
XSB\config\x86-pc-windows\bin\xsb.lib
```

which is distributed with XSB. In VC++, this library should be added as part of the linkage specification. In addition, the following directories for included header files must be specified as part of the preprocessor setup:

```
XSB\config\x86-pc-windows
XSB\prolog_includes
XSB\emu
```

In VC++, make sure you check off the "No precompiled headers" box as part of the "Precompiled headers" specification. All these options are available through the `Project>>Settings` menu item.

2.7   Functions for Use in Foreign Code

In addition to functions for passing data between Prolog an C, XSB contains other functions that may be useful in Foreign C code. We mention a few here that pertain to throwing exceptions from C code (cf. Volume 1 Chapter 8: *Exception Handling*). These functions can be used by code that uses either the lower- or higher-level interface.

void xsb_domain_error(CTXTdeclc char *valid_domain,Cell culprit,char *pred,int arity,int arg)

Used to throw an ISO-style domain error from foreign code, indicating that `culprit` is not in domain `valid_domain` in argument `arg` of `pred/arity`.

Example: The code fragment

```
Cell num;
:
xsb_domain_error(CTXTc "not_less_than_zero",num,"atom_length",2,2);
```

in `atom_length/2` gives rise to the behavior

```
| ?- atom_length(abcde,-1).
++Error[XSB/Runtime/P]: [Domain (-1 not in domain not_less_than_zero)]
            in arg 2 of predicate atom_length/2)
```

void xsb_existence_error(CTXTdeclc char *objType,Cell culprit,char *pred,int arity,int arg)

Used to throw an ISO-style existence error from foreign code, indicating that an object `culprit` of type `objType` does not exist, in argument `arg` of `pred/arity`.

Example: The code fragment

```
    Cell tid;
    :
    xsb_existence_error(CTXTc "thread",reg[2],"xsb_thread_join",1,1);
```

in `thread_join/1` gives rise to a the behavior

```
| ?- thread_join(7).
++Error[XSB/Runtime/P]: [Existence (No thread 1 exists)]
             in arg 1 of predicate thread_join/1)
```

if a thread with thread id 7 does not exist.

void **xsb_instantiation_error**(CTXTdeclc char *pred,int arity,int arg,char *state)
Used to throw an ISO-style instantiation error from foreign code. If `state` is a NULL pointer, the message indicates that there is an instantiation error for argument `arg` of of `pred/arity`. If `state` is non-NULL, the message additionally indicates that argument `arg` must be `state`.

Example: The code fragment

```
    xsb_instantiation_error(CTXTc "atom_length",2,1,NULL);
```

in `atom_length/2` gives rise to a the behavior

```
| ?- atom_length(X,Y).
++Error[XSB/Runtime/P]: [Instantiation]  in arg 1 of predicate atom_length/2
```

void **xsb_misc_error**(CTXTdeclc char *message,char *pred,int arity)
Used to throw a non ISO-error from foreign code, printing `message` and indicating that the error arose in `pred/arity`.

void **xsb_permission_error**(CTXTdeclc char *op,char *obj,Cell culprit,char *pred,int arity)

Used to throw an ISO-style permission error from foreign code, indicating that an operation of type `op` on type `obj` is not permitted on `culprit`, in argument `arg` of `pred/arity`.

Example: The code fragment

```
xsb_permission_error(CTXTc "unlock mutex","mutex not held by thread",
                     xsb_thread_id,"mutex_unlock",2);
```

in `mutex_unlock/1` gives rise to a the behavior

```
| ?- mutex_unlock(mymut).
++Error[XSB/Runtime/P]: [Permission (Operation) unlock mutex on mutex not held
 by thread: 0] in predicate mutex_unlock/1)
```

if thread 0 does not own mutex `mymut`.

void xsb_resource_error(CTXTdeclc char *resource,char *pred,int arity)
>   Used to indicate that there are not sufficient resources of type `resource` for `pred/arity` to
>   succeed.
>
>   **Example**: The code fragment
>
>    ```
>    xsb_resource_error(th,"system threads","thread_create",2);
>    ```
>
>   in thread_create/1 gives rise to a the behavior
>
>    ```
>    | ?- thread_create(X).
>    ++Error[XSB/Runtime/P]: [Resource (system threads))] in predicate thread_create/2)
>    ```
>
>   If the number of system threads has been exceeded.

void xsb_type_error(CTXTdeclc char *valid_type,Cell culprit,char *pred,int arity,int arg)

>   Used to throw an ISO-style type error from foreign code, indicating that `culprit` is not
>   in ISO type `valid_type` in argument `arg` of `pred/arity`.
>
>   **Example**: The code fragment
>
>    ```
>        Cell num;
>        :
>        if (!isinteger(num)) xsb_type_error(CTXTc "integer",num,"atom_length",2,2);
>    ```
>
>   in atom_length/2 gives rise to the behavior
>
>    ```
>    | ?- atom_length(foo,a).
>    ++Error[XSB/Runtime/P]: [Type (a in place of integer)] in arg 2
>    of predicate atom_length/2)
>    ```

void xsb_throw(CTXTdeclc prolog_term Ball)
>   Used to throw a Prolog term from C code, when an ISO-style error is not required. The
>   term can be caught and handled by the Prolog predicate `catch/3` just as any other thrown
>   term; however if it is not caught, XSB's default error handler will treat it as an unhandled
>   exception.

Chapter 3

Embedding XSB in a Process

There are many situations in which it is desirable to use XSB as a rule- or constraint- processing subcomponent of a larger system that is written in another language. Depending on the intended architecture, it may be appropriate for XSB to reside in its own process, separate from other components of an application, and communicating through sockets, a database, or some other mechanism. However it is often useful for XSB to reside in the same process as other components. To do this, one wants to be able to *call* XSB from the host language, providing queries for XSB to evaluate, and retrieving back the answers. An interface for calling XSB from C is provided for this purpose and is described in this chapter. Based on this C interface, XSB can also be called from Java either through a JNI or a socket-based interface, as described in the documentation for Interprolog, available through `xsb.sourceforge.net`. To call XSB from Visual Basic, a DLL is created as described in this chapter, and additional declarations must be made in visual basic as described in the web page "How to use XSB DLL from Visual Basic" <http://xsb.sourceforge.net/vbdll.html>. In addition, the interface described in this chapter has also been extended to allow XSB to be called from Delphi and Ruby. However, since all of these interfaces – Java, Ruby, Delphi and Visual Basic – depend on XSB's C API, we refer in this chapter to C programs or threads calling XSB, although each of the examples suitably modified can be extended to other calling languages.

New to Version 3.1 are extensions to the C API to allow multiple XSB threads to be called from multiple C threads [1]. In this Chapter, we provide an overview of XSB's C API, and then elaborate its use through a series of examples, beginning with a single XSB thread called by a single C thread, then showing how a C thread can interact with multiple XSB threads, and finally discuss how multiple XSB threads can interact with multiple POSIX threads. Finally, Section 3.3 describes each C function in the API.

3.1   Calling XSB from C

XSB provides several C functions (declared in `$XSBDIR/emu/cinterf.h` and defined in `$XSBDIR/emu/cinterf.c`), which can be called from C to interact with XSB as a subroutine. These

---

[1]XSB's threading model is based on POSIX threads, which can be called in Windows through a variety of POSIX APIs – see Volume 1 chapter 8 *Multi-threaded Programming in XSB*.

functions allow a C program to interact with XSB in a number of ways.

- XSB may be initialized, using most of the parameters available from the command-line.

- XSB may then execute a series of *commands* or *queries*. A command is a deterministic query which simply succeeds or fails without performing any unification on the query term. On the other hand, a non-deterministic query can be evaluated so that its answer substitutions are retrieved one at a time, as they are produced, just as if XSB were called on a command line. Alternately a non-deterministic query can be closed in the case where not every answer to the query is needed. Only one query per thread can be active at a time. I.e., an application must completely finish processing one query to a given thread T (either by retrieving all the answers for it, or by issuing a call to `xsb_close_query()`, before trying to evaluate another using T.

- Finally, XSB can be closed, so that no more queries can be made to any XSB threads.

In general, while any functions in the C API to XSB can be intermixed, the functions can be classified as belonging to three different levels.

- *A VarString level* which uses an XSB-specific C-type definition for variable-length strings (Section 3.4), to return answers.

- *A fixed-string level* provides routines that return answers in fixed-length strings.

- *A register-oriented level* that requires users to set up queries by setting registers for XSB which are made globally available to calling functions. The mechanisms for this resemble the lower-level C interface discussed in Chapter 2. This level of interface should only be used for the single-threaded applications, as it is difficult to prevent race-conditions at this level of interface when multiple C threads are used to call XSB.

The appropriate level to use depends on the nature of the calling program, the speed desired, and the expertise of the programmer. By and large, functions in the `VarString` level are the the easiest and safest to use, but they depend on a C type definition that may not be available to all calling programs (e.g. it may be difficult to use if the calling program is not directly based on C, such as Visual Basic or Delphi). For such applications functions from the fixed-string level would need to be used instead. In general, most applications should use either functions from the `VarString` or the fixed-string level, rather than the register-oriented level. This latter level should only be used by programmers who are willing to work at a low interface level, when the utmost speed is needed by an application, and when multiple threads do not need to interact with XSB.

3.2   Examples of Calling XSB

We introduce a series of examples of how XSB would be called using the string-level interfaces. Simple examples of the register-level interface are given in the `XSB/examples/c_calling_XSB` sub-directory, in files `cmain.c`, `cmain2.c`, `ctest.P`, and `Makefile`, but are not discussed in this section.

We structure out discussion by first showing how to construct a C program to call the single-threaded engine alone in Section 3.2.1. This example is mostly pedagogic: with a small amount of extra coding a C program can be constructed to call both the single- and the multi-threaded engine, and these extensions are discussed in Section 3.2.2. Next, we show how to a C program can call and manage multiple XSB threads in Section 3.2.3. Finally, we show how multiple XSB threads can interact with multiple C threads in Section 3.2.3.

3.2.1 The XSB API for the Sequential Engine Only

We start with a simple program shown, in Figure 3.1, that will call the following XSB predicate

```
  p(a,b,c).
  p(1,2,3).
  p([1,2],[3,4],[5,6]).
  p(A,B,A).

r(c,b,a).
r(3,2,1).
r([5,6],[3,4],[1,2]).
r(_A,B,B).
```

and backtrack through unifying answers (cf. `$XSBDIR/examples/c_calling_xsb/edb.P`). . This example will only compile properly if the sequential engine is used, and its style is `not` recommended: it will be shown in Section 3.2.2 how to extend the style.

We discuss the program in Figure 3.1 in detail. This program, slightly modified so that it compiles with the multi-threaded engine is in `$XSBDIR/examples/c_calling_xsb/cvartest.c`. An executable for this program can be make most easily by calling `$XSBDIR/examples/c_calling_xsb/make.P`, which makes the executable `cvstest`.

The program begins by including some standard C headers: note that `string.h` is needed for string manipulation routines such as `strcpy`. In addition, the XSB library header `cinterf.h` is necessary for the XSB C API. Since the program in Figure 3.1 uses functions in the `VarString` interface, within `main()` the routine `XSB_StrDefine(return_string)` declares and initializes a structure of type `VarString`, named `return_string`.

The next order of business is to initialize XSB. In order to do this, `xsb_init_string()` needs to know the installation directory for XSB, which must be passed as part of the initialization string. In Figure 3.1 this is done by manipulating the path of the executable (`cvstest`) that calls XSB. In fact any other approach would also work as long as the XSB installation directory were passed. Within the initialization string, other command line arguments can be passed to XSB if desired with the following exceptions: the arguments `-B` (boot module), `-D` (command loop driver), `-i` (interpreter) and `-d` (disassembler) cannot be used when calling XSB from a foreign language [2]. As a final point on initialization, note that the function `xsb_init()` can also be used to initialize XSB based on an argument vector and count (see Section 3.3).

---

[2]In previous versions of XSB, initialization from the C level required a `-n` option to be passed. This is no longer required.

```
#include <stdio.h>
#include <string.h>

/* cinterf.h is necessary for the XSB API, as well as the path manipulation routines*/
#include "cinterf.h"

extern char *xsb_executable_full_path(char *);
extern char *strip_names_from_path(char*, int);

int main(int argc, char *argv[]) {

  char init_string[1024];
  int rc;
  XSB_StrDefine(return_string);

  /* xsb_init_string() relies on the calling program to pass the absolute or relative
     path name of the XSB installation directory. We assume that the current
     program is sitting in the directory ../examples/c_calling_xsb/
     To get the installation directory, we strip 3 file names from the path. */

  strcpy(init_string,strip_names_from_path(xsb_executable_full_path(argv[0]),3));

  if (xsb_init_string(init_string) == XSB_ERROR) {
    fprintf(stderr,"++initializing XSB: %s/%s\n",xsb_get_init_error_type(),
            xsb_get_init_error_message());
    exit(XSB_ERROR);
  }

  /* Create command to consult a file: edb.P, and send it. */
  if (xsb_command_string("consult('edb.P').") == XSB_ERROR)
    fprintf(stderr,"++Error consulting edb.P: %s/%s\n",xsb_get_error_type(),xsb_get_error_message());

  rc = xsb_query_string_string("p(X,Y,Z).",&return_string,"|");
  while (rc == XSB_SUCCESS) {
    printf("Return %s\n",(return_string.string));
    rc = xsb_next_string(&return_string,"|");
  }

 if (rc == XSB_ERROR)
    fprintf(stderr,"++Query Error: %s/%s\n"xsb_get_error_type(),xsb_get_error_message());

  xsb_close();
}
```

Figure 3.1: Calling the Sequential Engine Using the `VarString` Interface

Note that the calling program checks for any errors returned by `xsb_init_string()` and other API commands. In general, `xsb_init_string()` may throw an error if the XSB's installation directory has become corrupted, or for similar reasons. This mechanism for error handling is different than that used if XSB is called in its usual stand-alone mode, in which case such an error would cause XSB to exit). An error returned by XSB's API are similar to an error ball described in Volume 1 *Exception Handling* in that it has both a *type* and a *message*. For normal Prolog exceptions, XSB's API will throw the same kinds of errors as XSB called in a stand-alone (or server) mode, i.e. instantiation errors, type errors, etc. However XSB's API adds two new error types:

- `init_error` is used as the type of an error discovered upon initialization of XSB, before query and command processing has begun. If an `init_error` is raised, XSB has not been properly initialized and will not run.

- `unrecoverable_error` is used to indicate that XSB has encountered an error, (such as a memory allocation error), during command or query processing from which it cannot recover. Such an error would cause XSB to immediately exit if it were called in a stand-alone mode. In general the calling program should handle unrecoverable errors as fatal since there is a good chance that the error conditions will affect the calling program as well as XSB.

Errors raised by `xsb_init_string()` usually have type `init_type`.

and a string pointer to the associated message can be found by the function `xsb_get_init_error_message()`.

As can be seen from the example, handling errors from commands is done in manner similar to that of initialization. For non-initialization errors, a string pointer to the type can be obtained by `xsb_get_error_type()`, while a string pointer to the message can be obtained by `xsb_get_error_message()`.

Next in Figure 3.1 the file `edb.P` is consulted (containing the `p/3` and `r/3` predicates shown above). Note, that the argument to `xsb_command_string` must be a syntactically valid Prolog term ending with a period, otherwise a syntax error will be thrown, which may be displayed through `xsb_get_error_type()` and `xsb_get_error_message()` [3].

Queries to XSB are a little more complicated than commands. Since a query may return multiple solutions, a query should usually be called from inside a loop. In Figure 3.1, the query is opened with `xsb_query_string()`. If the query has at least one answer, `xsb_query_string()` will return `XSB_SUCCESS`; if the query fails, it will return `XSB_FAILURE`, and if there is an exception it will return `XSB_ERROR` as usual. Any answer will be returned as a string in the `VarString return_string`, and each argument of the query will be separated by the character `|`. Thus, in our example, the first answer will write the string

```
a|b|c
```

Once a query has been opened, subsequent answers can be obtained via `xsb_next_string()`. These answers are written to `return_string` in the same manner as `xsb_query_string_string()`.

---

[3]Most XSB errors are handled in this manner when XSB is called through its API. A few errors will print directly to `stderr` and some XSB warnings will print to `stdwarn` which upon startup is dup-ed to `stderr`.

```
1|2|3
[1,2]|[3,4]|[5,6]
_h102|_h116|_h102
```

A query is automatically closed when no more answers can be derived from it. Alternately, a query that may have answers remaining can be closed using the command xsb_close_query(). If the calling application will need to pass more queries or commands to XSB nothing need be done at this point: a new queries or commands can be invoked using one of the functions just discussed. However if the calling process is finished with XSB and will never need it again during the life of the process, it can call xsb_close().

An Example using Fixed Strings

```
  int retsize = 15;
  char *return_string;
  int anslen;

  return_string = malloc(retsize);

  rc = xsb_query_string_string_b(CTXTc "p(X,Y,Z).",return_string,retsize,&anslen,"|");

  while (rc == XSB_SUCCESS || rc == XSB_OVERFLOW) {

    if (rc == XSB_OVERFLOW) {
      return_string = (char *) realloc(return_string,anslen);
      return_size = anslen;
      rc = xsb_get_last_answer_string(CTXTc return_string,retsize,&anslen);
    }

    printf("Return %s %d\n",return_string,anslen);
    rc = xsb_next_string_b(CTXTc return_string,15,&anslen,"|");
  }
```

Figure 3.2: Calling XSB using the Fixed String Interface

Figure 3.2 shows a fragment of code indicating how the previous example would be modified if the fixed-string interface were used. Note that return_string now becomes a pointer to explicitly malloc-ed memory. To open the query p(X,Y,Z) the function xsb_query_string_string_b() is called, with the _b indicating that a fixed buffer is being used rather than a VarString. The call is similar to xsb_query_string_string(), except that the length anslen of the buffer pointed to by return_string is now also required. If the answer to be returned (including separators) is longer than anslen, xsb_query_string_string_b() will return XSB_OVERFLOW. If this happens, a new answer buffer can be used (here the old one is realloc-ed) and the answer retrieved via xsb_get_last_answer_string. Similarly, further answers are obtained via xsb_next_string_b() whose length must be checked. Thus the only difference between the fixed-string level and the VarString level is that the length of each answer should be checked and xsb_get_last_answer_string() called if necessary.

3.2.2 The General XSB API

The previous section showed how to use the XSB API with both the `VarString` type and without, but did not consider the multi-threaded engine. In fact, there are different ways to use XSB's multi-threading that can have advantages for various situations. In the first mode, threads are managed from Prolog, with a single XSB thread called from the API; that XSB thread can then create another XSB thread that does work, and the first thread can return almost immediately to handle more requests from the API's caller. A second model allows the caller to manipulate a pool of several XSB threads, so that different XSB threads may be called from different threads over the API. In this model each C, Java, Ruby, or other thread could a number of different Prolog threads. In this section we sketch how to use the API to illustrate the first model, and sketch the second model in the next section.

Figure 3.3 shows how relevant portions of the previous `VarString` example can be adapted to use the multi-threaded engine. The main change is that a new variable is introduced on the C side that points to the context of the main thread. As pointed out in Chapter 2, each thread in the multi-threaded engine has a *context* in which is kept much of its thread-specific data (excluding tables and dynamic code). Of the threads running in the multi-threaded engine the thread created upon the call to `xsb_init()` is designated as the `main thread`, and is closed only upon calling `xsb_close()`.

Within the multi-threaded engine, a call to an API function such as `xsb_query_string_string()` is actually a call to a specific thread to do some work (using a thread context pointer). Accordingly, since any errors produced will be specific to a given thread, all calls to error reporting functions are also thread-specific. If no specific thread is needed, it may be best just to use the main thread, which is what is done in Figure 3.3. The thread context pointer `th` is initialized to the main thread using the API macro `xsb_get_main_thread()`. Afterwards, this pointer is passed into the various interface functions by making use of XSB macros defined in `context.h` In the multi-threaded engine, these macros are defined as

```
#define CTXT th
#define CTXTc th,
```

while in the single-threaded engine they are defined as empty strings, as is `xsb_get_main_thread()`. As a result the code in Figure 3.3 will compile and run properly both for the single-threaded and the multi-threaded engines.

At this stage, suppose one wanted a new thread to execute a specific command, say `do_foo`. In this case, a C call such as

```
xsb_query_string_string(CTXTc "thread_create(do_foo,Id).",&return_string,"|")
```

creates a thread to execute the command, and returns the thread id of the newly created thread in `return_string`. The behavior of this newly created thread is exactly the same as if it were created from the XSB command line: in particular the newly created thread will automatically exit upon completion of its command. As a somewhat technical point, there are two different ways of referring to XSB threads. The foreign language interfaces described in Chapter 2 and here use

```
.....

/* context.h is necessary for the type of a thread context. */
#include "context.h"

int main(int argc, char *argv[])
{

  char init_string[MAXPATHLEN];
  int rc;
  XSB_StrDefine(return_string);

  strcpy(init_string,strip_names_from_path(xsb_executable_full_path(argv[0]),3));
  if (xsb_init_string(init_string) == XSB_ERROR) {
    fprintf(stderr,"++initializing XSB: %s/%s\n",xsb_get_init_error_type(),
            xsb_get_init_error_message());
    exit(XSB_ERROR);
  }

#ifdef MULTI_THREAD
  th_context *th = xsb_get_main_thread();
#endif

  /* Create command to consult a file: edb.P, and send it. */
  if (xsb_command_string(CTXTc "consult('edb.P').") == XSB_ERROR)
    fprintf(stderr,"++Error consulting edb.P: %s/%s\n",xsb_get_error_type(CTXT),
            xsb_get_error_message(CTXT));

  rc = xsb_query_string_string(CTXTc "p(X,Y,Z).",&return_string,"|");
  while (rc == XSB_SUCCESS) {
    printf("Return %s\n",(return_string.string));
    rc = xsb_next_string(CTXTc &return_string,"|");
  }

 if (rc == XSB_ERROR)
    fprintf(stderr,"++Query Error: %s/%s\n",xsb_get_error_type(CTXT),xsb_get_error_message(CTXT));

  xsb_close();
}
```

Figure 3.3: Calling the Single- or Multi-Threaded Engine Using the `VarString` Interface

pointers to thread contexts so that the interfaces use much of the same code as the XSB engine. However Prolog refers to threads using thread identifiers. The two different forms can be converted into each other by the functions `xsb_thread_id_to_context()` and `xsb_thread_context_to_id()`.

3.2.3 Managing Multiple XSB Threads through the API

The ability to pass thread contexts into query and command functions allows a great deal of flexibility [4]. Once XSB is initialized, XSB threads can be created from C and can execute independently of each other, effectively giving the ability for different calling threads to query XSB in a mechanism reminiscent of database cursors.

Figure 3.4 illustrates a very simple example of this. XSB is initialized and the file `edb.P` consulted exactly as in Figure 3.4. However, the function `xsb_ccall_thread_create()` causes the XSB thread `p_th` to create a new thread, causes the new thread to call the same command loop as the main thread, and sets `r_th` to point to the context of the new thread. The new thread `r_th` can be used for commands or queries just as `p_th`. Figure 3.4 shows that queries to the two threads can be interleaved, and errors for both threads can be checked and reported independently.

It is important to note that since each thread created by `xsb_ccall_thread_create()` goes into a command-loop similar to the command loop, it will stay around until it is explicitly killed or until XSB is closed. The call

```
xsb_kill_thread(r_th);
```

is needed to make `r_th` to exit. Once a thread is exited, all of its data structures will be freed, including those that support `xsb_get_error_type()` and `xsb_get_error_message()` [5].

3.2.4 Calling Multiple XSB Threads using Multiple C Threads

Figure 3.4 shows how two XSB threads can be created, can receive different queries and can interleave their backtracking and answer return. Although Figure 3.4 demonstrated only backtracking through simple predicates, the mechanism employed works for complicated examples using tabling, dynamic code, and other features. All this provides a sophisticated interface, but it is not "fully" multi-threaded in the following sense. When a C thread T causes XSB to execute a command or query the thread must wait until the calling function returns before proceeding. In certain applications it may be useful, for example, for T to create a C thread T_{new} which runs asynchronously from T, executing the XSB command or query and then exiting. Alternately, an application may want to have a pool of C threads that can interact with a pool of XSB threads.

XSB's C API has been designed to support these features. Figure 3.5 shows fragments of Figure 3.4 rewritten so that the routines to print out the answers to the queries `p(X,Y,Z)` and `r(X,Y,Z)` can be called from C threads specially designed for this purpose. More specifically, the routine `query_ps()` calls `p_th` to query `p(X,Y,Z)` and backtrack through its answers – its use of a

---

[4]For the sake of brevity, we sometimes abuse notation and do not always distinguish between thread-contexts and their pointers.

[5]Note that causing XSB's main thread to exit will cause the entire process to exit – not just XSB.

```
.....
/* context.h is necessary for the type of a thread context. */
#include "context.h"

int main(int argc, char *argv[])
{

  static th_context *p_th, *r_th;
  char init_string[MAXPATHLEN];
  int rcp, rcr;
  XSB_StrDefine(p_return_string);
  XSB_StrDefine(r_return_string);

  strcpy(init_string,strip_names_from_path(xsb_executable_full_path(argv[0]),3));

  if (xsb_init_string(init_string)) {
    fprintf(stderr,"%s initializing XSB: %s/%s\n",xsb_get_init_error_type(),
            xsb_get_init_error_message());cin
    exit(XSB_ERROR);
  }

  p_th = xsb_get_main_thread();

  /* Create command to consult a file: edb.P, and send it. */
  if (xsb_command_string(p_th, "consult('edb.P').") == XSB_ERROR)
    fprintf(stderr,"++Error consulting edb.P: %s/%s\n",xsb_get_error_type(p_th),
    xsb_get_error_message(p_th));

  xsb_ccall_thread_create(p_th,&r_th);

  rcp = xsb_query_string_string(p_th,"p(X,Y,Z).",&p_return_string,"|");
  rcr = xsb_query_string_string(r_th,"r(X,Y,Z).",&r_return_string,"|");

  while (rcp == XSB_SUCCESS && rcr == XSB_SUCCESS) {

    printf("Return p %s\n",(p_return_string.string));
    rcp = xsb_next_string(p_th, &p_return_string,"|");

    printf("Return r %s\n",(r_return_string.string));
    rcr = xsb_next_string(r_th, &r_return_string,"|");
  }

 if (rcp == XSB_ERROR)
   fprintf(stderr,"++Query Error p: %s/%s\n",xsb_get_error_type(p_th),xsb_get_error_message(p_th));
 if (rcr == XSB_ERROR)
   fprintf(stderr,"++Query Error r: %s/%s\n",xsb_get_error_type(r_th),xsb_get_error_message(r_th));

 xsb_close();
}
```

Figure 3.4: Manipulating Multiple Threads Using the `VarString` Interface

single `void *` argument and a `void *` return reflect the requirements of functions that are to be called using `pthread_create()`.

We note several points about this example. First the XSB API is a low-level API that can be used to build application specific interfaces, and some experience with pthread programming is useful if multiple XSB threads are called from multiple C threads. For instance, one issue is fairness. When called from the C API each XSB thread X_T makes use of mutexes to ensure that it answers only one query or command at a time. If multiple C threads are are waiting for X_T to respond to requests or queries, there is no guarantee that the requests will be processed in any sort of order, or even that a request will eventually be handled (In order to ensure this, the calling program would have to use a queue or some other scheduling mechanism to send requests to the XSB thread). In addition, it is important to note that, *the main XSB thread should only be called from the C thread that initialized XSB.*. This restriction is due to the current design of synchronizing an XSB thread with calling threads, and may be lifted in the future.

Protected and Non-Protected API Functions

Example 3.5 shows that, when the `Varstring` functions are used, if a single calling thread opens a query to an XSB thread X_T, X_T will be protected from queries and commands posed by other C threads until the query is closed, failed out of, or exits via an error. In fact, queries (and commands) are protected when the `Varstring` or fixed string interfaces are used. However, consider what may happen when the register level interface is used. In this case, a calling thread may call one or more API functions to set up the registers, execute a command or query, call several more API functions to obtain the output, and so on. For this reason, if an application uses API commands that depend on user manipulation of registers (`xsb_command()`,`xsb_query()`,`xsb_query_string()`, and `xsb_next()`) the user must ensure that only one calling thread interacts with an XSB thread when that thread in the course of executing a command or query. See `$XSB_DIR/examples/c_calling_xsb/cregs_thread` for an example of how mutexes can be used to protect XSB threads.

3.3   A C API for XSB

3.3.1   Initializing and Closing XSB

`int xsb_init_string(char *options)`
> This function is used to initialize XSB via an initialization string `*options`, and must be called before any other calls can be made. The initialization string must include the path to the XSB directory installation directory `$XSB_DIR`, which is expanded to an absolute path by XSB. Any other command line options may be included just as in a command line except `-D`, `-d`, `-B` and `-i`. For example, a call from an executable in a sibling directory of XSB might have the form
>
>   `xsb_init_string("../XSB -e startup.");`
>
> which initializes XSB with the goal `?- startup.`
>
> **Return Codes**

```
.....

void *query_ps(void * arg) {
  int rc;
  th_context *p_th;
  XSB_StrDefine(p_return_string);
  p_th = (th_context *)arg;

  rc = xsb_query_string_string(p_th,"p(X,Y,Z).",&p_return_string,"|");
  while (rc == XSB_SUCCESS) {
    printf("Return p %s\n",(p_return_string.string));
    rc = xsb_next_string(p_th, &p_return_string,"|");
  }

 if (rc == XSB_ERROR)
   fprintf(stderr,"++Query Error p: %s/%s\n",xsb_get_error_type(p_th),xsb_get_error_message(p_th));
 return NULL;
}


int main(int argc, char *argv[]) {

  char init_string[MAXPATHLEN];
  static th_context *p_th, *r_th;
  int pstatus, rstatus;
  pthread_t pthread_id,rthread_id;
  XSB_StrDefine(p_return_string);
  XSB_StrDefine(r_return_string);

.....

  main_th = xsb_get_main_thread();

  /* Create command to consult a file: edb.P, and send it. */
  if (xsb_command_string(xsb_get_main_thread(), "consult('edb.P').") == XSB_ERROR)
    fprintf(stderr,"++Error consulting edb.P: %s/%s\n",xsb_get_error_type(main_th),
    xsb_get_error_message(main_th));

  xsb_ccall_thread_create(main_th,&r_th);
  xsb_ccall_thread_create(main_th,&p_th);

  pthread_create(&rthread_id,NULL,command_rs,r_th);
  pthread_create(&pthread_id,NULL,command_ps,p_th);
  pthread_create(&rthread_id,NULL,command_rs,r_th);
  pthread_create(&pthread_id,NULL,command_ps,p_th);

  rstatus = pthread_join(rthread_id,&rreturn);
  if (rstatus != 0) fprintf(stderr,"R join returns status %d\n",rstatus);
  pstatus = pthread_join(pthread_id,&preturn);
  if (pstatus != 0) fprintf(stderr,"P join returns status %d\n",pstatus);

 xsb_kill_thread(r_th);
 xsb_close();
}
```

Figure 3.5: Manipulating Multiple XSB Threads Using Multiple C Threads

- `XSB_SUCCESS` indicates that initialization returned successfully.

- `XSB_ERROR`

  - `init_error` if any error occurred during initialization.

  - `permission_error` if `xsb_init_string()` is called after XSB has already been correctly initialized.

`int xsb_init(int argc, char *argv[])`

This function is a variant of `xsb_init_string()` which passes initialization arguments as an argument vector: `argc` is the count of the number of arguments in the `argv` vector. The `argv` vector is exactly as would be passed from the command line to XSB.

- `argv[0]` must be an absolute or relative path name of the XSB installation directory (*i.e.*, `$XSB_DIR`). Here is an example, which assumes that we invoke the C program from the XSB installation directory.

```
int main(int argc, char *argv[])
{
  int myargc = 1;
  char *myargv[1];

  /* XSB_init relies on the calling program to pass the addr of the XSB
     installation directory. From here, it will find all the libraries */
  myargv[0] = ".";

  /* Initialize xsb */
  xsb_init(myargc,myargv);
}
```

The return codes for `xsb_init()` are the same as those for `xsb_init_string()`.

`int xsb_close()`

This routine closes the entire connection to XSB . After this, no more calls can be made (not even calls to `xsb_init_string()` or `xsb_init()`). In Version 3.2, no guarantee is made that all space used by XSB will be restored to the process (even when the process has dynamically linked to XSB), but space for any XSB tables is freed.

Return Codes

- `XSB_SUCCESS` indicates that XSB was closed successfully.

- `XSB_ERROR`

  - `permission_error` if `xsb_closed()` when XSB has not been (correctly) initialized.

3.3.2 Passing Commands to XSB

`int xsb_command_string(th_context *th, char *cmd)`

This function passes a command to the XSB thread designated by `th` (the first argument is

not used in the single-threaded engine). No query can be active in `th` when the command is called. The command is a string consisting of a Prolog (or HiLog) term terminated by a period (.).

When used in the multi-threaded engine, `xsb_command_string` protects the called thread from API calls from other pthreads until the command is finished.

Return Codes

- `XSB_SUCCESS` indicates that the command succeeded.
- `XSB_FAILURE` indicates that the command failed.
- `XSB_ERROR`
  - `permission_error` if `xsb_command_string()` is called while a query is open in `th`.
  - Otherwise, any queries thrown during execution of the command are accessable through `xsb_get_error_type(th)` and `xsb_get_error_message(th)`.

`int xsb_command(th_context *th)`

This function passes a command to the XSB thread designated by `th` (the first argument is not used in the single-threaded engine). Any previous query must have already been closed. Before calling `xsb_command()`, the calling program must construct the term representing the command in register 1 in the XSB thread's space. This can be done by using the `c2p_*` (and `p2p_*`) routines, which are described in Section 2.2.3 below. Register 2 may also be set before the call to `xsb_query()` (using `xsb_make_vars(int)` and `xsb_set_var_*()`) in which case any variables set to values in the `ret/n` term will be so bound in the call to the command goal. `xsb_command` invokes the command represented in register 1 and returns `XSB_SUCCESS` if the command succeeds, `XSB_FAILURE` if it fails, and `XSB_ERROR` if an error is thrown while executing the command.

When used in the multi-threaded engine, `xsb_command_string` *does not protect* the called thread from API calls from other pthreads until the command is finished. It is the user's responsibility to protect the XSB thread, using a mutex or other concurrency control, from the time the goal begins to be constructed in the register 1 until the command has completed.

Apart from the steps necessary to formulate the query and the lack of protection of the XSB thread, the behavior of `xsb_command()` is similar to that of `xsb_command_string()`, including its return codes.

3.3.3   Querying XSB

`int xsb_query_string_string(th_context *th, char *query, VarString *buff, char *sep)`

This function opens a query to the XSB thread designated by `th` (the first argument is not used in the single-threaded engine); it returns the first answer (if there is one) as a `VarString`. Any previous query to `th` must have already been closed. Any query may return multiple data answers. The first is found and made available to the caller as a result of this call. To get any subsequent answers, `xsb_next_string()` must be called. An example call is:

```
rc = xsb_query_string_string(th, "append(X,Y,[a,b,c]).",buff,";");
```

The second argument is the period-terminated query string. The third argument is a pointer to a variable string buffer in which the subroutine returns the answer (if any.) The variable string data type `VarString` is explained in Section 3.4. (Use `xsb_query_string_string_b()` if you cannot declare a parameter of this type in your programming language.) The last argument is a string provided by the caller, which is used to separate arguments in the returned answer. For the example query, `buff` would be set to the string:

```
[];[a,b,c]
```

which is the first answer to the append query. There are two fields of this answer, corresponding to the two variables in the query, `X` and `Y`. The bindings of those variables make up the answer and the individual fields are separated by the `sep` string, here the semicolon (`;`). In the answer string, XSB atoms are printed without quotes. Complex terms are printed in a canonical form, with atoms quoted if necessary, and lists produced in the normal list notation.

When used in the multi-threaded engine, `xsb_query_string_string` protects the called thread from API calls from other pthreads until the entire query is finished.

Return Codes

- `XSB_SUCCESS` indicates that the query succeeded.
- `XSB_FAILURE` indicates that the query failed.
- `XSB_ERROR`
  - `permission_error` if `xsb_query_string_string()` is called while a query to `th` is open.
  - Otherwise, any errors thrown during execution of the query are accessible through `xsb_get_error_type()` and `xsb_get_error_message()`.

int xsb_query_string_string_b(th_context *th,char *query,char *buff,int bufflen,int *anslen,char *sep)

This function provides a lower-level alternative to `xsb_query_string_string` (not using the `VarString` type), which makes it easier for non-C callers (such as Visual Basic or Delphi) to access XSB functionality. Any previous query to `th` must have already been closed. Any query may return possibly multiple data answers. The first is found and made available to the caller as a result of this call. To get any subsequent answers, `xsb_next_string_b()` or a similar function must be called. The first and last arguments are the same as in `xsb_query_string_string()`. The buff, bufflen, and `anslen` parameters are used to pass the answer (if any) back to the caller. `buff` is a character array provided by the caller in which the answer is returned. `bufflen` is the length of the buffer (`buff`) and is provided by the caller. `anslen` is returned by this routine and is the length of the computed answer. If that length is less than `bufflen`, then the answer is put in `buff` (and null-terminated). If the answer is longer than will fit in the buffer (including the null terminator), then the answer is not copied to the buffer and `XSB_OVERFLOW` is returned. In this case the caller can retrieve the answer by providing a bigger buffer (of size greater than the returned `anslen`) in a call to `xsb_get_last_answer_string()`.

When used in the multi-threaded engine, `xsb_query_string_string_b` protects the called thread from API calls from other pthreads until the entire query is finished.

Return Codes

- XSB_SUCCESS indicates that the query succeeded.

- XSB_FAILURE indicates that the query failed.

- XSB_ERROR

  - permission_error if xsb_query_string_string_b() is called while a query to th is open.

  - Otherwise, any queries thrown during execution of the command are accessible through xsb_get_error_type() and xsb_get_error_message().

- XSB_OVERFLOW indicates that the query succeeded, but the answer was too long for the buffer.

int xsb_query(th_context *th)

This function passes a query to the XSB thread th. Any previous query to th must have already been closed. Any query may return possibly multiple data answers. The first is found and made available to the caller as a result of this call. To get any subsequent answers, xsb_next() or a similar function must be called. Before calling xsb_query() the caller must construct the term representing the query in the XSB thread's register 1 (using routines described in Section 2.2.3 below.) If the query has no answers (i.e., just fails), register 1 is set back to a free variable and xsb_query() returns XSB_FAILURE. If the query has at least one answer, the variables in the query term in register 1 are bound to those answers and xsb_query() returns XSB_SUCCESS. In addition, register 2 is bound to a term whose main functor symbol is ret/n, where n is the number of variables in the query. The main subfields of this term are set to the variable values for the first answer. (These fields can be accessed by the functions p2c_*, or the functions xsb_var_*, described in Section 2.2.3 below.) Thus there are two places the answers are returned. Register 2 is used to make it easier to access them. Register 2 may also be set before the call to xsb_query() (using xsb_make_vars(int) and xsb_set_var_*()) in which case any variables set to values in the ret/n term will be so bound in the call to the goal.

When used in the multi-threaded engine, xsb_query *does not protect* the called thread from API calls from other pthreads until the query is finished, or even when the registers are being accessed. It is the user's responsibility to protect the XSB thread, using a mutex or other concurrency control, from the time the goal begins to be constructed in the register 1 until the query is closed, failed, or exited upon error.

int xsb_get_last_answer_string(th_context *th, char *buff, int bufflen, int *anslen)

This function is used only when a call xsb_query_string_string_b() or xsb_next_string_b() to th returns XSB_OVERFLOW, indicating that the buffer provided was not big enough to contain the computed answer. In that case the user may allocate a larger buffer and then call this routine to retrieve the answer (that had been saved.) Only one answer is saved per thread, so this routine must called immediately after the failing call in order to get the right answer. The parameters are the same as the 2nd through 4th parameters of xsb_query_string_string_b().

Return Codes

- **XSB_OVERFLOW** indicates that the answer was `still` too long for the buffer.

int xsb_query_string(th_context *th,char *query)

This function passes a query to the XSB thread `th`. The query is a string consisting of a term that can be read by the XSB reader. The string must be terminated with a period (.). Any previous query must have already been closed. In all other respects, **xsb_query_string()** is similar to **xsb_query()**, except the only way to retrieve answers is through Register 2. The ability to create the return structure and bind variables in it is particularly useful in this function.

When used in the multi-threaded engine, **xsb_query_string** *does not protect* the called thread from API calls from other pthreads until the query is finished, or even when the registers are being accessed. It is the user's responsibility to protect the XSB thread, using a mutex or other concurrency control, from the time the goal begins to be constructed in the register 1 until the query is closed, failed, or exited upon error.

Return Codes

- **XSB_SUCCESS** indicates that the query succeeded.
- **XSB_FAILURE** indicates that the query failed.
- **XSB_ERROR** indicates that an error occurred while executing the query.

int xsb_next_string(th_context *th,VarString *buff,char *sep)

This routine is called after **xsb_query_string()** to retrieve a subsequent answer in `buff`. If a query is not open in `th`, an error is returned. This function treats answers just as **xsb_query_string_string()**. For example after the example call

```
rc = xsb_query_string_string(th,"append(X,Y,[a,b,c]).",buff,";");
```

which returns with buff set to

```
[];[a,b,c]
```

Then a call:

```
rc = xsb_next_string(th,buff,";");
```

returns with buff set to

```
[a];[b,c]
```

the second answer to the indicated query.

In the multi-threaded engine, **xsb_next_string()** protects the XSB thread from concurrent access by other threads as long as the query was invoked by **xsb_query_string_string(_b)**.

Return Codes

- **XSB_SUCCESS** indicates that the query succeeded.

- **XSB_FAILURE** indicates that the query failed.
- **XSB_ERROR** indicates that an error occurred while executing the query.

int xsb_next_string_b(th_context *th, char *buff, int bufflen, int *anslen, char *sep)

   This function is a variant of **xsb_next_string()** that does not use the **VarString** type. Its parameters are the same as the 3rd through 6th parameters of **xsb_query_string_string_b()**. The next answer to the current query is returned in **buff**, if there is enough space. If the buffer would overflow, this routine returns **XSB_OVERFLOW**, and the answer can be retrieved by providing a larger buffer in a call to **xsb_get_last_answer_string_b()**. In any case, the length of the answer is returned in **anslen**.

In the multi-threaded engine, **xsb_next_string()** protects the XSB thread from concurrent access by other threads as long as the query was invoked by **xsb_query_string_string(_b)**.

Return Codes

- **XSB_SUCCESS** indicates that backtracking into the query succeeded.
- **XSB_FAILURE** indicates that backtracking into the query failed.
- **XSB_ERROR** indicates that an error occurred while further executing the query.
- **XSB_OVERFLOW** indicates that backtracking into the query succeeded, but the new answer was too long for the buffer.

int xsb_next(th_context *)
   This function is called after **xsb_query()** (which must have returned **XSB_SUCCESS**) to retrieve more answers. It rebinds the query variables in the term in register 1 and rebinds the argument fields of the **ret/n** answer term in register 2 to reflect the next answer to the query. Its return codes are as with **xsb_next_string()**.

When used in the multi-threaded engine, **xsb_next** *does not protect* the called thread from API calls from other pthreads until the query is finished, or even when the registers are being accessed. It is the user's responsibility to protect the XSB thread, using a mutex or other concurrency control, through the time that registers are accessed by the calling program.

int xsb_close_query(th_context *th)
   This function allows a user to close a query to **th** before all its answers have been retrieved. Since XSB is (usually) a tuple-at-a-time system, answers that are not retrieved are not computed so that closing a query may save time. If a given query Q is open, it is an error to open a new query without closing Q either by retrieving all its answers or explicitly calling **xsb_close_query()** to close Q. Calling **xsb_close_query()** when no query is open gives an error message, but otherwise has no effect.

Return Codes

- **XSB_SUCCESS** indicates that the current query was closed.
- **XSB_ERROR**
    - **permission_error** if **xsb_close_query()** is called while no query is open.

3.3.4   Obtaining Information about Errors

char * xsb_get_init_error_message()

 Used to find error messages if xsb_init_string() or xsb_init() returns XSB_ERROR. Any
errors returned by these functions have type init_error. Because initialization errors occur
before XSB or any of its threads have been initialized, initialization errors do not require a
thread context for input.

char * xsb_get_error_type(th_context *th)

 If a function called for th returned XSB_ERROR this function provides a pointer to a string rep-
resenting the type of the error. Types are as in Volume 1 *Exception Handling* with the addition
of init_error for errors that occur during initialization of XSB, and unrecoverable_error
for errors from which no recovery is possible for XSB (e.g. inability to allocate new memory).

char *xsb_get_error_message(th_context *th)

 If a function called for th returned XSB_ERROR this function provides a pointer to a string
representing a message associated with the error. For errors raised within the Prolog portion
of execution, messages are as in Volume 1 *Exception Handling.*

3.3.5   Thread Management from Calling Programs

int xsb_ccall_thread_create(th_context *callingThread, th_context **newThread)

 Causes callingThread to create a thread pointed to by newThread. newThread runs exactly
the same interpreter loop as callingThread and all API functions will work on newThread
just as on the main thread, or any other thread. newThread will be non-detached, and will
inherit any private parameters from callingThread. To create a thread to do a specific task
or a detached thread, rather than one that executes a command loop, simply call the query
thread_create/[2,3] from one of the query functions.

th_context *xsb_get_main_thread()

 Returns a pointer to the thread context of XSB's main thread. If XSB has not been initialized
or has been closed this function returns 0.

xsb_tid xsb_thread_id_to_context(th_context *th)

th_context *xsb_thread_context_to_id(xsb_tid id)

3.4   The Variable-length String Data Type

XSB uses variable-length strings to communicate with certain C subroutines when the size of the
output that needs to be passed from the Prolog side to the C side is not known. Variable-length
strings adjust themselves depending on the size of the data they must hold and are ideal for this situ-
ation. For instance, as we have seem the two subroutines xsb_query_string_string(query,buff,sep)

and `xsb_next_string(buff,sep)` use the variable string data type, `VarString`, for their second argument. To use this data type, make sure that

```
#include "cinterf.h"
```

appears at the top of the program file. Variables of the `VarString` type are declared using a macro that must appear in the declaration section of the program:

```
XSB_StrDefine(buf);
```

There is one important consideration concerning VarString with the *automatic* storage class: they must be *destroyed* on exit (see `XSB_StrDestroy`, below) from the procedure that defines them, or else there will be a memory leak. It is not necessary to destroy static `VarString`'s.

The public attributes of the type are `int length` and `char *string`. Thus, `buf.string` represents the actual contents of the buffer and `buf.length` is the length of that data. Although the length and the contents of a `VarString` string is readily accessible, the user **must not** modify these items directly. Instead, he should use the macros provided for that purpose:

- `XSB_StrSet(VarString *vstr, char *str)`: Assign the value of the regular null-terminated C string to the `VarString` `vstr`. The size of `vstr` is adjusted automatically.

- `XSB_StrSetV(VarString *vstr1, VarString *vstr2)`: Like `XSB_StrSet`, but the second argument is a variable-length string, not a regular C string.

- `XSB_StrAppend(VarString *vstr, char *str)`: Append the null-terminated string `str` to the `VarString` `vstr`. The size of `vstr` is adjusted.

- `XSB_StrPrepend(VarString *vstr, char *str)`: Like `XSB_StrAppend`, except that `str` is prepended.

- `XSB_StrAppendV(VarString *vstr1, VarString *vstr2)`: Like `XSB_StrAppend`, except that the second string is also a `VarString`.

- `XSB_StrPrependV(VarString *vstr1, VarString *vstr2)`: Like `XSB_StrAppendV`, except that the second string is prepended.

- `XSB_StrCompare(VarString *vstr1, VarString *vstr2)`: Compares two `VarString`. If the first one is lexicographically larger, then the result is positive; if the first string is smaller, than the result is negative; if the two strings have the same content (*i.e.*, `vstr1->string` equals `vstr2->string` then the result is zero.

- `XSB_StrCmp(VarString *vstr, char *str)`: Like `XSB_StrCompare` but the second argument is a regular, null-terminated string.

- `XSB_StrAppendBlk(VarString *vstr, char *blk, int size)`: This is like `XSB_StrAppend`, but the second argument is not assumed to be null-terminated. Instead, `size` characters pointed to by `blk` are appended to `vstr`. The size of `vstr` is adjusted, but the content is *not* null terminated.

- XSB_StrPrependBlk(VarString *vstr, char *blk, int size): Like XSB_StrPrepend, but blk is not assumed to point to a null-terminated string. Instead, size characters from the region pointed to by blk are prepended to vstr.

- XSB_StrNullTerminate(VarString *vstr): Null-terminates the VarString string vstr. This is used in conjunction with XSB_StrAppendBlk, because the latter does not null-terminate variable-length strings.

- XSB_StrEnsureSize(VarString *vstr, int minsize): Ensure that the string has room for at least minsize bytes. This is a low-level routine, which is used to interface to procedures that do not use VarString internally. If the string is larger than minsize, the size might actually shrink to the nearest increment that is larger minsize.

- XSB_StrShrink(VarString *vstr, int increment): Shrink the size of vstr to the minimum necessary to hold the data. increment becomes the new increment by which vstr is adjusted. Since VarString is automatically shrunk by XSB_StrSet, it is rarely necessary to shrink a VarString explicitly. However, one might want to change the adjustment increment using this macro (the default increment is 128).

- XSB_StrDestroy(VarString *vstr): Destroys a VarString. Explicit destruction is necessary for VarString's with the automatic storage class. Otherwise, memory leak is possible.

3.5 Passing Data into an XSB Module

The previous chapter described the low-level XSB/C interface that supports passing the data of arbitrary complexity between XSB and C. However, in cases when data needs to be passed into an executable XSB module by the main C program, the following higher-level interface should suffice. (This interface is actually implemented using macros that call the lower level functions.) These routines can be used to construct commands and queries into XSB 's register 1, which is necessary before calling xsb_query() or xsb_command().

void xsb_make_vars((int) N)
    xsb_make_vars creates a return structure of arity N in Register 2. So this routine may called before calling any of xsb_query, xsb_query_string, xsb_command, or xsb_command_string if parameters are to be set to be sent to the goal. It must be called before calling one of the xsb_set_var_* routines can be called. N must be the number of variables in the query that is to be evaluated.

void xsb_set_var_int((int) Val, (int) N)
    set_and_int sets the N^{th} field in the return structure to the integer value Val. It is used to set the value of the N^{th} variable in a query before calling xsb_query or xsb_query_string. When called in XSB, the query will have the N^{th} variable set to this value.

void xsb_set_var_string((char *) Val, (int) N)
    set_and_string sets the N^{th} field in the return structure to the atom with name Val. It is used to set the value of the N^{th} variable in a query before calling xsb_query or xsb_query_string. When called in XSB, the query will have the N^{th} variable set to this value.

```
void xsb_set_var_float((float) Val, (int) N)
```
> set_and_float sets the N^{th} field in the return structure to the floating point number with value Val. It is used to set the value of the N^{th} variable in a query before calling xsb_query or xsb_query_string. When called in XSB, the query will have the N^{th} variable set to this value.

```
prolog_int xsb_var_int((int) N)
```
> xsb_var_int is called after xsb_query or xsb_query_string returns an answer. It returns the value of the N^{th} variable in the query as set in the returned answer. This variable must have an integer value (which is cast to long in a 64-bit architecture).

```
char* xsb_var_string((int) N)
```
> xsb_var_string is called after xsb_query or xsb_query_string returns an answer. It returns the value of the N^{th} variable in the query as set in the returned answer. This variable must have an atom value.

```
prolog_float xsb_var_float((int) N)
```
> xsb_var_float is called after xsb_query or xsb_query_string returns an answer. It returns the value of the N^{th} variable in the query as set in the returned answer. This variable must have a floating point value (which is cast to double in a 64-bit architecture).

3.6 Creating an XSB Module that Can be Called from C

To create an executable that includes calls to the above C functions, these routines, and the XSB routines that they call, must be included in the link (ld) step.

Unix instructions: You must link your C program, which should include the main procedure, with the XSB object file located in

```
$XSBDIR/config/<your-system-architecture>/saved.o/xsb.o
```

Your program should include the file cinterf.h located in the XSB/emu subdirectory, which defines the routines described earlier, which you will need to use in order to talk to XSB. It is therefore recommended to compile your program with the option -I$XSB_DIR/XSB/emu.

The file $XSB_DIR/config/your-system-architecture/modMakefile is a makefile you can use to build your programs and link them with XSB. It is generated automatically and contains all the right settings for your architecture, but you will have to fill in the name of your program, etc.

It is also possible to compile and link your program with XSB using XSB itself as follows:

```
:- xsb_configuration(compiler_flags,CFLAGS),
      xsb_configuration(loader_flags,LDFLAGS),
      xsb_configuration(config_dir,CONFDIR),
      xsb_configuration(emudir,EMUDIR),
      xsb_configuration(compiler,Compiler),
```

```
        str_cat(CONFDIR, '/saved.o/', ObjDir),
        write('Compiling myprog.c ... '),
        shell([Compiler, ' -I', EMUDIR, ' -c ', CFLAGS, ' myprog.c ']),
        shell([Compiler, CFLAGE, ' -o ', './myprog ',
               ObjDir, 'xsb.o ', ' myprog.o ', LDFLAGS]),
        writeln(done).
```

This works for every architecture and is often more convenient than using the make files [6]. There are simple examples of C programs calling XSB in the $XSB_DIR/examples/c_calling_XSB directory, in files cmain.c, ctest.P, cmain2.c.

Windows instructions: To call XSB from C, you must build it as a DLL, which is done as follows:

```
  cd $XSB_DIR\XSB\build
  makexsb_wind DLL="yes"
```

The DLL, which you can call dynamically from your program is then found in

$$\text{\$XSB\_DIR}\backslash\text{config}\backslash\text{x86-pc-windows}\backslash\text{bin}\backslash\text{xsb.dll}$$

Since your program must include the file cinterf.h, it is recommended to compile it with the option /I$XSB_DIR\XSB\emu.

---

[6]The variable CFLAGS is needed in the linking stage in order to ensure that the appropriate memory option is passed if XSB is configured --with-bits32 or --with-bits64 to override the default on a 64-bit platform.

Chapter 4

XSB-ODBC Interface

By Baoqiu Cui, Lily Dong, and David S. Warren [1].

4.1 Introduction

The XSB-ODBC interface is subsystem that allows XSB users to access databases through ODBC connections. This is mostly of interest to Microsoft Windows users. The interface allows XSB users to access data in any ODBC compliant database management system (DBMS). Using this uniform interface, information in different DBMS's can be accessed as though it existed as Prolog facts. The XSB-ODBC interface provides users with three levels of interaction: an *SQL level*, a *relation level* and a *view level*. The *SQL level* allows users to write explicit SQL statements to be passed to the interface to retrieve data from a connected database. The *relation level* allows users to declare XSB predicates that connect to individual tables in a connected database, and which when executed support tuple-at-a-time retrieval from the base table. The *view level* allows users to use a complex XSB query, including conjunction, negation and aggregates, to specify a database query. A listing of the features that the XSB-ODBC interface provides is as follows:

- Concurrent access from multiple XSB processes to a single DBMS

- Access from a single XSB process to multiple ODBC DBMS's

- Full data access and cursor transparency including support for

    - Full data recursion through XSB's tabling mechanism (depending on the capabilities of the underlying ODBC driver.

    - Runtime type checking

    - Automatic handling of NULL values for insertion, deletion and querying

- Full access to data source including

---

[1]This interface was partly based on the XSB-Oracle Interface by Hassan Davulcu, Ernie Johnson and Terrance Swift.

- – Transaction support

- – Cursor reuse for cached SQL statements with bind variables (thereby avoiding re-parsing and re-optimizing).

- – Caching compiler generated SQL statements with bind variables and efficient cursor management for cached statements

- A powerful Prolog / SQL compiler based on [9].

- Full source code availability

- Independence from database schema by the *relation level* interface

- Performance as SQL by employing a *view level*

- No mode specification is required for optimized view compilation

We use the `Hospital` database as our example to illustrate the usage of XSB-ODBC interface in this manual. We assume the basic knowledge of Microsoft ODBC interface and its ODBC administrator throughout the text. Please refer to "Inside Windows^{TM} 95" (or more recent documentation) for information on this topic.

4.2   Using the Interface

The XSB-ODBC module is a module and as such exports the predicates it supports. In order to use any predicate defined below, **it must be imported** from `odbc_call`. For example, before you can use the predicate to open a data source, you must include:

```
:- import odbc_open/3 from odbc_call.
```

4.2.1   Connecting to and Disconnecting from Data Sources

Assuming that the data source to be connected to is available, i.e. it has an entry in `ODBC.INI` file which can be checked by running Microsoft ODBC Administrator, it can be connected to in the following way:

```
| ?- odbc_open(data_source_name, username, passwd).
```

If the connection is successfully made, the predicate invocation will succeed. This step is necessary before anything can be done with the data sources since it gives XSB the opportunity to initialize system resources for the session.

This is an executable predicate, but you may want to put it as a query in a file that declares a database interface and will be loaded.

To close the current session use:

```
| ?- odbc_close.
```

and XSB will give all the resources it allocated for this session back to the system.

If you are connecting to only one data source at a time, the predicates above are sufficient. However, if you want to connect to multiple data sources at the same time, you must use extended versions of the predicates above. When connecting to multiple sources, you must give an atomic name to each source you want to connect to, and use that name whenever referring to that source. The names may be chosen arbitrarily but must be used consistently. The extended versions are:

```
| ?- odbc_open(data_source_name, username, passwd, connectionName).
```

and

```
| ?- odbc_close(connectionName).
```

A list of existing Data Source Names and descriptions can be obtained by backtracking through odbc_data_sources/2. For example:

```
| ?- odbc_data_sources(DSN,DSNDescr).

DSN = mycdf
DSNDescr = MySQL driver;

DSN = mywincdf
DSNDescr = TDS driver (Sybase/MS SQL);
```

4.2.2   Accessing Tables in Data Sources Using SQL

There are several ways that can be used to extract information from or modify a table in a data source. The most basic way is to use predicates that pass an SQL statement directly to the ODBC driver. The basic call is:

```
| ?- odbc_sql(BindVals,SQLStmt,ResultRow).
```

where `BindVals` is a list of (ground) values that correspond to the parameter indicators in the SQL statement (the '?'s); `SQLStmt` is an atom containing an SQL statement; and `ResultRow` is a returned list of values constituting a row from the result set returned by the SQL query. Thus for a select SQL statement, this call is nondeterministic, returning each retrieved row in turn.

The `BindVals` list should have a length corresponding to the number of parameters in the query, in particular being the empty list ([]) if `SQLStmt` contains no '?'s. If `SQLStmt` is not a select statement returning a result set, then `ResultRow` will be the empty list, and the call is deterministic. Thus this predicate can be used to do updates, DDL statements, indeed any SQL statement.

SQLStmt need not be an atom, but can be a (nested) list of atoms which flattens (and concatenates) to form an SQL statement.

BindVals is normally a list of values of primitive Prolog types: atoms, integers, or floats. The values are converted to the types of the corresponding database fields. However, complex Prolog values can also be stored in a database field. If a term of the form term(VAL) appears in the BindVal list, then VAL (a Prolog term) will be written in canonical form (as produced by write_canonical) to the corresponding database field (which must be CHAR or BYTE). If a term of the form string(CODELIST) appears in BindVal, then CODELIST must be a list of ascii-codes (as produced by atom_codes) and these codes will be converted to a CHAR or BYTE database type.

ResultRow for a select statement is normally a list of variables that will nondeterministically be bound to the values of the fields of the tuples returned by the execution of the select statement. The Prolog types of the values returned will be determined by the database types of the corresponding fields. A CHAR or BYTE database type will be returned as a Prolog atom; an INTEGER database field will be returned as a Prolog integer, and similarly for floats. However, the user can request that CHAR and BYTE database fields be returned as something other than an atom. If the term string(VAR) appears in ResultRow, then the corresponding database field must be CHAR or BYTE, and in this case, the variable VAR will be bound to the list of ascii-codes that make up the database field. This allows an XSB programmer to avoid adding an atom to the atom table unnecessarily. If the term term(VAR) appears in ResultRow, then the corresponding database field value is assumed to be a Prolog term in canonical form, i.e., can be read by read_canonical/1. The corresponding value will be converted into a Prolog term and bound to VAR. This allows a programmer to store complex Prolog terms in a database. Variables in such a term are local only to that term.

When connecting to multiple data sources, you should use the form:

```
| ?- odbc_sql(ConnectionName,BindVals,SQLStmt,ResultRow).
```

For example, we can define a predicate, get_test_name_price, which given a test ID, retrieves the name and price of that test from the test table in the hospital database:

```
get_test_name_price(Id,Nam,Pri) :-
      odbc_sql([Id],'SELECT TName,Price FROM Test WHERE TId = ?', [Nam,Pri]).
```

The interface uses a cursor to retrieve this result and caches the cursor, so that if the same query is needed in the future, it does not need to be re-parsed, and re-optimized. Thus, if this predicate were to be called several times, the above form is more efficient than the following form, which must be parsed and optimized for each and every call:

```
get_test_name_price(Id,Nam,Pri) :-
    odbc_sql([],['SELECT TName,Price FROM Test WHERE TId = ''',Id,''''], [Nam,Pri]).
```

Note that to include a quote (') in an atom, it must be represented by using two quotes.

There is also a predicate:

```
| ?- odbc_sql_cnt(ConnectionName,BindVals,SQLStmt,Count).
```

This predicate is very similar to odbc_slq/4 except that it can only be used for UPDATE, INSERT, and DELETE SQL statements. The first three arguments are just as in odbc_slq/4; the fourth must be a variable in which is returned the integer count of the number of rows affected by the SQL operation.

4.2.3   Cursor Management

The XSB-ODBC interface is limited to using 100 open cursors. When XSB systems use database accesses in a complicated manner, management of open cursors can be a problem due to the tuple-at-a-time access of databases from Prolog, and due to leakage of cursors through cuts and throws. Often, it is more efficient to call the database through set-at-a-time predicates such as findall/3, and then to backtrack through the returned information. For instance, the predicate findall_odbc_sql/4 can be defined as:

```
findall_odbc_sql(ConnName,BindVals,SQLStmt,ResultRow):-
            findall(Res,odbc_sql(ConnName,BindVals,SQLStmt,Res),Results),
            member(ResultRow,Results).
```

As a convenience, therefore, the predicates findall_odbc_sql/3 and findall_odbc_sql/4 are defined in the ODBC interface.

4.2.4   Accessing Tables in Data Sources through the Relation Level

While all access to a database is possible using SQL as described above, the XSB-ODBC interface supports higher-level interaction for which the user need not know or write SQL statements; that is done as necessary by the interface. With the relation level interface, users can simply declare a predicate to access a table and the system generates the necessary underlying code, generating specialized code for each mode in which the predicate is called.

To declare a predicate to access a database table, a user must use the odbc_import/2 interface predicate.

The syntax of odbc_import/2 is as follows:

```
| ?- odbc_import('TableName'('FIELD1', 'FIELD2', ..., 'FIELDn'), 'PredicateName').
```

where 'TableName' is the name of the database table to be accessed and 'PredicateName' is the name of the XSB predicate through which access will be made. 'FIELD1', 'FIELD2', ... , 'FIELDn' are the exact attribute names(case sensitive) as defined in the database table schema. The chosen columns define the view and the order of arguments for the database predicate 'PredicateName'.

For example, to create a link to the Test table through the 'test' predicate:

```
| ?- odbc_import('Test'('TId','TName','Length','Price'),test).
```

```
yes
```

When connecting to multiple data sources, you should use the form:

```
| ?- odbc_import(ConnectionName,
                 'TableName'('FIELD1', 'FIELD2', ..., 'FIELDn'),
                 'PredicateName').
```

4.2.5 Using the Relation Level Interface

Once the links between tables and predicates have been successfully established, information can then be extracted from these tables using the corresponding predicates. Continuing from the above example, now rows from the table Test can be obtained:

```
| ?- test(TId, TName, L, P).
```

```
TId = t001
TName = X-Ray
L = 5
P = 100
```

Backtracking can then be used to retrieve the next row of the table Test.

Records with particular field values may be selected in the same way as in Prolog; no mode specification for database predicates is required. For example:

```
| ?- test(TId, 'X-Ray', L, P).
```

will automatically generate the query:

```
SELECT rel1.TId, rel1.TName, rel1.Length, rel1.Price
FROM Test rel1
WHERE rel1.TName = ?
```

and

```
| ?- test('NULL'(_), 'X-Ray',  L, P).
```

generates: (See Section 4.2.6)

```
SELECT NULL , rel1.TName, rel1.Length, rel1.Price
FROM Test rel1
WHERE rel1.TId IS NULL AND rel1.TName = ?
```

During the execution of this query the bind variable `?` will be bound to the value `'X-Ray'`.

Of course, the same considerations about cursors noted in Section 4.2.3 apply to the relation-level interface. Accordingly, the ODBC interface also defines the predicate `odbc_import/4` which allows the user to specify that rows are to be fetched through `findall/3`. For example, the call

```
odbc_import('Test'('TId','TName','Length','Price'),test,[findall(true)]).
```

will behave as described above *but* will make all database calls through `findall/3` and return rows by backtracking through a list rather than maintaining open cursors.

Also as a courtesy to Quintus Prolog users we have provided compatibility support for some PRODBI predicates which access tables at a relational level [2].

```
| ?- odbc_attach(PredicateName, table(TableName)).
```

eg. invoke

```
| ?- odbc_attach(test2, table('Test')).
```

and then execute

```
| ?- test2(TId, TName, L, P).
```

to retrieve the rows.

4.2.6   Handling NULL values

The interface treats NULL's by introducing a single valued function `'NULL'/1` whose single value is a unique (Skolem) constant. For example a NULL value may be represented by

```
'NULL'(null123245)
```

Under this representation, two distinct NULL values will not unify. On the other hand, the search condition `IS NULL Field` can be represented in XSB as `Field = 'NULL'(_)`

Using this representation of NULL's the following protocol for queries and updates is established.

Queries

$$| \ ?- \ \mathtt{dept('NULL'(\_),\_,\_).}$$

Generates the query:

---

[2]This predicate is obsolescent and `odbc_import/{2,3,4}` should be used instead.

```
SELECT NULL , rel1.DNAME , rel1.LOC
  FROM DEPT rel1
  WHERE rel1.DEPTNO IS NULL;
```

Hence, 'NULL'(_) can be used to retrieve rows with NULL values at any field.

'NULL'/1 fails the predicate whenever it is used with a bound argument.

> | ?- dept('NULL'(null2745),_,_).   → fails always.

Query Results

When returning NULL's as field values, the interface returns NULL/1 function with a unique integer argument serving as a skolem constant.

Notice that the above guarantees the expected semantics for the join statements. In the following example, even if Deptno is NULL for some rows in emp or dept tables, the query still evaluates the join successfully.

> | ?- emp(Ename,_,_,_,Deptno),dept(Deptno,Dname,Loc)..

Inserts

To insert rows with NULL values you can use Field = 'NULL'(_) or Field = 'NULL'(null2346). For example:

> | ?- emp_ins('NULL'(_), ...).   → inserts a NULL value for ENAME

> | ?- emp_ins('NULL'('bound'), ...)   → inserts a NULL value for ENAME.

Deletes

To delete rows with NULL values at any particular FIELD use Field = 'NULL'(_), 'NULL'/1 with a free argument. When 'NULL'/1 's argument is bound it fails the delete predicate always. For example:

| ?- emp_del('NULL'(_), ..).   → adds ENAME IS NULL to the generated SQL statement

> | ?- emp_del('NULL'('bound'), ...).   → fails always

The reason for the above protocol is to preserve the semantics of deletes, when some free arguments of a delete predicate get bound by some preceding predicates. For example in the following clause, the semantics is preserved even if the Deptno field is NULL for some rows.

> | ?- emp(_,_,_,_,Deptno), dept_del(Deptno).

4.2.7 The View Level Interface

The view level interface can be used to define XSB queries which include only imported database predicates (by using the relation level interface) described above and aggregate predicates (defined below). When these queries are invoked, they are translated into complex database queries, which are then executed taking advantage of the query processing ability of the DBMS.

One can use the view level interface through the predicate `odbc_query/2`:

```
| ?- odbc_query('QueryName'(ARG1, ..., ARGn), DatabaseGoal).
```

All arguments are standard XSB terms. `ARG1`, `ARG2`, ..., `ARGn` define the attributes to be retrieved from the database, while `DatabaseGoal` is an XSB goal (i.e. a possible body of a rule) that defines the selection restrictions and join conditions.

The compiler is a simple extension of [9] which generates SQL queries with bind variables and handles NULL values as described in Section 4.2.6. It allows negation, the expression of arithmetic functions, and higher-order constructs such as grouping, sorting, and aggregate functions.

Database goals are translated according to the following rules from [9]:

- Disjunctive goals translate to distinct SQL queries connected through the UNION operator.

- Goal conjunctions translate to joins.

- Negated goals translate to negated EXISTS subqueries.

- Variables with single occurrences in the body are not translated.

- Free variables translate to grouping attributes.

- Shared variables in goals translate to equi-join conditions.

- Constants translate to equality comparisons of an attribute and the constant value.

- Nulls are translated to `IS NULL` conditions.

For more examples and implementation details see [9].

In the following, we show the definition of a simple join view between the two database predicates *Room* and *Floor*.

Assuming the declarations:

```
| ?- odbc_import('Room'('RoomNo','CostPerDay','Capacity','FId'),room).

| ?- odbc_import('Floor'('FId','','FName'),floor).
```

use

```
| ?- odbc_query(query1(RoomNo,FName),
                  (room(RoomNo,_,_,FId),floor(FId,_,FName))).
yes
```

```
| ?- query1(RoomNo,FloorName).
```

Prolog/SQL compiler generates the SQL statement:

```
SELECT rel1.RoomNo , rel2.FName FROM Room rel1 , Floor rel2
WHERE rel2.FId = rel1.FId;
```

Backtracking can then be used to retrieve the next row of the view.

```
| ?- query1('101','NULL'(_)).
```

generates the SQL statement:

```
SELECT rel1.RoomNo, NULL
FROM Room rel1 , Floor rel2
WHERE rel1.RoomId = ? AND rel2.FId = rel1.FId AND rel2.FName IS NULL;
```

The view interface also supports aggregate functions such as sum, avg, count, min and max. For example

```
| ?- odbc_import('Doctor'('DId', 'FId', 'DName','PhoneNo','ChargePerMin'),doctor).

yes
| ?- odbc_query(avgchargepermin(X),
                (X is avg(ChargePerMin, A1 ^ A2 ^ A3 ^ A4 ^
                          doctor(A1,A2, A3,A4,ChargePerMin)))).

yes
| ?- avgchargepermin(X).

SELECT AVG(rel1.ChargePerMin)
FROM doctor rel1;

X = 1.64

yes
```

A more complicated example is the following:

```
| ?- odbc_query(nonsense(A,B,C,D,E),
                (doctor(A, B, C, D, E),
                 not floor('First Floor', B),
                 not (A = 'd001'),
                 E > avg(ChargePerMin, A1 ^ A2 ^ A3 ^ A4 ^
                         (doctor(A1, A2, A3, A4, ChargePerMin))))).

| ?- nonsense(A,'4',C,D,E).

SELECT rel1.DId , rel1.FId , rel1.DName , rel1.PhoneNo , rel1.ChargePerMin
FROM doctor rel1
WHERE rel1.FId = ? AND NOT EXISTS
(SELECT *
FROM Floor rel2
WHERE rel2.FName = 'First Floor' and rel2.FId = rel1.FId
) AND rel1.Did <> 'd001' AND rel1.ChargePerMin >
(SELECT AVG(rel3.ChargePerMin)
FROM Doctor rel3
);


A = d004
C = Tom Wilson
D = 516-252-100
E = 2.5
```

All database queries defined by odbc_query/{2,3} can be queried with any mode.

Note that at each call to a database relation or rule, the communication takes place through bind variables. The corresponding restrictive SQL query is generated, and if this is the first call with that adornment, it is cached. A second call with same adornment would try to use the same database cursor if still available, without reparsing the respective SQL statement. Otherwise, it would find an unused cursor and retrieve the results. In this way efficient access methods for relations and database rules can be maintained throughout the session.

If connecting to multiple data sources, use the form:

```
 :- odbc_query(connectionName,'QueryName'(ARG1, ..., ARGn), DatabaseGoal).
```

4.2.8   Insertions and Deletions of Rows through the Relational Level

Insertion and deletion operations can also be performed on an imported table. The two predicates to accomplish these operations are odbc_insert/2 and odbc_delete/2. The syntax of odbc_insert/2 is as follows: the first argument is the declared database predicate for insertions and the second

argument is some imported data source relation. The second argument can be declared with some of its arguments bound to constants. For example after `Room` is imported through `odbc_import`:

```
|?- odbc_import('Room'('RoomNo','CostPerDay','Capacity','FId'), room).
yes
```

Now we can do

```
| ?- odbc_insert(room_ins(A1,A2,A3),(room(A1,A2,A3,'3'))).

yes
| ?- room_ins('306','NULL'(_),2).

yes
```

This will insert the row: ('306',NULL, 2,'3') into the table `Room`. Note that any call to `room_ins/7` should have all its arguments bound.

See Section 4.2.6) for information about NULL value handling.

The first argument of `odbc_delete/2` predicate is the declared delete predicate and the second argument is the imported data source relation with the condition for requested deletes, if any. The condition is limited to simple comparisons. For example assuming `Room/3` has been imported as above:

```
| ?- odbc_delete(room_del(A), (room('306',A,B,C), A > 2)).

yes
```

After this declaration you can use:

```
| ?- room_del(3).
```

to generate the SQL statement:

```
DELETE From Room rel1
WHERE rel1.RoomNo = '306' AND rel1.CostPerDay = ? AND ? > 2
;
```

Note that you have to commit your inserts or deletes to tables to make them permanent. (See section 4.2.11).

These predicates also have the form in which an additional first argument indicates a connection, for use with multiple data sources.

Also, some ODBC drivers have been found that do not accept the form of SQL generated for deletes. In these cases, you must use the lower-level interface: `odbc_sql`.

4.2.9   Access to Data Dictionaries

The following utility predicates provide users with tools to access data dictionaries [3]. A brief description of these predicates is as follows:

odbc_show_schema(accessible(Owner)) Shows the names of all accessible tables that are owned by Owner. (This list can be long!) If Owner is a variable, all tables will be shown, grouped by owner.

odbc_show_schema(user) Shows just those tables that belongs to user.

odbc_show_schema(tuples('Table')) Shows all rows of the database table named `'Table'`.

odbc_show_schema(arity('Table')) The number of fields in the table `'Table'`.

odbc_show_schema(columns('Table')) The field names of a table.

For retrieving above information use:

- `odbc_get_schema(accessible(Owner),List)`

- `odbc_get_schema(user,List)`

- `odbc_get_schema(arity('Table'),List)`

- `odbc_get_schema(columns('Table'),List)`

The results of above are returned in List as a list.

4.2.10   Other Database Operations

odbc_create_table('TableName','FIELDs') `FIELDS` is the field specification as in SQL.

```
eg. odbc_create_table('MyTable', 'Col1 NUMBER,
                        Col2 TEXT(50),
                        Col3 TEXT(13)').
```

odbc_create_index('TableName','IndexName', index(_,Fields)) `Fields` is the list of columns for which an index is requested. For example:

---

[3]Users of Quintus Prolog may note that these predicates are all PRODBI compatible.

```
    odbc_create_index('Doctor', 'DocKey', index(_,'DId')).
```

odbc_delete_table('TableName') To delete a table named 'TableName'

odbc_delete_view('ViewName') To delete a view named 'ViewName'

odbc_delete_index('IndexName') To delete an index named 'IndexName'

4.2.11   Transaction Management

Depending on how the transaction options are set in ODBC.INI for data sources, changes to the data source tables may not be committed (i.e., the changes become permanent) until the user explicitly issues a commit statement. Some ODBC drivers support autocommit, which, if on, means that every update operation is immediately committed upon execution. If autocommit is off, then an explicit commit (or rollback) must be done by the program to ensure the updates become permanent (or are ignored.).

The predicate `odbc_transaction/1` supports these operations.

odbc_transaction(autocommit(on)) Turns on autocommit, so that all update operations will be immediately committed on completion.

odbc_transaction(autocommit(off)) Turns off autocommit, so that all update operations will not be committed until explicitly done so by the program (using one of the following operations.)

odbc_transaction(commit) Commits all transactions up to this point. (Only has an effect if autocommit is off).

odbc_transaction(rollback) Rolls back all update operations done since the last commit point. (Only has an effect if autocommit is off).

4.2.12   Interface Flags

Users are given the option to monitor control aspects of the ODBC interface by setting ODBC flags via the predicates **set_odbc_flag/2** and **odbc_flag/2**.

The first aspect that can be controlled is whether to display SQL statements for SQL queries. This is done by the `show_query` flag. For example:

```
| ?- odbc_flag(show_query,Val).

    Val = on
```

Indicates that SQL statements will now be displayed for all SQL queries, and is the default value for the ODBC interface. To turn it off execute the command `set_odbc_flag(show_query,on)`.

The second aspect that can be controlled is the action taken upon ODBC errors. Three possible actions may be useful in different contexts and with different drivers. First, the error may be ignored, so that a database call succeeds; second the error cause the predicate to fail, and third the error may cause an exception to be thrown to be handled by a catcher (or the default system error handler, see Volume 1).

| ?- `odbc_flag(fail_on_error, ignore)` Ignores all ODBC errors, apart from writing a warning. In this case, it's the users' users' responsibility to check each of their actions and do error handling.

| ?- `odbc_flag(fail_on_error, fail)` Interface fails whenever error occurs.

| ?- `odbc_flag(fail_on_error, throw)` Throws an error-term of the form `error(odbc_error,Message,Backtr` in which `Message` is a textual description of the ODBC error, and `Backtrace` is a list of the continuations of the call. These continuations may be printed out by the error handler.

The default value of fail_on_error is on.

4.2.13 Datalog

Users can write recursive Datalog queries with exactly the same semantics as in XSB using imported database predicates or database rules. For example assuming `odbc_parent/2` is an imported database predicate, the following recursive query computes its transitive closure.

```
:- table(ancestor/2).
ancestor(X,Y) :- odbc_parent(X,Y).
ancestor(X,Z) :- ancestor(X,Y), odbc_parent(Y,Z).
```

This works with drivers that support multiple open cursors to the same connection at the same time. (Sadly, some don't.) In the case of drivers that don't support multiple open cursors, one can often replace each odbc_import-ed predicate call

```
...,predForTable(A,B,C),...
```

by

```
...,findall([A,B,C],predForTable(A,B,C),PredList),
    member([A,B,C],PredList)...
```

and get the desired effect.

4.3   Error messages

ERR - DB: Connection failed For some reason the attempt to connect to data source failed.

- Diagnosis: Try to see if the data source has been registered with Microsoft ODBC Administrator, the username and password are correct and MAXCURSORNUM is not set to a very large number.

ERR - DB: Parse error The SQL statement generated by the Interface or the first argument to `odbc_sql/1` or `odbc_sql_select/2` can not be parsed by the data source driver.

- Diagnosis: Check the SQL statement. If our interface generated the erroneous statement please contact us at `xsb-contact@cs.sunysb.edu`.

ERR - DB: No more cursors left Interface run out of non-active cursors either because of a leak or no more free cursors left.

- Diagnosis: System fails always with this error. odbc_transaction(rollback) or odbc_transaction(commit) should resolve this by freeing all cursors.

ERR - DB: FETCH failed Normally this error should not occur if the interface running properly.

- Diagnosis: Please contact us at xsb-contact@cs.sunysb.edu

4.4   Notes on specific ODBC drivers

MyODBC The ODBC driver for MySQL is called MyODBC, and it presents some particularities that should be noted.

First, MySQL, as of version 3.23.55, does not support strings of length greater than 255 characters. XSB's ODBC interface has been updated to allow the use of the BLOB datatype to encode larger strings.

More importantly, MyODBC implements SQLDescribeCol such that, by default, it returns actual lengths of columns in the result table, instead of the formal lengths in the tables. For example, suppose you have, in table A, a field f declared as "VARCHAR (200)". Now, you create a query of the form "SELECT f FROM A WHERE ..." If, in the result set, the largest size of f is 52, that's the length that SQLDescribeCol will return. This breaks XSB's caching of query-related data-structures. In order to prevent this behavior, you should configure your DSN setup so that you pass "Option=1" to MyODBC.

Chapter 5

The New XSB-Database Interface

By Saikat Mukherjee, Michael Kifer and Hui Wan

5.1 Introduction

The XSB-DB interface is a package that allows XSB users to access databases through various drivers. Using this interface, information in different DBMSs can be accessed by SQL queries. The interface defines Prolog predicates which makes it easy to connect to databases, query them, and disconnect from the databases. Central to the concept of a connection to a database is the notion of a *handle*. A connection handle describes a particular connection to a database. Similar to a connection handle is the notion of a query handle which describes a particular query statement. As a consequence of the handles, it is possible to open multiple database connections (to the same or different databases) and keep alive multiple queries (again from the same or different connections). The interface also supports dynamic loading of drivers. As a result, it is possible to query databases using different drivers concurrently [1].

Currently, this package provides drivers for ODBC, a native MySQL driver, and a driver for the embedded MySQL server.

5.2 Configuring the Interface

Generally, each driver has to be configured separately, but if the database packages such as ODBC, MySql, etc., are installed in standard places then the XSB configuration mechanism will do the job automatically.

Under Windows, first make sure that XSB is configured and built correctly for Windows, and that it runs. As part of that building process, the command

```
makexsb_wind
```

---

[1]In Version 3.2, this package has not been ported to the multi-threaded engine.

must have been executed in the directory `XSB\build`. It will normally configure the ODBC driver without problems. For the MySQL driver one has to edit the file

    packages\dbdrivers\mysql\cc\NMakefile.mak

to indicate where MySQL is installed. To build the embedded MySQL driver under Windows, the file

    packages\dbdrivers\mysqlenbedded\cc\NMakefile.mak

might need to be edited. Then you should either rebuild XSB using the `makexsb_wind` command or by running

    nmake /f NMakefile.mak

in the appropriate directories (`dbdrivers\mysql\cc` or `dbdrivers\mysqlenbedded\cc`). Note that you need a C++ compiler and `nmake` installed on your system for this to work.[2]

Under Unix, the `configure` script will build the drivers automatically if the `--with-dbdrivers` option is specified. If, however, ODBC and MySQL are not installed in their standard places, you will have to provide the following parameters to the configure script:

- `--with-odbc-libdir=LibDIR` – `LibDIR` is the directory where the library libodbc.so lives on your system.

- `--with-odbc-incdir=IncludeDIR` – `IncludeDIR` is the directory where the ODBC header files, such as `sql.h` live.

- `--with-mysql-libdir=MySQLlibdir` – `MySQLlibdir` is the directory where MySQL's shared libraries live on your system.

- `--with-mysql-incdir=MySQLincludeDir` – `MySQLincludeDir` is the directory where MySQL's header files live.

If you are also using the embedded MySQL server and want to take advantage of the corresponding XSB driver, you need to provide the following directories to tell XSB where the copy of MySQL that supports the embedded server is installed. This has to be done *only* if that copy is not in a standard place, like `/usr/lib/mysql`.

- `--with-mysqlembedded-libdir=MySQLlibdir` – `MySQLlibdir` is the directory where MySQL's shared libraries live on your system. This copy of MySQL must be configured with support for the embedded server.

- `--with-mysqlembedded-incdir=MySQLincludeDir` – `MySQLincludeDir` is the directory where MySQL's header files live.

---

[2] http://www.microsoft.com/express/vc/
http://download.microsoft.com/download/vc15/Patch/1.52/W95/EN-US/Nmake15.exe

Under Cygwin, the ODBC libraries come with the distribution; they are located in the directory `/cygdrive/c/cygwin/lib/w32api/` and are called `odbc32.a` and `odbccp32.a`. (Check if your installation is complete and has these libraries!) Otherwise, the configuration of the interface under Cygwin is same as in unix (you do not need to provide any ODBC-specific parameters to the configure script under Cygwin).

If at the time of configuring XSB some database packages (*e.g.*, MySQL) are not installed on your system, you can install them later and configure the XSB interface to them then. For instance, to configure the ODBC interface separately, you can type

```
cd packages/dbdrivers/odbc
configure
```

Again, if ODBC is installed in a non-standard location, you might need to supply the options `--with-odbc-libdir` and `--with-odbc-incdir` to the configure script. Under Cygwin ODBC is always installed in a standard place, and `configure` needs no additional parameters.

Under Windows, separate configuration of the XSB-DB interfaces is also possible, but you need Visual Studio installed. For instance, to configure the MySQL interface, type

```
cd packages\dbdrivers\mysql\cc
nmake /f NMakefile.mak
```

As before, you might need to edit the NMakefile.mak script to tell the compiler where the required MySQL's libraries are. You also need the file `packages\dbdrivers\mysql\mysql_init.P` with the following content:

```
:- export mysql_info/2.
mysql_info(support, 'yes').
mysql_info(libdir,  '').
mysql_info(ccflags, '').
mysql_info(ldflags, '').
```

Similarly, to configure the ODBC interface, do

```
cd packages\dbdrivers\odbc\cc
nmake /f NMakefile.mak
```

You will also need to create the file `packages\dbdrivers\odbc\odbc_init.P` with the following contents:

```
:- export odbc_info/2.
odbc_info(support, 'yes').
odbc_info(libdir,  '').
odbc_info(ccflags, '').
odbc_info(ldflags, '').
```

5.3   Using the Interface

We use the `student` database as our example to illustrate the usage of the XSB-DB interface in this manual. The schema of the student database contains three columns viz. the student name, the student id, and the name of the advisor of the student.

The XSB-DB package has to be first loaded before using any of the predicates. This is done by the call:

```
| ?- [dbdrivers].
```

Next, the driver to be used for connecting to the database has to be loaded. Currently, the interface has support for a native MySQL driver (using the MySQL C API), and an ODBC driver. For example, to load the ODBC driver call:

```
| ?- load_driver(odbc).
```

Similarly, to load the mysql driver call:

```
| ?- load_driver(mysql).
```

or

```
| ?- load_driver(mysqlembedded).
```

5.3.1   Connecting to and Disconnecting from Databases

There are two predicates for connecting to databases, `db_connect/5` and `db_connect/6`. The `db_connect/5` predicate is for ODBC connections, while `db_connect/6` is for other (non-ODBC) database drivers.

```
| ?- db_connect(+Handle, +Driver, +DSN, +User, +Password).
| ?- db_connect(+Handle, +Driver, +Server, +Database, +User, +Password).
```

The `db_connect/5` predicate assumes that an entry for a data source name (DSN) exists in the `odbc.ini` file. The `Handle` is the handle name used for the connection. The `Driver` is the driver being used for the connection. The `User` and `Password` are the user name and password being used for the connection. The user is responsible for giving the name to the handle. To connect to the data source mydb using the user name xsb and password xsb with the odbc driver, the call is as follows:

```
| ?- db_connect(ha, odbc, mydb, xsb, xsb).
```

where `ha` is the user-chosen handle name (a Prolog atom) for the connection.

The `db_connect/6` predicate is used for drivers other than ODBC. The arguments `Handle`, `Driver`, `User`, and `Password` are the same as for `db_connect/5`. The `Server` and `Database` arguments specify the server and database to connect to. For example, for a connection to a database called `test` located on the server `wolfe` with the user name `xsb`, the password `foo`, and using the `mysql` driver, the call is:

```
| ?- db_connect(ha, mysql, wolfe, test, xsb, foo).
```

where `ha` is the handle name the user chose for the connection.

If the connection is successfully made, the predicate invocation will succeed. This step is necessary before anything can be done with the data sources since it gives XSB the opportunity to initialize system resources for the session.

To close a database connection use:

```
| ?- db_disconnect(Handle).
```

where handle is the handle name for the connection. For example, to close the connection to above mysql database call:

```
| ?- db_disconnect(ha).
```

and XSB will give all the resources it allocated for this session back to the system.

5.3.2   Querying Databases

The interface supports two types of querying. In direct querying, the query statement is not prepared while in prepared querying the query statement is prepared before being executed. The results from both types of querying are retrieved tuple at a time. Direct querying is done by the predicate:

```
| ?- db_query(ConnectionHandle, QueryHandle, SQLQueryList, ReturnList).
```

ConnectionHandle is the name of the handle used for the database connection. QueryHandle is the name of the handle for this particular query. Currently, the query handle is being used only for prepared queries. However, in future versions, the query handle can be used in direct queries to retrieve arbitrary tuples from a result set using cursors. Also, it will be possible to combine arbitrary tuples from different queries to the same database using the query handle. The SQLQueryList is a list of terms which is used to build the SQL query. The terms in this list are ground atoms. ReturnList is a list of variables each of which correspond to a return value in the query. It is upto the user to specify the correct number of return variables corresponding to the query. Also, as in the case of a connection handle, the user is responsible for giving the name to the query handle. For example, a query on the student database to select all the students for a given advisor is accomplished by the call:

```
| ?- X = adv,
    db_query(ha, qa, ['select T.name from student T where T.advisor = ', X], [P]),
    fail.
```

where `ha` and `qa` are respectively the connection handle and query handle name the user chose.

Observe that the query list is composed of the SQL string and a ground value for the advisor. The return list is made of one variable corresponding to the student name. The failure drive loop retrieves all the tuples.

Preparing a query is done by the call to the predicate:

```
| ?- db_prepare(ConnectionHandle, QueryHandle, SQLQueryList).
```

As before, ConnectionHandle and QueryHandle specify the handles for the connection and the query. The SQLQueryList is a list of terms which build up the query string. The placeholder '?' is used for values which have to be bound during the execution of the statement. For example, to prepare a query for selecting the advisor name for a student name using our student database:

```
| ?- db_prepare(ha, qa, ['select T.advisor from student T where T.name = ?']).
```

A prepared statement is executed using the predicate:

```
| ?- db_prepare_execute(QueryHandle, BindList, ReturnList).
```

The BindList contains the ground values corresponding to the '?' in the prepared statement. The ReturnList is a list of variables for each argument in a tuple of the result set.

For direct querying, the statement handle is closed automatically when all the tuples in the result set have been retrieved. In order to explicitly close a statement handle, and free all the resources associated with the handle, a call is made to the predicate:

```
| ?- db_statement_close(QueryHandle).
```

where QueryHandle is the query handle for the statement to be closed.

The interface is also able to transparently handle Prolog terms. Users can both save and retrieve terms without any special processing.

5.4   Error Handling

Each predicate in the XSB-DB interface throws an exception with the functor

```
xsb_error(database(Number), Message)
```

where Number is a string with the error number and Message is a string with a slightly detailed error message. It is upto the user to catch this exception and proceed with error handling. This is done by the throw-catch error handling mechanism in XSB. For example, in order to catch the error which will be thrown when the user attempts to close a database connection for a handle (`ha`) which does not exist:

```
| ?- catch(db_disconnect(ha),
     xsb_error(database(Number), Message), handler(Number, Message)).
```

It is the user's responsibility to define the handler predicate which can be as simple as printing out the error number and message or may involve more complicated processing.

A list of error numbers and messages that are thrown by the XSB-DB interface is given below:

- **XSB_DBI_001: XSB_DBI ERROR: Driver already registered**
  This error is thrown when the user tries to load a driver, using the `load_driver` predicate, which has already been loaded previously.

- **XSB_DBI_002: XSB_DBI ERROR: Driver does not exist**
  This error is thrown when the user tries to connect to a database, using `db_connect`, with a driver which has not been loaded.

- **XSB_DBI_003: XSB_DBI ERROR: Function does not exist in this driver**
  This error is thrown when the user tries to use a function support for which does not exist in the corresponding driver. For example, this error is generated if the user tries to use `db_prepare` for a connection established with the mysql driver.

- **XSB_DBI_004: XSB_DBI ERROR: No such connection handle**
  This error is thrown when the user tries to use a connection handle which has not been created.

- **XSB_DBI_005: XSB_DBI ERROR: No such query handle**
  This error is thrown when the user tries to use a query handle which has not been created.

- **XSB_DBI_006: XSB_DBI ERROR: Connection handle already exists**
  This error is thrown when the user tries to create a connection handle in `db_connect` using a name which already exists as a connection handle.

- **XSB_DBI_007: XSB_DBI ERROR: Query handle already exists**
  This error is thrown when the user tries to create a query handle, in `db_query` or `db_prepare`, using a name which already exists as a query handle for a different query.

- **XSB_DBI_008: XSB_DBI ERROR: Not all parameters supplied**
  This error is thrown when the user tries to execute a prepared statement, using `db_prepare_execute`, without supplying values for all the parameters in the statement.

- **XSB_DBI_009: XSB_DBI ERROR: Unbound variable in parameter list**
  This error is thrown when the user tries to execute a prepared statement, using `db_prepare_execute`, without binding all the parameters of the statement.

5.5   Notes on specific drivers

ODBC Driver

The ODBC driver has been tested in Linux using the `unixodbc` driver manager. It currently supports the following functionality: (a) connecting to a database using a DSN, (b) direct querying of the database, (c) using prepared statements to query the database, (d) closing a statement handle, and (d) disconnecting from the database. The ODBC driver has also been tested under Windows and Cygwin.

MySQL Driver

The MySQL driver provides access to the native MySQL C API. Currently, it has support for the following functionality: (a) connecting to a database using `db_connect`, (b) direct querying of the database, (c) using prepared statements to query the database, (d) closing a statement handle, and (e) disconnecting from the database.

The MySQL driver has been tested under Linux and Windows.

Driver for the Embedded MySQL Server

This driver provides access to the Embedded MySQL Server Library `libmysqld`. Currently, it has support for the following functionality: (a) connecting to a database `db_connect`, (b) direct querying of the database, (c) using prepared statements to query the database, (d) closing a statement handle, and (e) disconnecting from the database.

The MySQL driver for Embedded MySQL Server has been tested under Linux.

In order to use this driver, you will need:

- MySQL with Embedded Server installed on your machine. If your don't have a precompiled binary distribution of MySQL, which was configured with libmysqld support (the embedded server library), you will need to build MySQL from sources and configure it with the `--with-embedded-server` option.

- append to `/etc/my.cnf` (or `/etc/mysql/my.cnf` – whichever is used on your machine) or ∼/`.my.cnf`:

```
[mysqlembedded_driver_SERVER]
language = /usr/local/mysql/share/mysql/english
datadir = .....
```

  You will probably need to replace `/usr/local/mysql/share/mysql/english` with a directory appropriate for your MySQL installation.

You might also need to set the `datadir` option to specify the directory where the databases managed by the embedded server are to be kept. This has to be done if there is a possibility of running the embedded MySQL server alongside the regular MySQL server. In that case, the `datadir` directory of the embedded server must be different from the `datadir` directory of the regular server (which is likely to be specified using the `datadir` option in `/etc/my.cnf` or `/etc/mysql/my.cnf`. This is because specifying the same directory might lead to a corruption of your databases. See `http://dev.mysql.com/doc/refman/5.1/en/multiple-servers.html` for further details on running multiple servers.

Please note that loading the embedded MySQL driver increases the memory footprint of XSB. This additional memory is released automatically when XSB exits. If you need to release the memory before exiting XSB, you can call `driverMySQLEmbedded_lib_end` after disconnecting from MySQL. Note that once `driverMySQLEmbedded_lib_end` is called, no further connections to MySQL are allowed from the currently running session of XSB (or else XSB will exit abnormally).

Chapter 6

Introduction to XSB Packages

An XSB package is a piece of software that extends XSB functionality but is not critical to programming in XSB. Around a dozen packages are distributed with XSB, ranging from simple meta-interpreters to complex software systems. Some packages provide interfaces from XSB to other software systems, such as Perl, SModels or Web interfaces (as in the `libwww` package). Others, such as the CHR and Flora packages, extend XSB to different programming paradigms.

Each package is distributed in the `$XSB_DIR/packages` subdirectory, and has two parts: an initialization file, and a subdirectory in which package source code files and executables are kept. For example, the `xsbdoc` package has files `xsbdoc.P`, `xsbdoc.xwam`, and a subdirectory, `xsbdoc`. If a user doesn't want to retain `xsbdoc` (or any other package) he or she may simply remove the initialization files and the associated subdirectory without affecting the core parts of the XSB system.

Several of the packages are documented in this manual in the various chapters that follow. However, many of the packages contain their own manuals. For these packages, we provide only a summary of their functionality in Chapter 14.

Chapter 7

POSIX Regular Expression and Wildcard Matching

By Michael Kifer

XSB has an efficient interface to POSIX pattern regular expression and wildcard matching functions. To take advantage of these features, you must build XSB using a C compiler that supports POSIX 1.0 (for regular expression matching) and the forthcoming POSIX 2.0 (for wildcard matching). The recent versions of GCC and SunPro compiler will do, as probably will many other compilers. This also works under Windows, provided you install Cygnus' CygWin and use GCC to compile [1].

7.1 `regmatch`: Regular Expression Matching and Substitution

The following discussion assumes that you are familiar with the syntax of regular expressions and have a reasonably good idea about their capabilities. One easily accessible description of POSIX regular expressions is found in the on-line Emacs manual.

The regular expression matching functionality is provided by the package called `Regmatch`. To use it interactively, type:

```
:- [regmatch].
```

If you are planning to use pattern matching from within an XSB program, then you need to include the following directive:

```
:- import re_match/5, re_bulkmatch/5,
          re_substitute/4, re_substring/4,
          re_charlist_to_string/2
    from regmatch.
```

---

[1]This package has not yet been ported to the multi-threaded engine.

Matching.   The predicates `re_match/5` and `re_bulkmatch/5` perform regular expression match-
ing. The predicate `re_substitute/4` replaces substrings in a list with strings from another list
and returns the resulting new string.

The `re_match/5` predicate has the following calling sequence:

```
re_match(+Regexp, +InputStr, +Offset, ?IgnoreCase, -MatchList)
```

`Regexp` is a regular expression, *e.g.*, "`abc([^;,]*); (dd|ee)*;`". It can be a Prolog atom or string
(*i.e.*, a list of characters). The above expression matches any substring that has "abc" followed by
a sequence of characters none of which is a ";" or a ",", followed by a "; ", followed by a sequence
that consists of zero or more of "dd" or "ee" segments, followed by a ";". An example of a string
where such a match can be found is "`123abc&*^; ddeedd;poi`".

`InputStr` is the string to be matched against. It can be a Prolog atom or a string (list of
characters). `Offset` is an integer offset into the string. The matching process starts at this offset.
`IgnoreCase` indicates whether the case of the letters is to be ignored. If this argument is an
uninstantiated variable, then the case is *not* ignored. If this argument is bound to a non-variable,
then the case *is* ignored.

The last argument, `MatchList`, is used to return the results. It must unify with a list of the
form:

```
[match(beg_off0,end_off0), match(beg_off1,end_off1), ...]
```

The term `match(beg_off0,end_off0)` represents the substring that matches the *entire* regular
expression, and the terms `match(beg_off1,end_off1)`, ..., represent the matches corresponding
to the *parenthesized subexpressions* of the regular expression. The terms **beg_off** and **end_off** above
are integers that specify beginning and ending offsets of the various matches. Thus, **beg_off0** is
the offset into `InputStr` that points to the start of the maximal substring that matches the entire
regular expression; **end_off0** points to the end of such a substring. In our case, the maximal
matching substring is "abc&*^; ddeedd;" and the first term in the list returned by

```
| ?- re_match('abc([^;,]*); (dd|ee)*;', '123abc&*^; ddeedd;poi', 0, _,L).
```

is `match(3,18)`.

The most powerful feature of POSIX pattern matching is the ability to remember and return
substrings matched by parenthesized subexpressions. When the above predicate succeeds, the terms
2,3, etc., in the above list represent the offsets for the matches corresponding to the parenthesized
expressions 1,2,etc. For instance, our earlier regular expression "`abc([^;,]*); (dd|ee)*;`" has
two parenthetical subexpressions, which match "&*^" and "dd, respectively. So, the complete
output from the above call is:

```
L = [match(3,18),match(6,9),match(15,17)]
```

The maximal number of parenthetical expressions supported by the Regmatch package is 30.
Partial matches to parenthetical expressions 31 and over are discarded.

The match-terms corresponding to parenthetical expressions can sometimes report "*no-use.*" This is possible when the regular expression specifies that zero or more occurrences of the parenthesized subexpression must be matched, and the match was made using zero subexpressions. In this case, the corresponding match term is `match(-1,-1)`. For instance,

```
| ?- re_match('ab(de)*', 'abcd',0,_,L).
L = [match(0,2),match(-1,-1)]
yes
```

Here the match that was found is the substring "ab" and the parenthesized subexpression "de" was not used. This fact is reported using the special match term `match(-1,-1)`.

Here is one more example of the power of POSIX regular expression matching:

```
| ?- re_match("a(b*|e*)cd\\1",'abbbcdbbbbbo', 0, _, M).
```

Here the result is:

```
M = [match(0,9),match(1,4)]
```

The interesting features here are the positional parameter \\1 and the alternating parenthetical expression `a(b*|e*)`. The alternating parenthetical expression here can match any sequence of b's *or* any sequence of e's. Note that if the string to be matched is not known when we write the program, we will not know a priori which sequence will be matched: a sequence of b's or a sequence of e's. Moreover, we do not even know the length of that sequence.

Now, suppose, we want to make sure that the matching substrings look like this:

```
abbbcdbbb
aeeeecdeeee
abbbbbbcdbbbbbb
```

How can we make sure that the suffix that follows "cd" is exactly the same string that is stuck between "a" and "cd"? This is what \\1 precisely does: it represents the substring matched by the first parenthetical expression. Similarly, you can use \\2, etc., if the regular expression contains more than one parenthetical expression.

The following example illustrates the use of the offset argument:

```
| ?- re_match("a(b*|e*)cd\\1",'abbbcdbbbbboabbbcdbbbbbb',2,_,M).
```

```
M = [match(12,21),match(13,16)]
```

Here, the string to be matched is double the string from the previous example. However, because we said that matching should start at offset 2, the first half of the string is not matched.

The `re_match/5` predicate fails if `Regexp` does not match `InputStr` or if the term specified in `MatchList` does not unify with the result produced by the match. Otherwise, it succeeds.

We should also note that parenthetical expressions can be represented using the \(...\) notation. What if you want to match a "(" then? You must escape it with a "\\" then:

```
| ?- re_match("a(b*)cd\\(",'abbbcd(bbo', 0, _, M).
```

```
M = [match(0,7),match(1,4)]
```

Now, what about matching the backslash itself? Try harder: you need four backslashes:

```
| ?- re_match("a(b*)cd\\\\",'abbbcd\bbo', 0, _, M).
```

```
M = [match(0,7),match(1,4)]
```

The predicate `re_bulkmatch/5` has the same calling sequence as `re_match/5`, and the meaning of the arguments is the same, except the last (output) argument. The difference is that `re_bulkmatch/5` ignores parenthesized subexpressions in the regular expression and instead of returning the matches corresponding to these parenthesized subexpressions it returns the list of all matches for the top-level regular expression. For instance,

```
| ?- re_bulkmatch('[^a-zA-Z0-9]+', '123&*-456 )7890% 123', 0, 1, X).
```

```
X = [match(3,6),match(9,11),match(15,17)]
```

Extracting the matches.   The predicate `re_match/5` provides us with the offsets. How can we actually get the matched substrings? This is done with the help of the predicate `re_substring/4`:

```
    re_substring(+String, +BeginOffset, +EndOffset, -Result).
```

This predicate works exactly like `substring/4` described in Section 1.5, except that the resulting substring is not interned (if it is an atom). All you can do with this string is to immediately convert it into a list (using `atom_codes/2`) or into a true atom (using `intern_string/2`, which must be imported from module machine).

The reason for these complications is to allow the user to control the size of the atom table. At present, XSB does not have atom table garbage collection, so heavy use of string manipulation functions can result in atom table overflow. This danger is particularly severe when XSB is used for processing HTML pages. This predicate will become an alias to `substring/4` when atom garbage collection will be added to XSB.

On the other hand, converting strings into lists (without interning them first) is safe, because lists are garbage-collected in XSB Version 2.0.

Here is a complete example that shows matching followed by a subsequent extraction of the matches:

```
| ?- import intern_string/2 from machine.
```

```
| ?- Str = 'abbbcd\bbo',
      re_match("a(b*)cd\\\\",Str,0,_,[match(X,Y), match(V,W)|L]),
      re_substring(Str,X,Y,UninternedMatch),
      intern_string(UninternedMatch,Match),
      re_substring(Str,V,W,UninternedParen1),
      atom_codes(UninternedParen1,Paren1).

Str = abbbcd\bbo
X = 0
Y = 7
V = 1
W = 4
L = []
UninternedMatch = abbbcd\
Match = abbbcd\
UninternedParen1 = bbb
Paren1 = [98,98,98]
```

Note that the strings `UninternedMatch` and `UninternedParen1` cannot be used by themselves. In the first case, we converted the string into a Prolog atom and in the second case into a string. The resulting objects (`Match` and `Paren1`) can be used in further computations.

Observe that XSB reports that `UninternedMatch` and `UninternedParen1` are both equal the string "bbb", while `Match` — the atom obtained from `UninternedMatch` — is different. This is because `UninternedMatch` and `UninternedParen1` are uninterned and both occupy the same physical space. Thus, the second call to `re_substring/4` overrides the value stored in this location by the first call.

Substitution. The predicate `re_substitute/4` has the following invocation:

```
re_substitute(+InputStr, +SubstrList, +SubstitutionList, -OutStr)
```

This predicate works exactly like `string_substitute/4` described in Section 1.5, except that the result of the substitution is not interned (for the same reason as in `re_substring/4`. This predicate will become an alias to `string_substitute/4` when atom garbage collection will be added to XSB.

```
| ?- re_bulkmatch('[^a-zA-Z0-9]+', '123&*-456 )7890| 123', 0, _, X),
     re_substitute('123&*-456 )7890| 123', X, ['+++'], Y).

X = [match(3,6),match(9,11),match(15,17)]
Y = 123+++456+++7890+++123
```

Efficiency considerations.

- Try not to work with too many regular expressions at once. Before a regular expression can be used, it must be compiled (which **re_match/5** does automatically). **re_match/5** maintains a cache of compiled regular expressions, so they do not need to be compiled each time they are used. However, if more than 8–10 expressions are used simultaneously, repeated recompilation might result.

- When a list of characters is passed to any one of the above predicates, it is converted into a C string. This can be expensive, if done too often for the same string.

  One way to circumvent the problem is to use **atom_codes/2** to first convert the list into an atom and then use that atom repeatedly in the match operations. One problem here might be the aforementioned overflow of the atom table. So, if this is a concern, the following predicate (which always succeeds) can help:

  ```
  re_charlist_to_string(+ListOfCharacters, -String)
  ```

  This predicate converts lists of characters into uninterned strings, which can be used without the fear of atom table overflow:

  ```
  | ?- re_charlist_to_string("abcdefg",L).
  ```

  ```
  L = abcdefg
  ```

  The resulting string can be passed to **re_match/5**, **re_substitute/5**, and **re_substring/4** for further processing.

  Note, however: you cannot call **re_charlist_to_string** before you finished working with the string generated by the previous call: all calls to this function use the same static buffer to hold the output string, so each subsequent call to **re_charlist_to_string** will override the previously generated strings.

7.2 `wildmatch`: Wildcard Matching and Globing

These interfaces are implemented using the `Wildmatch` package of XSB. This package provides the following functionality:

1. Telling whether a wildcard, like the ones used in Unix shells, match against a given string. Wildcards supported are of the kind available in tcsh or bash. Alternating characters (*e.g.*, "[abc]" or "[^abc]") are supported.

2. Finding the list of all file names in a given directory that match a given wildcard. This facility generalizes `directory/2` (in module `directory`), and it is much more efficient.

3. String conversion to lower and upper case.

To use this package, you need to type:

```
| ?- [wildmatch].
```

If you are planning to use it in an XSB program, you need this directive:

```
:- import glob_directory/4, wildmatch/3, convert_string/3 from wildmatch.
```

The calling sequence for `glob_directory/4` is:

```
glob_directory(+Wildcard, +Directory, ?MarkDirs, -FileList)
```

The parameter `Wildcard` can be either a Prolog atom or a Prolog string. `Directory` is also an atom or a string; it specifies the directory to be globbed. `MarkDirs` indicates whether directory names should be decorated with a trailing slash: if `MarkDirs` is bound, then directories will be so decorated. If MarkDirs is an unbound variable, then trailing slashes will not be added.

`FileList` gets the list of files in `Directory` that match `Wildcard`. If `Directory` is bound to an atom, then `FileList` gets bound to a list of atoms; if `Directory` is a Prolog string, then `FileList` will be bound to a list of strings as well.

This predicate succeeds is at least one match is found. If no matches are found or if `Directory` does not exist or cannot be read, then the predicate fails.

The calling sequence for `wildmatch/3` is as follows:

```
wildmatch(+Wildcard, +String, ?IgnoreCase)
```

`Wildcard` is the same as before. `String` represents the string to be matched against `Wildcard`. Like `Wildcard`, `String` can be an atom or a string. `IgnoreCase` indicates whether case of letters should be ignored during matching. Namely, if this argument is bound to a non-variable, then the case of letters is ignored. Otherwise, if `IgnoreCase` is a variable, then the case of letters is preserved.

This predicate succeeds when `Wildcard` matches `String` and fails otherwise.

The calling sequence for `convert_string/3` is as follows:

```
convert_string(+InputString, +OutputString, +ConversionFlag)
```

The input string must be an atom or a character list. The output string must be unbound. Its type will "atom" if so was the input and it will be a character list if so was the input string. The conversion flag must be the atom `tolower` or `toupper`.

This predicate always succeeds, unless there was an error, such as wrong type argument passed as a parameter.

Chapter 8

`perlmatch`: Perl-based String Manipulation

By Michael Kifer and Jin Yu

XSB has an efficient interface to the Perl interpreter, which allows XSB programs to use powerful Perl pattern matching capabilities. this interface is provided by the `Perlmatch` package. You need Perl 5.004 or later to be able to take advantage of this service [1].

This package is mostly superseded by the the more efficient POSIX `Regmatch` package described in the previous section. However, Perl regular expressions provide certain features not available in the `Regmatch` package, such as the ability to perform global replacements of matched substrings. Also, the `Perlmatch` package has a different programming interface, which is modeled after the interface provided by Perl itself. So, if you are a big fan of Perl, this package is for you.

The following discussion assumes that you are familiar with the syntax of Perl regular expressions and have a reasonably good idea about the capabilities of Perl matching and string substitution functions.

In the interactive mode, you must first load the `Perlmatch` package:

```
:- [perlmatch].
```

In a program, you must import the package predicates:

```
:- import bulk_match/3, get_match_result/2, try_match/2,
   next_match/0, perl_substitute/3, load_perl/0, unload_perl/0
   from perlmatch.
```

---

[1]This packages is not yet supported by the multi-threaded engine.

8.1 Iterative Pattern Matching

There are two ways to do matching. One is to first do the matching operation and then count the beans. To find the first match, do:

```
:- try_match( +String, +Pattern ).
```

Both arguments must be of the XSB string data types. If there is a match in the string, the submatches $1, $2, etc., the prematch substring $' (*i.e.*, the part before the match), the post-match substring $' (*i.e.*, the part after the match), the whole matched substring $&, and the last parentheses match substring $+ will be stored into a global data structure, and the predicate try_match(string, pattern) succeeds. If no match is found, this predicate fails.

The ability to return parts of the match using the Perl variables $+, $1, $2, etc., is an extremely powerful feature of Perl. As we said, a familiarity with Perl matching is assumed, but we'll give an example to stimulate the interest in learning these capabilities of Perl. For instance, m/(\d+)\.?(\d*)(+)/— matches a valid number. Moreover, if the number had the form 'xx.yy', then the Perl variable $1 will hold 'xx' and $1 will hold 'yy'. If the number was of the form '.zz', then $1 and $2 will be empty, and $3 will hold 'zz'.

XSB-Perl interface provides access to all these special variables using the predicate `get_match_result()`. The input variables string and pattern are of XSB string data types. For example:

For instance,

```
:- try_match('First 11.22 ; next:  3.4', 'm/(+)?(*)/g').
yes.
```

finds the character which precedes by 's' in the string 'this is a test'. The first match is 'is'.

Now, we can use `get_match_result()` to find the submatches. The first argument is a tag, which can be 1 through 20, denoting the Perl variables $1 – $20. In addition, the entire match can be found using the tag `match`, the part before that is hanging off the tag `prematch` and the part of the string to the right of the match is fetched with the tag `postmatch`. For instance, in the above, we shall have:

This function is used to fetch the pattern match results $1, $2, etc., $', $', $& and $+, as follows:

```
:- get_match_result(1,X).
X=11
yes
:- get_match_result(2,X).
X=22
yes
:- get_match_result(3,X).
no
:- get_match_result(4,X).
```

```
no
:- get_match_result(prematch,X)
X=First  (including 1 trailing space)
yes
:- get_match_result(postmatch,X)
X= ; next 3.4
yes
:- get_match_result(match,X)
11.22
yes
```

As you noticed, if a tag fetches nothing (like in the case of Perl variables $3, $4, etc.), then the predicate fails.

The above is not the only possible match, of course. If we want more, we can call:

```
:- next_match.
```

This will match the second number in the string. Correspondingly, we shall have:

```
:- get_match_result(1,X).
X=3
yes
:- get_match_result(2,X).
X=4
yes
:- get_match_result(3,X).
no
:- get_match_result(4,X).
no
:- get_match_result(prematch,X)
X=First 11.22 ; next
yes
:- get_match_result(postmatch,X)
no
:- get_match_result(match,X)
3.4
yes
```

The next call to `next_match` would fail, because there are no more matches in the given string. Note that `next_match` and `get_match_result` do not take a string and a pattern as argument—they work off of the previous `try_match`. If you want to change the string or the pattern, call `try_match` again (with different parameters).

Note: To be able to iterate using `next_match`, the perl pattern must be *global i.e.*, it must have the modifier 'g' at the end (*e.g.*, `m/a.b*/g`). Otherwise, `next_match` simply fails and only one (first)

match will be returned.

8.2 Bulk Matching

XSB-perl interface also supports *bulk* matching through the predicate `bulk_match/3`. Here, you get all the substrings that match the patters at once, in a list, and then you can process the matches as you wish. However, this does not give you full access to submatches. More precisely, if you use parenthesized expressions, then you get the list of non-null values in the variables $1, $2, etc. If you do not use parenthesized regular expressions, you get the result of the full match. For instance,

```
:- bulk_match('First 11.22 ; next: 3.4', 'm/(\d+)\.?(\d*)/g', Lst).
Lst=[11,22,3,4]
yes
:- bulk_match('First 11.22 ; next: 3.4', 'm/\d+\.?\d*/g', Lst).
Lst=[11.22,3.4]
yes
```

`bulk_match/3` never fails. If there is no match, it returns an empty list.

Please note that you must specify 'g' at the end of the pattern in order to get something useful. This ia Perl thing! If you do not, instead of returning a list of matches, Perl will think that you just want to test if there is a match or not, and it will return `[1]` or `[]`, depending on the outcome.

8.3 String Substitution

The last feature of the XSB-Perl interface is string substitution based on pattern matching. This is achieved through the predicate `string_substitute/3`:

```
:- perl_substitute(+String, +PerlSubstitutionExpr, -ResultString).
```

We assume you are familiar with the syntax of Perl substitution expressions. Here we just give an example of what kind of things are possible:

```
:- perl_substitute('this is fun', 's/(this) (is)(.*)/\2 \1\3?/', Str).
Str=is this fun?
```

8.4 Unloading Perl

Playing with Perl is nice, but this also means that both XSB and the Perl interpreter are loaded in the main memory. If you do not need Perl for some time and memory is at premium, you can unload the Perl interpreter:

```
:- unload_perl.
```

This predicate always succeeds. If you need Perl matching features later, you can always come back to it: it is loaded automatically each time you use a pattern matching or a string substitution predicate.

5

Chapter 9

sgml and xpath: SGML/XML/HTML Parsers and XPath

By Rohan Shirwaikar

9.1   Introduction

This suite of packages consists of the `sgml` package, which can parse XML, HTML, XHTML, and even SGML documents and the `xpath` package, which supports XPath queries on XML documents. The `sgml` package is an adaptation of a similar package in SWI Prolog and a port of SWI's codebase with some minor changes. The `xpath` package provides an interface to the popular `libxml2` library, which supports XPath and XML parsing, and is used in Mozilla based browsers. At present, the XML parsing capabilities of `libxml2` are not utilized explicitly in XSB, but such support might be provided in the future. The `sgml` package does not rely on `libxml2` [1].

Installation and configuration.   The `sgml` package does not require any installation steps under Unix-based systems or under Cygwin. Under native Windows, if you downloaded XSB from CVS, you need to compile the package as follows:

```
cd XSB\packages\sgml\cc
nmake /f NMakefile.mak
```

You need MS Visual Studio for that. If you downloaded a prebuilt version of XSB, then the `sgml` package should have already been compiled for you and no installation is required.

The details of the `xpath` package and the corresponding configuration instructions appear in Section 9.4.

---

[1]This package has not yet been tested for thread-safety

9.2   Overview of the SGML Parser

The `sgml` package accepts input in the form of files, URLs and Prolog atoms. To load the `sgml` parser, the user should type

```
 ?- [sgml].
```

at the prompt. If `test.html` is a file with the following contents

```
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 3.2//EN">

<html>
<head>
<title>Demo</title>
</head>
<body>

<h1 align=center>This is a demo</title>

<p>Paragraphs in HTML need not be closed.

<p>This is called 'omitted-tag' handling.
</body>
</html>
```

then the following call

```
?- load_html_structure(file('test.html'), Term, Warn).
```

will parse the document and bind `Term` to the following Prolog term:

```
[ element(html,
          [],
          [ element(head,
                    [],
                    [ element(title,
                              [],
                              [ 'Demo'
                              ])
                    ]),
            element(body,
                    [],
                    [ '\n',
                      element(h1,
                              [ align = center
```

```
                                ],
                                [ 'This is a demo'
                                ]),
                        '\n\n',
                        element(p,
                                [],
                                [ 'Paragraphs in HTML need not be closed.\n'
                                ]),
                        element(p,
                                [],
                                [ 'This is called 'omitted-tag\' handling.'
                                ])
                ])
        ])
].
```

The XML document is converted into a list of Prolog terms of the form `element(Name,Attributes,Content)`. Each term corresponds to an XML element. *Name* represents the name of the element. *Attributes* is a list of attribute-value pairs of the element. *Content* is a list of child-elements and CDATA. For instance,

```
<aaa>fooo<bbb>foo1</bbb></aaa>
```

will be parsed as

```
element(aaa,[],[fooo, element(bbb,[],[foo1])])
```

Entities (e.g. `<`) are returned as part of `CDATA`, unless they cannot be represented. See **load_sgml_structure/3** for details.

9.3   Predicate Reference

9.3.1   Loading Structured Documents

SGML, HTML, and XML documents are parsed by the predicate **load_structure/4**, which has many options. For convenience, a number of commonly used shorthands are provided to parse SGML, XML, HTML, and XHTML documents respectively.

load_sgml_structure(*+Source, -Content, -Warn*)

load_xml_structure(*+Source, -Content, -Warn*)

load_html_structure(*+Source, -Content, -Warn*)

load_xhtml_structure(*+Source, -Content, -Warn*)

The parameters of these predicates have the same meaning as those in **load_structure/4**, and are described below.

The above predicates (in fact, just `load_xml_structure/3` and `load_html_structure/3`) are the most commonly used predicates of the `sgml` package. The other predicates described in this section are needed only for advanced uses of the package.

load_structure(+*Source, -Content, +Options, -Warn*)
> *Source* can have one of the following forms: `url(`*url*`)`, `file(`*file name*`)`, `string(`*'document as a Prolog atom'*`)`. The parsed document is returned in *Content*. *Warn* is bound to a (possibly empty) list of warnings generated during the parsing process. *Options* is a list of parameters that control parsing, which are described later.
>
> The list *Content* can have the following members:
>
> **A Prolog atom**
> > Atoms are used to represent character strings, i.e., `CDATA`.
>
> **element(**ature*Name, Attributes, Content*)
> > *Name* is the name of the element tag. Since SGML is case-insensitive, all element names are returned as lowercase atoms.
> >
> > *Attributes* is a list of pairs the form *Name=Value*, where *Name* is the name of an attribute and *Value* is its value. Values of type `CDATA` are represented as atoms. The values of multi-valued attributes (`NAMES`, *etc.*) are represented as a lists of atoms. Handling of the attributes of types `NUMBER` and `NUMBERS` depends on the setting of the `number(+NumberMode)` option of **set_sgml_parser/2** or **load_structure/3** (see later). By default the values of such attributes are represented as atoms, but the `number(...)` option can also specify that these values must be converted to Prolog integers.
> >
> > *Content* is a list that represents the content for the element.
>
> **entity(**austa*Code*)
> > If a character entity (e.g., `Α`) is encountered that cannot be represented in the Prolog character set, this term is returned. It represents the code of the encountered character (e.g., `entity(913)`).
>
> **entity(**austa*Name*)
> > This is a special case of `entity(Code)`, intended to handle special symbols by their name rather than character code. If an entity refers to a character entity holding a single character, but this character cannot be represented in the Prolog character set, this term is returned. For example, if the contents of an element is `Α < Β` then it will be represented as follows:
> >
> > `[ entity('Alpha'), ' < ', entity('Beta') ]`
> >
> > Note that entity names are case sensitive in both SGML and XML.
>
> **sdata(**austa*Text*)
> > If an entity with declared content-type `SDATA` is encountered, this term is used. The data of the entity instantiates *Text*.

ndata(*Text***)**

> If an entity with declared content-type `NDATA` is encountered, this term is used. The data instantiates *Text*.

pi(*Text***)**

> If a processing instruction is encountered (`<?...?>`), *Text* holds the text of the processing instruction. Please note that the `<?xml ...?>` instruction is ignored and is not treated as a processing instruction.

The *Options* parameter is a list that controls parsing. Members of that list can be of the following form:

dtd(*?DTD***)**

> Reference to a DTD object. If specified, the `<!DOCTYPE ...>` declaration supplied with the document is ignored and the document is parsed and validated against the provided DTD. If the DTD argument is a variable, then a the variable *DTD* gets bound to the DTD object created out of the DTD supplied with the document.

dialect(*+Dialect***)**

> Specify the parsing dialect. The supported dialects are `sgml` (default), `xml` and `xmlns`.

space(*+SpaceMode***)**

> Sets the space handling mode for the initial environment. This mode is inherited by the other environments, which can override the inherited value using the XML reserved attribute **xml:space**. See Section 9.3.2 for details.

number(*+NumberMode***)**

> Determines how attributes of type `NUMBER` and `NUMBERS` are handled. If `token` is specified (the default) they are passed as an atom. If `integer` is specified the parser attempts to convert the value to an integer. If conversion is successful, the attribute is represented as a Prolog integer. Otherwise the value is represented as an atom. Note that SGML defines a numeric attribute to be a sequence of digits. The `-` (minus) sign is not allowed and `1` is different from `01`. For this reason the default is to handle numeric attributes as tokens. If conversion to integer is enabled, negative values are silently accepted and the minus sign is ignored.

defaults(*+Bool***)**

> Determines how default and fixed attributes from the DTD are used. By default, defaults are included in the output if they do not appear in the source. If `false`, only the attributes occurring in the source are emitted.

file(*+Name***)**

> Sets the name of the input file for error reporting. This is useful if the input is a stream that is not coming from a file. In this case, errors and warnings will not have the file name in them, and this option allows one to force inclusion of a file name in such messages.

line(*+Line***)**

> Sets the starting line-number for reporting errors. For instance, if `line(10)` is specified and an error is found at line X then the error message will say that the error occurred at line X+10. This option is used when the input stream does not start with the first line of a file.

max_errors(+*Max*)

>   Sets the maximum number of errors. The default is 50. If this number is reached, the
>   following exception is raised:
>
>   ```
>   error(limit_exceeded(max_errors, Max), _)
>   ```

9.3.2   Handling of White Spaces

Four modes for handling white-spaces are provided. The initial mode can be switched using the
`space(SpaceMode)` option to **load_structure/3** or **set_sgml_parser/2**. In XML mode, the mode
is further controlled by the **xml:space** attribute, which may be specified both in the DTD and in
the document. The defined modes are:

space(sgml)

>   Newlines at the start and end of an element are removed. This is the default mode for the
>   SGML dialect.

space(preserve)

>   White space is passed literally to the application. This mode leaves all white space handling
>   to the application. This is the default mode for the XML dialect.

space(default)

>   In addition to `sgml` space-mode, all consecutive whitespace is reduced to a single space-
>   character.

space(remove)

>   In addition to `default`, all leading and trailing white-space is removed from `CDATA` objects.
>   If, as a result, the `CDATA` becomes empty, nothing is passed to the application. This mode is
>   especially handy for processing data-oriented documents, such as RDF. It is not suitable for
>   normal text documents. Consider the HTML fragment below. When processed in this mode,
>   the spaces surrounding the three elements in the example below are lost. This mode is not
>   part of any standard: XML 1.0 allows only `default` and `preserve`.
>
>   ```
>   Consider adjacent <b>bold</b> <ul>and</ul> <it>italic</it> words.
>   ```
>
>   The parsed term will be `['Consider adjacent',element(b,[],[bold]),element(ul,[],`
>   `[and]),element(it,[],[italics]),words]`.

9.3.3   XML documents

The parser can operate in two modes: the `sgml` mode and the `xml` mode, as defined by the
`dialect(Dialect)` option. HTML is a special case of the SGML mode with a particular DTD.
Regardless of this option, if the first line of the document reads as below, the parser is switched
automatically to the XML mode.

```
<?xml ... ?>
```

Switching to XML mode implies:

- *XML empty elements*
  The construct `<element attribute ... attribute/>` is recognized as an empty element.

- *Predefined entities*
  The following entities are predefined: `<` (<), `>` (>), `&` (&), `'` (') and `"` (").

- *Case sensitivity*
  In XML mode, names of tags and attributes are case-sensitive, except for the DTD reserved names (i.e. `ELEMENT`, *etc.*).

- *Character classes*
  In XML mode, underscore (`_`) and colon (`:`) are allowed in names.

- *White-space handling*
  White space mode is set to `preserve`. In addition, the XML reserved attribute **xml:space** is honored; it may appear both in the document and the DTD. The `remove` extension (see `space(remove)` earlier) is allowed as a value of the **xml:space** attribute. For example, the DTD statement below ensures that the **pre** element preserves space, regardless of the default processing mode.

```
<!ATTLIST pre xml:space nmtoken #fixed preserve>
```

XML Namespaces

Using the dialect `xmlns`, the parser will recognize XML namespace prefixes. In this case, the names of elements are returned as a term of the format

$$URL\!:\!LocalName$$

If an identifier has no namespace prefix and there is no default namespace, it is returned as a simple atom. If an identifier has a namespace prefix but this prefix is undeclared, the namespace prefix rather than the related URL is returned.

Attributes declaring namespaces (`xmlns:`*ns*`=`*url*) are represented in the translation as regular attributes.

9.3.4 DTD-Handling

The DTD (**D**ocument **T**ype **D**efinition) are internally represented as objects that can be created, freed, defined, and inspected. Like the parser itself, it is filled by opening it as a Prolog output stream and sending data to it. This section summarizes the predicates for handling the DTD.

new_dtd(*+DocType, -DTD, -Warn***)**

> Creates an empty DTD for the named *DocType*. The returned DTD-reference is an opaque term that can be used in the other predicates of this package. *Warn* is the list of warnings generated.

free_dtd(*+DTD, -Warn***)**

> Deallocate all resources associated to the DTD. Further use of *DTD* is invalid. *Warn* is the list of warnings generated.

open_dtd(*+DTD, +Options, -Warn***)**

> This opens and loads a DTD from a specified location (given in the *Options* parameter (see next). *DTD* represents the created DTD object after the source is loaded. *Options* is a list options. Currently the only option supported is *source(location)*, where *location* can be of one of these forms:

> > `url(`*url*`)`
> > `file(`*fileName*`)`
> > `string(`*'document as a Prolog atom'*`)`.

dtd(*+DocType, -DTD, -Warn***)**

> Certain DTDs are part of the system and have known doctypes. Currently, 'HTML' and 'XHTML' are the only recognized built-in doctypes. Such a DTD can be used for parsing simply by specifying the doctype. Thus, the `dtd/3` predicate takes the doctype name, finds the DTD associated with the given doctype, and creates a dtd object for it. *Warn* is the list of warnings generated.

dtd(*+DocType, -DTD, +DtdFile -Warn***)**

> The predicate parses the DTD present at the location *DtdFile* and creates the corresponding DTD object. *DtdFile* can have one of the following forms: `url(`*url*`)`, `file(`*fileName*`)`, `string(`*'document as a Prolog atom'*`)`.

9.3.5 Low-level Parsing Primitives

The following primitives are used only for more complex types of parsing, which might not be covered by the `load_structure/4` predicate.

new_sgml_parser(*-Parser, +Options, -Warn***)**

> Creates a new parser. *Warn* is the list of warnings generated. A parser can be used one or multiple times for parsing documents or parts thereof. It may be bound to a DTD or the DTD may be left implicit. In this case the DTD is created from the document prologue or (if it is not in the prologue) parsing is performed without a DTD. The *Options* list can contain the following parameters:

> **dtd(** *?DTD***)**

> > If *DTD* is bound to a DTD object, this DTD is used for parsing the document and

the document's prologue is ignored. If *DTD* is a variable, the variable gets bound to a created DTD. This DTD may be created from the document prologue or build implicitly from the document's content.

free_sgml_parser(*+Parser, -Warn***)**

Destroy all resources related to the parser. This does not destroy the DTD if the parser was created using the `dtd(`*DTD*`)` option. *Warn* is the list of warnings generated during parsing (can be empty).

set_sgml_parser(*+Parser, +Option, -Warn***)**

Sets attributes to the parser. *Warn* is the list of warnings generated. *Options* is a list that can contain the following members:

file(*File***)**

Sets the file for reporting errors and warnings. Sets the linenumber to 1.

line(*Line***)**

Sets the starting line for error reporting. Useful if the stream is not at the start of the (file) object for generating proper line-numbers. This option has the same meaning as in the `load_structure/4` predicate.

charpos(*Offset***)**

Sets the starting character location. See also the `file(File)` option. Used when the stream does not start from the beginning of a document.

dialect(*Dialect***)**

Set the markup dialect. Known dialects:

`sgml`

The default dialect. This implies markup is case-insensitive and standard SGML abbreviation is allowed (abbreviated attributes and omitted tags).

`xml`

This dialect is selected automatically if the processing instruction `<?xml ...>` is encountered.

`xmlns`

Process file as XML file with namespace support.

qualify_attributes(*Boolean***)**

Specifies how to handle unqualified attributes (i.e., without an explicit namespace) in XML namespace (`xmlns`) dialect. By default, such attributes are not qualified with namespace prefixes. If `true`, such attributes are qualified with the namespace of the element they appear in.

space(*SpaceMode***)**

Define the initial handling of white-space in `PCDATA`. This attribute is described in Section 9.3.2.

number(*NumberMode***)**

If `token` is specified (the default), attributes of type number are represented as a Prolog atom. If `integer` is specified, such attributes are translated into Prolog integers. If the

conversion fails (e.g., due to an overflow) a warning is issued and the value is represented as an atom.

doctype(*Element***)**

Defines the top-level element of the document. If a `<!DOCTYPE ...>` declaration has been parsed, this declaration is used. If there is no `DOCTYPE` declaration then the parser can be instructed to use the element given in `doctype(_)` as the top level element. This feature is useful when parsing part of a document (see the `parse` option to **sgml_parse/3**).

sgml_parse(*+Parser, +Options, -Warn***)**

Parse an XML file. The parser can operate in two input and two output modes. Output is a structured term as described with **load_structure/4**.

Warn is the list of warnings generated. A full description of *Options* is given below.

document(*+Term***)**

A variable that will be unified with a list describing the content of the document (see **load_structure/4**).

source(*+Source***)**

Source can have one of the following forms: `url(`*url*`)`, `file(`*fileName*`)`, `string(`*'document as a Prolog atom'*`)`. This option *must* be given.

content_length(*+Characters***)**

Stop parsing after the given number of *Characters*. This option is useful for parsing input embedded in *envelopes*, such as HTTP envelopes.

parse(*Unit***)**

Defines how much of the input is parsed. This option is used to parse only parts of a file.

file

Default. Parse everything upto the end of the input.

element

The parser stops after reading the first element. Using `source(`*Stream*`)`, this implies reading is stopped as soon as the element is complete, and another call may be issued on the same stream to read the next element.

declaration

This may be used to stop the parser after reading the first declaration. This is useful if we want to parse only the `doctype` declaration.

max_errors(*+MaxErrors***)**

Sets the maximum number of errors. If this number is exceeded, further writes to the stream will yield an I/O error exception. Printing of errors is suppressed after reaching this value. The default is 100.

syntax_errors(*+ErrorMode***)**

Defines how syntax errors are handled.

quiet

Suppress all messages.

print
> Default. Print messages.

9.3.6   External Entities

While processing an SGML document the document may refer to external data. This occurs in three places: external parameter entities, normal external entities and the `DOCTYPE` declaration. The current version of this tool deals rather primitively with external data. External entities can only be loaded from a file.

Two types of lines are recognized by this package:

```
DOCTYPE   doctype   file
```

```
PUBLIC   "Id "   file
```

The parser loads the entity from the file specified as *file*. The file can be local or a URL.

9.3.7   Exceptions

Exceptions are generated by the parser in two cases. The first case is when the user specifies wrong input. For example when specifying

```
load_structure( string('<m></m>'), Document, [line(xyz)], Warn)
```

The string `xyz` is not in the domain of `line`. Hence in this case a domain error exception will be thrown.

Exceptions are generated when XML being parsed is not well formed. For example if the input XML contains

```
'<m></m1>'
```

exceptions will be thrown.

In both cases the format of the exception is

```
error( sgml( error term),   error message)
warning( sgml( warning term),   warning message)
```

where *error term* or *warning term* can be of the form

- *pointer to the parser instance,*

- *line at which error occurred,*

- *error code.*

- *functor(argument)*, where *functor* and *argument* depend on the type of exception raised. For example,

  `resource-error(no-memory)` — if memory is unavailable

  `permission-error(file-name)` — no permission to read a file

  A `system-error(description) --- internal system error`

  `type-error(expected,actual)` — data type error

  `domain-error(functor,offending-value)` — the offending value is not in the domain of the functor. For instance, in `load_structure( string('<m></m>'), Document, [line(xyz)], Warn)`, `xyz` is not in the domain of `line`.

  `existence-error(resource)` — resource does not exist

  `limit-exceeded(limit,maxval)` — value exceeds the limit.

9.3.8   Unsupported features

The current parser is rather limited. While it is able to deal with many serious documents, it omits several less-used features of SGML and XML. Known missing SGML features include

- *NOTATION on entities*
  Though notation is parsed, notation attributes on external entity declarations are not represented in the output.

- *NOTATION attributes*
  SGML notations may have attributes, declared using `<!ATTLIST #NOT name attrib>`. Those data attributes are provided when you declare an external CDATA, NDATA, or SDATA entity. XML does not support external CDATA, NDATA, or SDATA entities, nor any of the other uses to which data attributes are put in SGML.

- *SGML declaration*
  The 'SGML declaration' is fixed, though most of the parameters are handled through indirections in the implementation.

- *The RANK feature*
  It is regarded as obsolete.

- *The LINK feature*
  It is regarded as too complicated.

- *The CONCUR feature*
  Concurrent markup allows a document to be tagged according to more than one DTD at the same time. It is not supported.

- *The Catalog files*
  Catalog files are not supported.

In the XML mode, the parser recognizes SGML constructs that are not allowed in XML. Also various extensions of XML over SGML are not yet realized. In particular, XInclude is not implemented.

9.3.9   Summary of Predicates

dtd/2	Find or build a DTD for a document type
free_dtd/1	Free a DTD object
free_sgml_parser/1	Destroy a parser
load_dtd/2	Read DTD information from a file
load_structure/4	Parse XML/SGML/HTML data into Prolog term
load_sgml_structure/3	Parse SGML file into Prolog term
load_html_structure/3	Parse HTML file into Prolog term
load_xml_structure/3	Parse XML file into Prolog term
load_xhtml_structure/3	Parse XHTML file into Prolog term
new_dtd/2	Create a DTD object
new_sgml_parser/2	Create a new parser
open_dtd/3	Open a DTD object as an output stream
set_sgml_parser/2	Set parser options (dialect, source, *etc.*)
sgml_parse/2	Parse the input
xml_name/1	Test atom for valid XML name
xml_quote_attribute/2	Quote text for use as an attribute
xml_quote_cdata/2	Quote text for use as PCDATA

9.4   XPath support

XPath is a query language for addressing parts of an XML document. In XSB, this support is provided by the `xpath` package. To use this package the `libxml2` XML parsing library must be installed on the machine. It comes with most Linux distributions, since it is part of the Gnome desktop, or one can download it from *http://xmlsoft.org/*. It is available for Linux, Solaris, Windows, and MacOS. Note that both the library itself and the `.h` files of that library must be installed. In some Linux distributions, the `.h` files might reside in a separate package from the package that contains the actual library. For instance, the library (`libxml2.so`) might be in the package called `libxml2` (which is usually installed by default), while the `.h` files hight be in the package `libxml2-devel` (which is usually *not* in a default installation).

On Unix-based systems (and MacOS), the package needs to be configured using the `configure` script found in the XSB's `packages/xpath` directory. The `configure` script has several options such as `--exec-prefix`, `--prefix`, etc, which may need to be specified if the package `libxml2` is installed in a non-standard place. In this case, `--exec-prefix=`*directory-of-libxml2* must specify the directory where the `lib/libxml2.so` file is located and `--prefix=`*include-dir* where the `include/libxml2` directory is found. The default `exec-prefix` is `/usr/lib` and `prefix` is `/usr`. If the defaults are fine (they would be, if the Gnome desktop is installed on your machine), then one can configure the package using

```
./configure
```

Otherwise, use

```
./configure --prefix=/usr/share --exec-prefix=/usr/local
```

if, for example, `libxml2.so` is in `/usr/local/lib/libxml2.so` and the included `.h` files are in `/usr/share/include/libxml2`.

On Windows and under Cygwin, the `libxml2` library is already included in the XSB distribution and does not need to be downloaded. If you are using a prebuilt XSB distribution for Windows, then you do not need to do anything—the package has already been built for you.

For Cygwin, you only need to run the `./configure` script without any options. This needs to be done regardless of whether you downloaded XSB from CVS or a released prebuilt version.

If you downloaded XSB from CVS and want to use it under native Windows (not Cygwin), then you would need to compile the XPath package, and you need Microsoft's Visual Studio. To compile the package one should do the following:

```
cd packages\xpath\cc
nmake /f NMakefile.mak
```

The following section assumes that the reader is familiar with the syntax of XPath and its capabilities. To load the `xpath` package, type

```
:-[xpath].
```

The program needs to include the following directive:

```
:- import parse_xpath/4 from xpath.
```

XPath query evaluation is done by using the `parse_xpath` predicate.

parse_xpath(*+Source, +XPathQuery, -Output, +NamespacePrefixList* **)**
> *Source* is a term of the format `url(`*url*`)`, `file(`*filename*`)` or `string(`*'XML-document-as-a-string'*`)`. It specifies that the input XML document is contained in a file, can be fetched from a URL, or is given directly as a Prolog atom.
>
> *XPathQuery* is a standard XPath query which is to be evaluated on the XML document in *Source*.
>
> *Output* gets bound to the output term. It represents the XML element returned after the XPath query is evaluated on the XML document in *Source*. The output term is of the form `string(`*'XML-document'*`)`. It can then be parsed using the `sgml` package described earlier.
>
> *NamespacePrefixList* is a space separated list of pairs of the form *prefix = namespace*. This specifies the namespace prefixes that are used in the XPath query.
>
> For example if the xpath expression is `'/x:html/x:head/x:meta'` where `x` is a prefix that stands for `'http://www.w3.org/1999/xhtml'`, then `x` would have to be defined as follows:

```
?- parse_xpath(url('http://w3.org'), '/x:html/x:head/x:meta', O4,
                     'x=http://www.w3.org/1999/xhtml').
```

In the above, the xpath query is `'/x:html/x:head/x:meta'` and the prefix has been defined as `'x=http://www.w3.org/1999/xhtml'`.

Chapter 10

`libwww`: The XSB Internet Access Package

By Michael Kifer

10.1 Features and Configuration

This package was inspired by the PiLLoW project. The XSB Libwww package offers much better performance and a superset of the PiLLoW functionality as related to the HTTP protocol, but this package does not implement the part of PiLLoW that deals with construction of Web pages [1].

The XSB Libwww is implemented in C and relies on the basic HTTP functions provided by the Libwww library developed by the WWW Consortium (`http://www.w3c.org/Library`). Therefore, this library must be installed in order for the XSB Libwww package to work. In addition, XSB must be configured to work with the Libwww library as follows:

```
configure --with-libwww=directory-where-Libwww-is-installed
```

If you forgot to configure XSB for use with the Libwww library, you do not need to reconfigure and recompile XSB. Instead, you can `cd` to `packages/libwww` and type the above command in that directory. This will enable the previously compiled version of XSB to work with Libwww.

Note that the Libwww library **cannot** be moved after you configured it. If it is moved, XSB will not be able to find it, because hard path names are built into some parts of this library at configuration time. Use the `--prefix` argument of `configure` to specify where you want to install the Libwww library.

One of the most important aspects of the Libwww package is that it allows XSB to dispatch multiple HTTP requests, which interleave their Web access phases. This can be a significant performance boost. Furthermore, the HTML, the XML, and the RDF parsers begin their work as

---

[1]This package has not yet been ported to the multi-threaded engine.

the fragments of pages arrive, so by the time the page is fully accepted, it is also parsed. Here is a list of features provided by the XSB Libwww package:

- HTML-4 parser.

- XML parser (non-validating).

- RDF parser.

- Page fetching (without parsing).

- Form handling.

- HTTP header information.

- Multiple, interleaved HTTP requests.

- Basic and digest authentication.

- Redirection and proxies.

10.2 Accessing Internet with Libwww

To start using the package, you must load it first:

```
:- [libwww].
```

10.2.1 Syntax

The general form of a Web call is as follows:

```
:- libwww_request([request1, request2, ..., request_n]).
```

Each request has the following syntax:

```
request_type(+URL, +RequestParams, -ResponseParams, -Result, -Status)
```

The request type functor must be either `htmlparse`, `xmlparse`, `rdfparse`, `fetch`, or `header`. The first two are requests to parse HTML/XML pages, respectively. `Fetch` is a request to bring in a page without parsing, and `header` is a request to retrieve only the header information (which is returned in the `ResponseParams` argument—see below). The URL must be an atom or a string (list of characters).[2]

---

[2] The string feature will be deprecated when XSB will have working atom garbage collection. When URL is a list of characters, then Result is also a list of characters, which eases the burden on the atom table and allows XSB to work longer before memory is exhausted.

10.2.2 Request Parameters

Request parameters must be either a variable (in which case the request is considered to not have special parameters) or a list. The following terms are allowed in that list:

- `timeout(+Secs)` – request timeout. If it is not specified, a default value (5 seconds) is used. Only the first request in a list should have the timeout value set. Timeouts that appear in subsequent requests are ignored.

- `authentication([c(+Realm,+Username,+Pasword),...])` – If the site requires authentication, you should specify it in a list as an argument to the `authentication/1` functor. `Realm` is a string that the servers return to let applications know which username/password pair to send (in case the application works with several pages that require different authentication). They are used as page identifiers for the authentication purposes. If `Realm` is an atom (*e.g.*, `authentication('FooSite', boo, moo)`), then when a Web server requests authentication for the `FooSite` realm, the Libwww package will send the foo/moo user-password pair.

  If `Realm` is a variable, then it is considered to match every realm. The Libwww package searches for matching authentication triples in the order they appear in the authentication list. Thus, the triple where `Realm` is a variable should appear last and serve as a default username/password pair.

- `formdata([attval_pair1, attval_pair2, ...])` — list of attribute/value pairs to fill out a form (in case `URL` is a CGI script). Each attribute/value pair must be an atom of the form `attr=val`.

- `selection(Taglist1,Taglist2,Taglist3)` — if the request is `htmlparse` or `xmlparse`, then this term provides control over which tags to parse. `Taglist1` is a list of tags that specifies inside which tags to parse. For instance, if it is `[ul,form]` than parsing will be done only inside these elements. Other elements will be ignored. `Taglist2` tells the system to stop parsing inside the corresponding elements. For instance, `[table]` means that parsing should be done only inside `ul` and `form` elements. However, if we hit a table during parsing, then parsing should stop unless we hit `ul` or `form` inside the `table` element. This switching of parsing on and off can continue to arbitrary depth. `Taglist3` is a list of tags that are to be ignored completely. That is, the parsing process will simply strip these tags (but not the text inside them). For instance, if `Taglist3` is `[p,i]` and the page contains "`<p>foo <i>moo</i>`" then parsing will be done as if the page contained just "`foo moo`".

- `if_modified_since(Date)` — fetch the document only if it has been modified after the specified date. The date must be given as an atom string in the GMT format, *e.g.*, `'Tue, 21 Sep 1999 14:46:36 GMT'`. Otherwise, the status code `WWW_EXPIRED_DOC` is returned.

10.2.3 Response Parameters

The `ResponseParams` argument is a list of terms returned by the `libwww_request` call. It contains two kinds of information: header information and sub-request information. The header information consists of terms like: `header('Content-Type', 'text/html')`, `header('Server',`

'Netscape-Enterprise/3.6 SP2'), etc., as defined by the HTTP protocol (`header/2` is a functor and its arguments are atoms). The sub-request information consists of terms of the form: `subrequest('http://www.foo.org/test/file.html',-401)`. It indicates that during processing of the current request, it was necessary to access another page, `http://www.foo.org/test/file.html`, but the server responded with the error code -401 (authentication error). Such sub-requests might be spawned during XML parsing.

10.2.4   Result of a Libwww Call

The type of the `Result` returned by a `libwww_request` call depends on the request type.

Page fetching.   In case of `fetch`, `Result` is an atom or a list of characters (depending on whether `URL` was specified as an atom or a list of characters), or it might be an unbound variable in case of an error. For `header` requests, `Result` is always an unbound variable.

HTML and XML parsing.   For `htmlparse` and `xmlparse`, `Result` is a variable in case of an error and a complex term otherwise. In the latter case, it is a list of the form `[elt1,...,elt_n]`, where each `elt_i` is of the form:

```
elt(tag, [attval(attrname,value),...], [elt1',...,elt'_m])
```

The second argument here represents the list of attribute-value pairs. In HTML, some attributes, like `checked`, can be binary, in which case the corresponding value will be unbound. The third argument represents HTML or XML elements that are within the scope of `tag`. These elements have the same syntax as the parent element: `elt(tag',attrs,sub-elements)`. If a tag has no attributes or if it does not have sub-elements, the corresponding lists will be empty. One special tag, `pcdata`, is introduced to represent pieces of text that appear in the document. This tag is our own creation—neither HTML nor XML use tags to represent text. One important difference between `pcdata` and other tags is that the third argument in `elt(pcdata,...,...)` is an atom or a list of characters, not a list (unlike other tags). If URL was specified as an atom, then the third argument of the `pcdata`-element is an atom as well. If URL is a character list, then so is the corresponding argument in the `pcdata`-element.

RDF parsing.   The form of the `Result` returned by an `rdfparse` call is a list of terms of the form:
$$[\texttt{triple}_1, \texttt{triple}_2, ..., \texttt{triple}_n]$$
where each triple is of the form:

$$\texttt{rdftriple}(\texttt{predicate}, \texttt{subject}, \texttt{object})$$

which is the usual representation of RDF elements.

10.2.5 Status Code

Finally, `Status` is bound to an integer that represents the return code from the HTTP request. A complete list of return codes is given in `XSB/prolog_includes/http_errors.h`. If you need to refer to error codes in your Prolog application, it is advisable to use symbolic notation (*e.g.*, `HT_TIMEOUT`) rather than specific numbers (*e.g.*, -905). To make this happen, put the following lines at the top of your program:

```
:- compiler_options([xpp_on]).
#include "http_errors.h"
```

The Libwww package also includes a predicate, `http_liberr/3`, that is convenient for providing English language explanations to the errors:

```
:- import http_liberr/3 from usermod.
```

The first argument of this predicate is the status code returned by an earlier call to `Libwww_request`, the second is an English explanation of the status, and the last is an atom that describes the class of the error (*e.g.*, internal, server error, client error, etc.). For full details see `XSB/packages/libwww/http_liberr.P`. Note that the status code for a successful call is HT_LOADED (=200), not 0 or 1!

10.3 Example

Here is a complete example:

```
| ?- libwww_request([xmlparse('http://public.org/test/simple1.xml',[timeout(4)],
                             P,Y,Z)]),
     http_liberr(Z,Explanation,Class).

P = [header(Content-Type,text/html),
     subrequest('http://public.org/secret/001.ent',-401)]
Y = [elt(doc,[],[elt(pcdata,[],' '),
                elt(foo,[attval(att1,123),attval(att2,ppp)],
                        [elt(pcdata,[],'Test1 '),elt(pcdata,[],' Test2 '),
                         elt(pcdata,[],adsdd),elt(pcdata,[], )]),
                elt(pcdata,[],  ),
                elt(a,[],[elt(pcdata,[], aaaaaaaaaa),
                elt(pcdata,[],' '),
                elt(b,[attval(att,1)],[]),
                elt(pcdata,[],' '),
                elt(c,[attval(att,2)],[]),
                elt(pcdata,[],' '),
                elt(d,[],[elt(pcdata,[],dddddddd),
                elt(f,[],[elt(pcdata,[],kkkkkkk)]),
```

```
                                   elt(pcdata,[],abc)]),elt(pcdata,[],' ')])])]
        Z = 200
        Explanation = 'OK'
        Class = 'success'
```

The above is a successful (because of the return code 200) request to parse an XML page. This page apparently had a reference to an external entity that was located in a protected domain. Since we did not supply authentication information, the call returned authentication failure for that subrequest (as indicated by the term `subrequest('http://public.org/secret/001.ent',-401)` in the fourth argument).

10.4   Special Notes about Parsing XML

XML documents can have macros (called *entities*), which are defined in the DTD and their occurrences are expanded according to the definition. XML specification allows entities to refer to an external document. When this happens, the libwww package must issue a subrequest to fetch the document containing the macro expansion. Such a subrequest can fail for a variety of reasons (usually network related). Status code of each subrequest is recorded in the third argument of the corresponding request term (that is passed to `libwww_request/1`), as shown in an earlier example. If the status indicates an error, the corresponding place in the request result (the fourth argument) will contain a term of the form `unexpanded_entity(Url,Status)`. For instance:

```
 ?- libwww_request([xml('http://localhost/simple.xml',
        [if_modified_since('Tue, 21 Sep 1999 14:46:36 GMT')],
        P,R,S)]).


P = [subrequest(http://localhost/001.ent,-404),
     subrequest(http://localhost/sub/001.ent,200),
     header(Content-Type,text/xml),header(Content-Length,505),
     header(Accept-Ranges,bytes),header(ETag,"5089-1f9-38f012dc),
     header(Last-Modified,Sun, 09 Apr 2000 05:19:24 GMT),
     header(Server,Apache/1.3.9 (Unix)  (Red Hat/Linux)),
     header(Date,Tue, 11 Apr 2000 05:01:55 GMT)]
R = [elt(doc,[],[elt(foo,[attval(att1,123),attval(att2,abc)],
                     [elt(pcdata,[],test),
                      unexpanded_entity(http://localhost/001.ent,-404),
                      elt(pcdata,[], Test1 ),
                      elt(pcdata,[], Test2 )]),
                      elt(pcdata,[],  )])]
S = 200
```

Another observation is that due to a bug in the Libwww library, requests to fetch local files (documents with URLs of the form `file://abc`) may take a long time or even abort, if the request requires spawning subrequests to fetch other local files (for instance, a request to parse a local XML

document that contains external references to local files). However, this is not a serious problem in practice, since the Web interface is rarely used to access local files.

Chapter 11

Constraint Packages

Constraint packages are an important part of modern logic programming, but approaches to constraints differ both in their semantics and in their implementation. At a semantic level, *Constraint Logic Programming* associates constraints with logical variables, and attempts to determine solutions that are inconsistent with or entailed by those constaints. At an implementational level, the constraints can either be manipulated by accessing attributed variables or by adding *constraint handling rules* to a program. The former approach of attributed variables can be much more efficient than constraint handling rules (which are themselves implemented through attributed variables) but are much more difficult to use than constraint handling rules. These variable-based approaches differ from that of *Answer Set Programming* in which a constraint problem is formulated as a set of rules, which are consistent if a stable model can be constructed for them.

XSB supports all of these approaches. Two packages based on attributed variables are presented in this chapter: CLP(R) and the `bounds` package, which provides a simple library for handling finite domains. XSB's CHR package is described in Chapter 12, and XSB's Answer Set Programming Package, `XASP` is described in Chapter 13.

Before describing the individual packages, we note that these packages can be freely used with variant tabling, the mechanisms for which handle attributed variables. However in Version 3.2, calling a predicate P that is tabled using call subsumption will raise an error if the call to P contains any constrained variables (attributed variables).

11.1 `clpr`: The CPL(R) package

The CLP(R) library supports solutions of linear equations and inequalities over the real numbers and the lazy treatment of nonlinear equations [1]. In displaying sets of equations and disequations, the library removes redundancies, performs projections, and provides for linear optimization. The goal of the XSB port is to provide the same CLP(R) functionality as in other platforms, but also to allow constraints to be used by tabled predicates. This section provides a general introduction

---

[1] The CLP(R) package is based on the clpqr package included in SWI Prolog version 5.6.49. This package was originally written by Christian Holzbaur and ported to SWI by Leslie De Konick. Terrance Swift ported the package to XSB and and wrote this XSB manual section.

to the CLP(R) functionality available in XSB, for further information on the API described in Section 11.1.1 see `http://www.ai.univie.ac.at/clpqr`, or the Sicstus Prolog manual (the CLP(R) library should behave similarly on XSB and Sicstus at the level of this API).

The `clpr` package may be loaded by the command `[clpr]`. Loading the package imports exported predicates from the various files in the `clpr` package into `usermod` (see Volume 1, Section 3.3) so that they may be used in the interpreter. Modules that use the exported predicates need to explicitly import them from the files in which they are defined (e.g. `bv`, as shown below).

XSB's tabling engine supports the use of attributed variables (Section 1.2), which in turn have been used to port real constraints to XSB under the CLP(R) library of Christian Holzbauer [13]. Constraint equations are represented using the Prolog syntax for evaluable functions (Volume 1, Section 6.2.1). Formally:

ConstraintSet –>	*C*	*C* , *C*	
C –>	*Expr* `=:=` *Expr*	equation
	| *Expr* `=` *Expr*	equation
	| *Expr* `<` *Expr*	strict inequation
	| *Expr* `>` *Expr*	strict inequation
	| *Expr* `=<` *Expr*	nonstrict inequation
	| *Expr* `>=` *Expr*	nonstrict inequation
	| *Expr* `=/=` *Expr*	disequation
Expr –>	*variable*	Prolog variable
	| *number*	floating point number
	| `+` *Expr*	
	| `-` *Expr*	
	| *Expr* `+` *Expr*	
	| *Expr* `-` *Expr*	
	| *Expr* `*` *Expr*	
	| *Expr* `/` *Expr*	
	| `abs(`*Expr*`)`	
	| `sin(`*Expr*`)`	
	| `cos(`*Expr*`)`	
	| `tan(`*Expr*`)`	
	| `pow(`*Expr*`,`*Expr*`)`	raise to the power
	| `exp(`*Expr*`,`*Expr*`)`	raise to the power
	| `min(`*Expr*`,`*Expr*`)`	minimum of two expressions
	| `max(`*Expr*`,`*Expr*`)`	maximum of two expressions
	| `#(`*Expr)*	symbolic numerical constants

11.1.1 The CLP(R) API

From the command line, it is usually easiest to load the `clpr` package and call the predicates below directly from `usermod` (the module implicitly used by the command line). However, when calling

```
:- import {}/1 from clpr.

root(N, R) :-
root(N, 1, R).
root(0, S, R) :- !, S=R.
root(N, S, R) :-
N1 is N-1,
{ S1 = S/2 + 1/S },
root(N1, S1, R).
```

Figure 11.1: Example of a file with a CLP(R) predicate

any of these predicates from compiled code, they must be explicitly imported from their modules
(e.g. {} must be explicitly imported from `clpr`). Figure 11.1.1 shows an example of how this is
done. '

{+Constraints}   module: clpr
     When the CLP(R) package is loaded, inclusion of equations in braces ({}) adds `Constraints`
     to the constraint store where they are checked for satisfiability.

     Example:

```
| ?- [clpr].
[clpr loaded]
[itf loaded]
[dump loaded]
[bv_r loaded]
[nf_r loaded]

yes

| ?- {X = Y+1, Y = 3*X}.

X = -0.5000
Y = -1.5000;

yes
```

     Error Cases

   - `Constraints` is not instantiated
       - `instantiation_error`
   - `Constraints` is not an equation, an inequation or a disequation
       - `domain_error('constraint relation',Rel)`
   - `Constraints` contains an expression `Expr` that is not a numeric expression
       - `domain_error('numeric expression',Expr)`

```
entailed(+Constraint)                                         module:  clpr
```
   Succeeds if `Constraint` is logically implied by the current constraint store. `entailed/1` does not change the constraint store.

   Example:

```
| ?- {A =< 4},entailed(A =\= 5).
 { A =< 4.0000 }

yes
```

   Error Cases

   - `Constraints` is not instantiated
     - `instantiation_error`
   - `Constraints` is not an equation, an inequation or a disequation
     - `domain_error('constraint relation',Rel)`

```
inf(+Expr,-Val)                                                          clpr
sup(+Expr,-Val)                                                          clpr
minimize(Expr)                                                           clpr
maximize(Expr)                                                 module:  clpr
```
   These four related predicates provide various mechanisms to compute the maximum and minimum of expressions over variables in a constraint store. In the case where the expression is not bounded from above over the reals `sup/2` and `maximize/1` will fail; similarly if the expression is not bounded from below `inf/2` and `minimize/1` will fail.

   Examples:

```
| ?- {X = 2*Y,Y >= 7},inf(X,F).
 { X >= 14.0000 }
 { Y = 0.5000 * X }

X = _h8841
Y = _h9506
F = 14.0000

| ?- {X = 2*Y,Y >= 7},minimize(X).
X = 14.0000
Y = 7.0000

| ?- {X = 2*Y,Y =< 7},maximize(X-2).

X = 14.0000
Y = 7.0000

| ?- {X = 2*Y,Y =< 7},sup(X-2,Z).
 { X =< 14.0000 }
 { Y = 0.5000 * X }
```

```
    X = _h8975
    Y = _h9640
    Z = 12.0000


    yes
    | ?- {X = 2*Y,Y =< 7},maximize(X-2).


    X = 14.0000
    Y = 7.0000


    yes
```

inf(+Expr,-Val, +Vector, -Vertex)   clpr
sup(+Expr,-Val, +Vector, -Vertex)   module: clpr
   These predicates work like `inf/2` and `sup/2` with the following addition. `Vector` is a list
   of Variables, and for each variable V in `Vector`, the value of V at the extremal point `Val` is
   returned in corresponding position in the list `Vertex`.

   Example:

```
    | ?= { 2*X+Y =< 16, X+2*Y =< 11,X+3*Y =< 15, Z = 30*X+50*Y},
        sup(Z, Sup, [X,Y], Vertex).
     { X + 3.0000 * Y =< 15.0000 }
     { X + 0.5000 * Y =< 8.0000 }
     { X + 2.0000 * Y =< 11.0000 }
     { Z = 30.0000 * X + 50.0000 * Y }


    X = _h816
    Y = _h869
    Z = _h2588
    Sup = 310.0000
    Vertex = [7.0000,2.0000]
```

bb_inf(+IntegerList,+Expr,-Inf,-Vertex, +Eps)   module: clpr
   Works like `inf/2` in `Expr` but assumes that all the variables in `IntegerList` have inte-
   gral values. `Eps` is a positive number between 0 and 0.5 that specifies how close an ele-
   ment of `IntegerList` must be to an integer to be considered integral – i.e. for such an X,
   `abs(round(X) - X) < Eps`. Upon success, `Vertex` is instantiated to the integral values of
   all variables in `IntegerList`. `bb_inf/5` works properly for non-strict inequalities only.

   Example:

```
    | ?- {X > Y + Z,Y > 1, Z > 1},bb_inf([Y,Z],X,Inf,Vertex,0).
     { Z > 1.0000 }
     { Y > 1.0000 }
     { X - Y - Z > 0.0000 }


    X = _h14286
    Y = _h10914
    Z = _h13553
```

```
Inf = 4.0000
Vertex = [2.0000,2.0000]

yes
```

Error Cases

- `IntegerList` is not instantiated

    - `instantiation_error`

`bb_inf(+IntegerList,+Expr,-Inf)`   module: clpr

Works like `bb_inf/5`, but with the neighborhood, `Eps`, set to 0.001.

Example

```
|?- {X >= Y+Z, Y > 1, Z > 1}, bb_inf([Y,Z],X,Inf)
 { Z > 1.0000 }
 { Y > 1.0000 }
 { X - Y - Z >= 0.0000 }

X = _h14289
Y = _h10913
Z = _h13556
Inf = 4.

yes
```

`dump(+Variables,+NewVars,-CodedVars`   module: dump

For a list of variables `Variables` and a list of variable names `NewVars`, returns in `CodedVars` the constraints on the variables, without affecting the constraint store.

Example:

```
| ?- {X > Y+1, Y > 2},
     dump([X,Y], [x,y], CS).
 { Y > 2.0000 }
 { X - Y > 1.0000 }

X = _h17748
Y = _h17139
CS = [y > 2.0000,x - y > 1.0000];
```

Error Cases

- `Variables` is not instantiated to a list of variables

    - `instantiation_error`

`projecting_assert(+Clause)`   module: dump

In XSB, when a subgoal is tabled, the tabling system automatically determines the relevant projected constraints for an answer and copies them into and out of a table. However,

when a clause with constrained variables is asserted, this predicate must be used rather than `assert/1` in order to project the relevant constraints. This predicate works with either standard or trie-indexed dynamic code.

Example:

```
| ?- {X > 3},projecting_assert(q(X)).
 { X > 3.0000 }

X = _h396

yes
| ?- listing(q/1).
q(A) :-
    clpr : {A > 3.0000}.

yes
| ?- q(X),entailed(X > 2).
 { X > 3.0000 }

X = _h358

yes
| ?- q(X),entailed(X > 4).

no
```

11.2 The `bounds` Package

Version 3.2 of XSB does not support a full-fledged CLP(FD) package. However it does support a simplified package that maintains an upper and lower bound for logical variables. `bounds` can thus be used for simple constraint problems in the style of finite domains, as long as these problems that do not rely on too heavily on propagation of information about constraint domains [2]

Perhaps the simplest way to explain the functionality of `bounds` is by example. The query

```
|?- X in 1..2,X #> 1.
```

first indicates via `X in 1..2` that the lower bound of `X` is `1` and the higher bound `2`, and then constraints `X`, which is not yet bound, to be greater than `1`. Applying this latter constraint to `X` forces the lower bound to equal the upper bound, instantiating `X`, so that the answer to this query is `X = 2`.

Next, consider the slightly more complex query

```
|?- X in 1..3,Y in 1..3,Z in 1..3,all_different([X,Y,Z]),X = 1, Y = 2.
```

---

[2]The `bounds` package was written by Tom Schrijvers, and ported to XSB from SWI Prolog version 5.6.49 by Terrance Swift, who also wrote this manual section.

all_different/3 constraints X, Y and Z each to be different, whatever their values may be. Accordingly, this constraint together with the bound restrictions, implies that instantiating X and Y also causes the instantiation of Z. In a similar manner, the query

```
|?- X in 1..3,Y in 1..3,Z in 1..3,sum([X,Y,Z],#=,9),
```

onstrains the sum of the three variables to equal 9 – and in this case assigns them a concrete value due to their domain restrictions.

In many constraint problems, it does not suffice to know whether a set of constraints is satisfiable; rather, concrete values may be needed that satisfy all constraints. One way to produce such values is through the predicate labelling/2

```
|?- X in 1..5,Y in 1..5,X #< Y,labeling([max(X)],[X,Y]))
```

In this query, it is specified that X and Y are both to be instantiated not just by any element of their domains, but by a value that assigns X to be the maximal element consistent with the constraints. Accordingly X is instantiated to 4 and Y to 5.

Because constraints in bounds are based on attributed variables which are handled by XSB's variant tabling mechanisms, constrained variables can be freely used with variant tabling as the folowing fragment shows:

```
table_test(X):- X in 2..3,p(X).

:- table p/1.
p(X):- X in 1..2.

?- table_test(Y).

Y = 2
```

For a more elaborate example, we turn to the *SEND MORE MONEY* example, , in which the problem is to assign numbers to each of the letters *S,E,N,D,M,O,R,Y* so that the number *SEND* plus the number *MORE* equals the number *MONEY*. Borrowing a solution from the SWI manual [24], the bounds package solves this problem as:

```
send([[S,E,N,D], [M,O,R,E], [M,O,N,E,Y]]) :-
    Digits = [S,E,N,D,M,O,R,Y],
    Carries = [C1,C2,C3,C4],
    Digits in 0..9,
    Carries in 0..1,
    M #= C4,
    O + 10 * C4 #= M + S + C3,
    N + 10 * C3 #= O + E + C2,
    E + 10 * C2 #= R + N + C1,
    Y + 10 * C1 #= E + D,
    M #>= 1,
    S #>= 1,
    all_different(Digits),
    label(Digits).
```

In many cases, it may be useful to test whether a given constraint is true or false. This can be done by unifying a variable with the truth value of a given constraint – i.e. by *reifying* the constraint. As an example, the query

```
|?-  X in 1..10, Y in 1..10,Z in 0..1,X #< Y, X #= Y #<=> Z,label([Z]).
```

sets the bounded variable Z to the truth value of X #= Y, or 0 [3].

A reader familiar with the finite domain library of Sicstus [16] will have noticed that the syntax of bounds is consistent with that library. It is important to note however, bounds maintains only the upper and lower bounds of a variables as its attributes, (along, of course with constraints on those variables) rather than an explicit vector of permissable values. As a result, bounds may not be suitable for large or complex constraint problems.

11.2.1 The bounds API

Note that bounds does not perform error checking, but instead relies on the error checking of lower-level comparison and arithmetic operators.

in(-Variable,+Bound)   bounds
> Adds the constraint Bound to Variable, where Bound should be of the form Low..High, with Low and High instantiated to integers. This constraint ensures that any value of Variable must be greater than or equal to Low and less than or equal to High. Unlike some finite-domain constraint systems, it does *not* materialize a vector of currently allowable values for Variable.
>
> Variables that have not had their domains explicitly constrained are considered to be in the range min_integer..max_integer.

#>(Expr1,Expr2)   bounds
#<(Expr1,Expr2)   bounds
#>=(Expr1,Expr2)   bounds
#=<(Expr1,Expr2)   bounds
#=(Expr1,Expr2)   bounds
#=(Expr1,Expr2)   bounds
> Ensures that a given relation holds between Expr1 and Expr2. Within these constraints, expressions may contain the functions +/2, -/2, */2, +/2, +/2, +/2, mod/2, and abs/1 in addition to integers and variables.

#<=>(Const1,Const2)   bounds
#=>(Const1,Const2)   bounds
#<=(Const1,Const2)   bounds
> Constrains the truth-value of Const1 to have the speficied logical relation ("iff", "only-if" or "if") to Const2, where Const1 and Const2 have one of the six relational operators above.

---

[3]The current version of the bounds package does not always seem to propagate entailment into the values of reified variables.

all_different(+VarList) bounds

> VarList must be a list of variables: constrains all variables in VarList to have different values.

sum(VarList,Op,?Value) bounds

> VarList must be a list of variables and Value an integer or variable: constrains the sum of all variables in VarList to have the relation Op to Value (see preceding example).

labeling(+Opts,+VarList bounds

> This predicate succeeds if it can assign a value to each variable in VarList such that no constraint is violated. Note that assigning a value to each constrained variable is equivalent to deriving a solution that satisfies all constraints on the variables, which may be intractible depending on the constraints. Opts allows some control over how value assignment is performed in deriving the solution.
>
> - **leftmost** Assigns values to variables in the order in which they occur. For example the query:
>
>   ```
>   |?- X in 1..4,Y in 1..3,X #< Y,labeling([leftmost],[X,Y]),writeln([X,Y]),fail.
>   [1,2]
>   [1,3]
>   [2,3]
>
>   no
>   ```
>
>   instantiates X and Y to all values that satisfy their constraints, and does so by considering each value in the domain of X, checking whether it violates any constraints, then considering each value of Y and checking whether it violates any constraints.
>
> - **ff** This "first-fail" strategy assignes values to variables based on the size of their domains, from smallest to largest. By adopting this strategy, it is possible to perform a smaller search for a satisfiable solution because the most constrained variables may be considered first (though the bounds of the variable are checked rather than a vector of allowable values).
>
> - **min** and **max** This strategy labels variables in the order of their minimal lower bound or maximal upper bound.
>
> - **min(Expr)** and **max(Expr)** This strategy labels the variables so that their assignment causes Expr to have a minimal or maximal value. Consider for example how these strategies would affect the labelling of the preceding query:
>
>   ```
>   |?- X in 1..4,Y in 1..3,X #< Y,labeling([min(Y)],[X,Y]),writeln([X,Y]),fail.
>   [1,2]
>
>   no
>   |?- X in 1..4,Y in 1..3,X #< Y,labeling([max(X)],[X,Y]),writeln([X,Y]),fail.
>   [2,3]
>
>   no
>   ```

label(+VarList) bounds

> Shorthand for labeling([leftmost],+VarList).

`indomain(?Var)` bounds

Unifies `Var` with an element of its domain, and upon sucessive backttrakcing, with all other elements of its domain.

`serialized(+BeginList,+Durations` bounds

`serialized/2` can be useful for scheduling problems. As input it takes a list of variables or integers representing the beginnings of temporal events, along with a list of non-negative intergers indicating the duration of each event in `BeginList`. The effect of this predicate is to constrain each of the events in `BeginList` to have a start time such that their durations do not overlap. As an example, consier the query

```
|?- X in 1..10, Y in 1..10, serialized([X,Y],[8,1]),label([X,Y]),writeln((X,Y)),fail.
```

In this query event `X` is taken to have duration of 8 units, while event `Y` is taken to have duration of 1 unit. Executing this query will instantiate `X` and `Y` to many different values, such as `(1,9)`, `(1,10)`, and `(2,10)` where `X` is less than `Y`, but also `(10,1)`, `(10,2)` and many others where `Y` is less than `X`. Refining the query as

```
X in 1..10, Y in 1..10, serialized([X,Y],[8,1]),X #< Y,label([X,Y]),writeln((X,Y)),fail.
```

removes all solutions where `Y` is less than `X`.

`lex_chain(+List)` bounds

`lex_chain/1` takes as input a list of lists of variables and integers, and enforces the constraint that each element in a given list is less than or equal to the elements in all succeeding lists. As an example, consider the query

```
|?- X in 1..3,Y in 1..3,lex_chain([[X],[2],[Y]]),label([X,Y]),writeln([X,Y]),fail.
[1,2]
[1,3]
[2,2]
[2,3]
```

`lex_chain/1` ensures that `X` is less than or equal to 2 which is less than or equal to `Y`.

Chapter 12

Constraint Handling Rules

12.1 Introduction

Constraint Handling Rules (CHR) is a committed-choice bottom-up language embedded in XSB. It is designed for writing constraint solvers and is particularly useful for providing application-specific constraints. It has been used in many kinds of applications, like scheduling, model checking, abduction, type checking among many others.

CHR has previously been implemented in other Prolog systems (SICStus, Eclipse, Yap, hProlog), Haskell and Java. The XSB CHR system is based on the hProlog CHR system.

In this documentation we restrict ourselves to giving a short overview of CHR in general and mainly focus on XSB-specific elements. For a more thorough review of CHR we refer the reader to [11]. More background on CHR can be found at [10].

In Section 12.2 we present the syntax of CHR in XSB and explain informally its operational semantics. Next, Section 12.3 deals with practical issues of writing and compiling XSB programs containing CHR. Section 12.4 provides a few useful predicates to inspect the constraint store and Section 12.5 illustrates CHR with two example programs. How to combine CHR with tabled predicates is covered in Section 12.6. Finally, Section 12.7 concludes with a few practical guidelines for using CHR.

12.2 Syntax and Semantics

12.2.1 Syntax

The syntax of CHR rules in XSB is the following:

```
rules --> rule, rules.
rules --> [].

rule --> name, actual_rule, pragma, [atom('.')].
```

```
name --> xsb_atom, [atom('@')].
name --> [].

actual_rule --> simplification_rule.
actual_rule --> propagation_rule.
actual_rule --> simpagation_rule.

simplification_rule --> constraints, [atom('<=>')], guard, body.
propagation_rule --> constraints, [atom('==>')], guard, body.
simpagation_rule --> constraints, [atom('\')], constraints, [atom('<=>')],
                     guard, body.

constraints --> constraint, constraint_id.
constraints --> constraint, [atom(',')], constraints.

constraint --> xsb_compound_term.

constraint_id --> [].
constraint_id --> [atom('#')], xsb_variable.

guard --> [].
guard --> xsb_goal, [atom('|')].

body --> xsb_goal.

pragma --> [].
pragma --> [atom('pragma')], actual_pragmas.

actual_pragmas --> actual_pragma.
actual_pragmas --> actual_pragma, [atom(',')], actual_pragmas.

actual_pragma --> [atom('passive(')], xsb_variable, [atom(')')].
```

Additional syntax-related terminology:

- **head:** the constraints in an **actual_rule** before the arrow (either <=> or ==>)

12.2.2   Semantics

In this subsection the operational semantics of CHR in XSB are presented informally. They do not differ essentially from other CHR systems.

When a constraint is called, it is considered an active constraint and the system will try to apply the rules to it. Rules are tried and executed sequentially in the order they are written.

A rule is conceptually tried for an active constraint in the following way. The active constraint is matched with a constraint in the head of the rule. If more constraints appear in the head they are looked for among the suspended constraints, which are called passive constraints in this context. If the necessary passive constraints can be found and all match with the head of the rule and the guard of the rule succeeds, then the rule is committed and the body of the rule executed. If not all the necessary passive constraint can be found, the matching fails or the guard fails, then the body is not executed and the process of trying and executing simply continues with the following rules. If for a rule, there are multiple constraints in the head, the active constraint will try the rule sequentially multiple times, each time trying to match with another constraint.

This process ends either when the active constraint disappears, i.e. it is removed by some rule, or after the last rule has been processed. In the latter case the active constraint becomes suspended.

A suspended constraint is eligible as a passive constraint for an active constraint. The other way it may interact again with the rules, is when a variable appearing in the constraint becomes bound to either a non-variable or another variable involved in one or more constraints. In that case the constraint is triggered, i.e. it becomes an active constraint and all the rules are tried.

Rule Types   There are three different kinds of rules, each with their specific semantics:

- `simplification:`

  The simplification rule removes the constraints in its head and calls its body.

- `propagation:`

  The propagation rule calls its body exactly once for the constraints in its head.

- `simpagation:`

  The simpagation rule removes the constraints in its head after the \ and then calls its body. It is an optimization of simplification rules of the form:

$$constraints_1, constraints_2 <=> constraints_1, body$$

  Namely, in the simpagation form:

$$constraints_1 \backslash constraints_2 <=> body$$

  The $constraints_1$ constraints are not called in the body.

Rule Names   Naming a rule is optional and has no semantical meaning. It only functions as documentation for the programmer.

Pragmas   The semantics of the pragmas are:

- **passive/1:** the constraint in the head of a rule with the identifier specified by the **passive/1** pragma can only act as a passive constraint in that rule.

Additional pragmas may be released in the future.

12.3   CHR in XSB Programs

12.3.1   Embedding in XSB Programs

Since chr is an XSB package, it must be explicitly loaded before being used.

```
?- [chr].
```

CHR rules are written in a tt .chr file. They should be preceded by a declaration of the constraints used:

```
:- constraints ConstraintSpec1, ConstraintSpec2, ...
```

where each `ConstraintSpec` is a functor description of the form name/arity pair. Ordinary code may be freely written between the CHR rules.

The CHR constraints defined in a particular .chr file are associated with a CHR module. The CHR module name can be any atom. The default module is `user`. A different module name can be declared as follows:

```
:- chr_module(modulename).
```

One should never load different files with the same CHR module name.

12.3.2   Compilation

Files containing CHR rules are required to have a `.chr` extension, and their compilation has two steps. First the `.chr` file is preprocessed into a `.P` file containing XSB code. This `.P` file can then be loaded in the XSB emulator and used normally.

load_chr(File)   chr_pp
>   load_chr/1 takes as input a file name whose extension is either `.chr` or that has no extension. It preprocesses `File` if the times of the CHR rule file is newer than that of the corresponding Prolog file, and then consults the Prolog file.

preprocess(File,PFile)   chr_pp
>   preprocess/2 takes as input a file name whose extension is either `.chr` or that has no extension. It preprocesses `File` if the times of the CHR rule file is newer than that of the corresponding Prolog file, but does not consult the Prolog file.

12.4   Useful Predicates

The `chr` module contains several useful predicates that allow inspecting and printing the content of the constraint store.

`show_store(+Mod)`   chr
>   Prints all suspended constraints of module `Mod` to the standard output.

`suspended_chr_constraints(+Mod,-List)`   chr
>   Returns the list of all suspended CHR constraints of the given module.

12.5   Examples

Here are two example constraint solvers written in CHR.

- The program below defines a solver with one constraint, `leq/2`, which is a less-than-or-equal constraint.

  ```
  :- chr_module(leq).

  :- export cycle/3.

  :- import length/2 from basics.

  :- constraints leq/2.
  reflexivity  @ leq(X,X) <=> true.
  antisymmetry @ leq(X,Y), leq(Y,X) <=> X = Y.
  idempotence  @ leq(X,Y) \ leq(X,Y) <=> true.
  transitivity @ leq(X,Y), leq(Y,Z) ==> leq(X,Z).

  cycle(X,Y,Z):-
          leq(X,Y),
          leq(Y,Z),
          leq(Z,X).
  ```

- The program below implements a simple finite domain constraint solver.

  ```
  :- chr_module(dom).

  :- import member/2 from basics.

  :- constraints dom/2.
  ```

```
    dom(X,[]) <=> fail.
    dom(X,[Y]) <=> X = Y.
    dom(X,L1), dom(X,L2) <=> intersection(L1,L2,L3), dom(X,L3).

    intersection([],_,[]).
    intersection([H|T],L2,[H|L3]) :-
            member(H,L2), !,
            intersection(T,L2,L3).
    intersection([_|T],L2,L3) :-
            intersection(T,L2,L3).
```

These and more examples can be found in the `examples/chr/` folder accompanying this XSB release.

12.6   CHR and Tabling

The advantage of CHR in XSB over other Prolog systems, is that CHR can be combined with tabling. Hence part of the constraint solving can be performed once and reused many times. This has already shown to be useful for applications of model checking with constraints.

However the use of CHR constraints is slightly more complicated for tabled predicates. This section covers how exactly to write a tabled predicate that has one or more arguments that also appear as arguments in suspended constraints. In the current release the CHR-related parts of the tabled predicates have to be written by hand. In a future release this may be substituted by an automatic transformation.

12.6.1   General Issues and Principles

The general issue is how call constraints should be passed in to the tabled predicate and how answer constraints are passed out of the predicate. Additionally, in some cases care has to be taken not to generate infinite programs.

The recommended approach is to write the desired tabled predicate as if no additional code is required to integrate it with CHR. Next transform the tabled predicate to take into account the combination of tabling and CHR. Currently this transformation step has to be done by hand. In the future we hope to replace this hand coding with programmer declarations that guide automated transformations.

Hence we depart from an ordinary tabled predicate, say `p/1`:

```
:- table p/1.

p(X) :-
    ... /* original body of p/1 */.
```

In the following we will present several transformations or extensions of this code to achieve a particular behavior. At least the transformation discussed in subsection 12.6.2 should be applied to obtain a working integration of CHR and tabling. Further extensions are optional.

12.6.2   Call Abstraction

Currently only one type of call abstraction is supported: full constraint abstraction, i.e. all constraints on variables in the call should be removed. The technique to accomplish this is to replace all variables in the call that have constraints on them with fresh variables. After the call, the original variables should be unified with the new ones.

In addition, the call environment constraint store should be replaced with an empty constraint store before the call and on return the answer store should be merged back into the call environment constraint store.

The previously mentioned tabled predicate `p/1` should be transformed to:

```
:- import merge_answer_store/1,
          get_chr_store/1,
          set_chr_store/1,
          get_chr_answer_store/2
   from chr.

:- table tabled_p/2.

p(X) :-
        tabled_p(X1,AnswerStore),
        merge_answer_store(AnswerStore),
        X1 = X.

tabled_p(X,AnswerStore) :-
        get_chr_store(CallStore),
        set_chr_store(_EmptyStore)
        orig_p(X),
        get_chr_answer_store(chrmod,AnswerStore),
        set_chr_store(CallStore).

orig_p(X) :-
   ... /* original body of p/1 */.
```

This example shows how to table the CHR constraints of a single CHR module `chrmod`. If multiple CHR modules are involved, one should add similar arguments for the other modules.

12.6.3   Answer Projection

To get rid of irrelevant constraints, most notably on local variables, the answer constraint store should in some cases be projected on the variables in the call. This is particularly important for

programs where otherwise an infinite number of answers with ever growing answer constraint stores could be generated.

The current technique of projection is to provide an additional `project/1` constraint to the CHR solver definition. The argument of this constraint is the list of variables to project on. Appropriate CHR rules should be written to describe the interaction of this `project/1` constraint with other constraints in the store. An additional rule should take care of removing the `project/1` constraint after all such interaction.

The `project/1` constraint should be posed before returning from the tabled predicate.

If this approach is not satisfactory or powerful enough to implement the desired projection operation, you should resort to manipulating the underlying constraint store representation. Contact the maintainer of XSB's CHR system for assistance.

Example   Take for example a predicate `p/1` with a less than or equal constraint `leq/2` on variables and integers. The predicate `p/1` has local variables, but when `p` returns we are not interested in any constraints involving local variables. Hence we project on the argument of `p/1` with a project constraint as follows:

```
:- import memberchk/2 from lists.

:- import merge_answer_store/1,
          get_chr_store/1,
          set_chr_store/1,
          get_chr_answer_store/2
   from chr.

:- table tabled_p/2.

:- constraints leq/2, project/1.

... /* other CHR rules */
project(L) \ leq(X,Y) <=>
        ( var(X), \+ memberchk(X,L)
        ; var(Y), \+ memberchk(Y,L)
        ) | true.

project(_) <=> true.

p(X) :-
        tabled_p(X1,AnswerStore),
        merge_answer_store(AnswerStore),
        X1 = X.

tabled_p(X,AnswerStore) :-
        get_chr_store(CallStore),
        set_chr_store(_EmptyStore)
        orig_p(X),
        project([X]),
```

```
        get_chr_answer_store(chrmod,AnswerStore),
        set_chr_store(CallStore).

orig_p(X) :-
   ... /* original body of p/1 */.
```

The example in the following subsection shows projection in a full application.

12.6.4   Answer Combination

Sometimes it is desirable to combine different answers to a tabled predicate into one single answer
or a subset of answers. Especially when otherwise there would be an infinite number of answers. If
the answers are expressed as constraints on some arguments and the logic of combining is encoded
as CHR rules, answers can be combined by merging the respective answer constraint stores.

Another case where this is useful is when optimization is desired. If the answer to a predicate
represents a valid solution, but an optimal solution is desired, the answer should be represented
as constraints on arguments. By combining the answer constraints, only the most constrained, or
optimal, answer is kept.

Example   An example of a program that combines answers for both termination and optimisation
is the shortest path program below:

```
:- chr_module(path).

:- import length/2 from lists.

:- import merge_chr_answer_store/1,
           get_chr_store/1,
           set_chr_store/1,
           get_chr_answer_store/2
   from chr.

breg_retskel(A,B,C,D) :- '_$builtin'(154).

:- constraints geq/2, plus/3, project/1.

geq(X,N) \ geq(X,M) <=> number(N), number(M), N =< M | true.

reflexivity  @ geq(X,X) <=> true.
antisymmetry @ geq(X,Y), geq(Y,X) <=> X = Y.
idempotence  @ geq(X,Y) \ geq(X,Y) <=> true.
transitivity @ geq(X,Y), geq(Y,Z) ==> var(Y) | geq(X,Z).

plus(A,B,C) <=> number(A), number(B) | C is A + B.
plus(A,B,C), geq(A,A1) ==> plus(A1,B,C1), geq(C,C1).
plus(A,B,C), geq(B,B1) ==> plus(A,B1,C1), geq(C,C1).
```

```
project(X) \ plus(_,_,_) # ID <=> true pragma passive(ID).
project(X) \ geq(Y,Z) # ID <=> (Y \== X ; var(Z) )| true pragma passive(ID).
project(_) <=> true.


path(X,Y,C) :-
tabled_path(X,Y,C1,AS),
merge_chr_answer_store(AS),
C = C1.


:- table tabled_path/4.

tabled_path(X,Y,C,AS) :-
'_$savecp'(Breg),
breg_retskel(Breg,4,Skel,Cs),
copy_term(p(X,Y,C,AS,Skel),p(OldX,OldY,OldC,OldAS,OldSkel)),
        get_chr_store(GS),
set_chr_store(_GS1),
orig_path(X,Y,C),
        project(C),
( get_returns(Cs,OldSkel,Leaf),
  OldX == X, OldY == Y ->
            merge_chr_answer_store(OldAS),
            C = OldC,
            get_chr_answer_store(path,MergedAS),
            sort(MergedAS,AS),
            ( AS = OldAs ->
                fail
            ;
                delete_return(Cs,Leaf)
            )
;
            get_chr_answer_store(path,UnsortedAS),
            sort(UnsortedAS,AS)
),
        set_chr_store(GS).

orig_path(X,Y,C) :- edge(X,Y,C1), geq(C,C1).
orig_path(X,Y,C) :- path(X,Z,C2), edge(Z,Y,C1), plus(C1,C2,C0), geq(C,C0).

edge(a,b,1).
edge(b,a,1).
edge(b,c,1).
edge(a,c,3).
edge(c,a,1).
```

The predicate `orig_path/3` specifies a possible path between two nodes in a graph. In `tabled_path/4` multiple possible paths are combined together into a single path with the shortest distance. Hence the tabling of the predicate will reject new answers that have a worse distance and will replace the old answer when a better answer is found. The final answer gives the optimal solution, the shortest path. It is also necessary for termination to keep only the best answer. When cycles appear in the

graph, paths with longer and longer distance could otherwise be put in the table, contributing to the generation of even longer paths. Failing for worse answers avoids this infinite build-up.

The predicate also includes a projection to remove constraints on local variables and only retain the bounds on the distance.

The sorting canonicalizes the answer stores, so that they can be compared.

12.6.5   Overview of Tabling-related Predicates

`merge_answer_store(+AnswerStore)`   chr
> Merges the given CHR answer store into the current global CHR constraint store.

`get_chr_store(-ConstraintStore)`   chr
> Returns the current global CHR constraint store.

`set_chr_store(?ConstraintStore)`   chr
> Set the current global CHR constraint store. If the argument is a fresh variable, the current global CHR constaint store is set to be an empty store.

`get_chr_answer_store(+Mod,-AnswerStore)`   chr
> Returns the part of the current global CHR constraint store of constraints in the specified CHR module, in the format of an answer store usable as a return argument of a tabled predicate.

12.7   Guidelines

In this section we cover several guidelines on how to use CHR to write constraint solvers and how to do so efficiently.

- **Set semantics:** The CHR system allows the presence of identical constraints, i.e. multiple constraints with the same functor, arity and arguments. For most constraint solvers, this is not desirable: it affects efficiency and possibly termination. Hence appropriate simpagation rules should be added of the form:

$$constraint \backslash constraint <=> true$$

- **Multi-headed rules:** Multi-headed rules are executed more efficiently when the constraints share one or more variables.

12.8   CHRd

An alternate implementation of CHR can be found in the CHRd package. The main objective of the CHRd package is to optimize processing of constraints in the environment where termination is guaranteed by the tabling engine, (and where termination benefits provided by the existing solver

are not critical). CHRd takes advantage of XSB's tabling to simplify CHR's underlying storage structures and solvers. Specifically, we entirely eliminate the thread-global constraint store in favor of a distributed one, realized as a collection of sets of constraints entirely associated with program variables. This decision limits the applicability of CHRd to a restricted class of CHR programs, refered to as direct-indexed CHR,in which all constraints in the head of a rule are connected by shared variables. Most CHR programs are direct-indexed, and other programs may be easily converted to fall into this class. Another advance of CHRd is its set-based semantics which removes the need to maintain the propagation history, thus allowing further simplicity in the representation of the constraints. The CHRd package itself is described in [19], and both the semantics of CHRd and the class of direct-indexed CHR are formally defined in [20].

Chapter 13

XASP: Answer Set Programming with XSB and Smodels

By Luis Castro, Terrance Swift, David S. Warren [1]

The term *Answer Set Programming (ASP)* describes a paradigm in which logic programs are interpreted using the (extended) stable model semantics. While the stable model semantics is quite elegant, it has radical differences from traditional program semantics based on Prolog. First, stable model semantics applies only to ground programs; second stable model semantics is not goal-oriented – determining whether a stable model is true in a program involves examining each clause in a program, regardless of whether the goal would depends on the clause in a traditional evaluation [2].

Despite (or perhaps because of) these differences, ASP has proven to be a useful paradigm for solving a variety of combinatorial programs. Indeed, determining a stable model for a logic program can be seen as an extension of the NP-complete problem of propositional satisfiability, so that satisfiability problems that can be naturally represented as logic programs can be solved using ASP.

The current generation of ASP systems are very efficient for determining whether a program has a stable model (analogous to whether the program, taken as a set of propositional axioms, is satisfiable). However, ASP systems have somewhat primitive file-based interfaces. XSB is a natural complement to ASP systems. Its basis in Prolog provides a procedural counterpart for ASP, as described in Chapter 5 of Volume 1 of this manual; and XSB's computation of the Well-founded semantics has a well-defined relationship to stable model semantics. Furthermore, deductive-database-like capabilities of XSB allow it to be an efficient and flexible grounder for many ASP problems.

The XASP package provides various mechanisms that allow tight linkage of XSB programs to the Smodels [18] stable model generator. The main interface is based on a store of clauses that can

---

[1] Thanks to Barry Evans for helping resuscitate the XASP installation procedure, and to Gonçalo Lopes for the installation procedure on Windows.

[2] In Version 3.2, the Smodels API has not been tested with the multi-threaded engine, and Smodels itself is not thread-safe.

be incrementally asserted or deleted by an XSB program. Clauses in this store can make use of all of the cardinality and weight constraint syntax supported by Smodels, in addition to default negation. When the user decides that the clauses in a store are a complete representation of a program whose stable model should be generated, the clauses are copied into Smodels buffers. Using the Smodels API, the generator is invoked, and information about any stable models generated are returned. This use of XASP is roughly analogous to building up a constraint store in CLP, and periodically evaluating that store, but integration with the store is less transparent in XASP than in CLP. In XASP, clauses must be explicitly added to a store and evaluated; furthermore clauses are not removed from the store upon backtracking, unlike constraints in CLP.

The XNMR interpreter provides a second, somewhat more implicit use of XASP. In the XNMR interface a query Q is evaluated as is any other query in XSB. However, conditional answers produced for Q and for its subgoals, upon user request, can be considered as clauses and sent to Smodels for evaluation. In backtracking through answers for Q, the user backtracks not only through answer substitutions for variables of Q, but also through the stable models produced for the various bindings.

13.1 Installing the Interface

Installing the Smodels interface of XASP sometimes can be tricky for two reasons. First, XSB must dynamically load the Smodels library, and dynamic loading introduces platform dependencies. Second since Smodels is written in C++ and XSB is written in C, the load must ensure that names are properly resolved and that C++ libraries are loaded, steps that may addressed differently by different compilers [3]. However, by following the steps outlined below in the section for Unix or Windows, XASP should be running in a matter of minutes.

13.1.1 Installing the Interface under Unix

In order to use the Smodels interface, several steps must be performed.

1. *Creating a library for Smodels.* Smodels itself must be compiled as a library. Unlike previous versions of XSB, which required a special configuration step for Smodels, Version 3.2 requires no special confiuration, since XSB includes source code for Smodels 2.32 as a subdirectory of the `$XSBDIR/packages/xasp` directory (denoted `$XASPDIR`). We suggest making Smodels out of this directory [4]. Thus, to make the Smodels library

   (a) Change directory to `$XASPDIR/smodels`
   (b) On systems other than OS X, type

$$\texttt{make lib}$$

---

[3]XSB's compiler can automatically call foreign compilers to compile modules written in C, but in Version 3.2 of XSB C++ modules must be compiled with external commands, such as the `make` command shown below.

[4]Although distributed with XSB, Smodels is distributed under the GNU General Public License, a license that is slightly stricter than the license XSB uses. Users distributing applications based on XASP should be aware of any restrictions imposed by GNU General Public License.

on OS X, type [5]

```
make -f Makefile.osx lib
```

If the compilation step ran successfully, there should be a file `libsmodels.so` (or `libsomodels.dylib` on MacOS X or `libsmodels.dll` on Windows...) in `$XASPDIR/smodels/.libs`

(c) Change directory back to `$XASPDIR`

2. *Compiling the XASP files* Next, platform-specific compilation of XASP files needs to be performed. This can be done by consulting `prologMake.P` and executing the goal

```
?- make.
```

It is important to note that under Version 3.2, code compiled by the single threaded engine will only be executable by the single threaded engine, and code compiled by the multi-threaded engine will only be executable by the multi-threaded engine.

3. *Checking the Installation* To see if the installation is working properly, cd to the subdirectory `tests` and type:

`sh testsuite.sh <$XSBDIR>`

If the test suite succeeded it will print out a message along the lines of

`PASSED testsuite for /Users/terranceswift/XSBNEW/XSB/config/powerpc-apple-darwin7.5.1/bin/xsb`

13.1.2 Installing XASP under Windows using Cygwin

To install XASP under Windows, you must use Version 3.2 of XSB or later and Version 2.31 or later of Smodels [6]. You should also have a recent version of Cygwin (e.g. 1.5.20 or later) with all the relevant development packages installed, such as `devel`, `make`, `automake`, `patchtools`, and possibly `x11` (for `makedepend`) Without an appropriate Cygwin build environment many of these steps will simply fail, sometimes with quite cryptic error messages.

1. *Patch and Compile Smodels* First, uncompress `smodels-2.31.tar.gz` in some directory, (for presentation purposes we use `/cygdrive/c/smodels-2.31` — that is, `c:\smodels-2.31`). After that, you must apply the patch provided with this package. This patch enables the creation of a DLL from Smodels. Below is a sample session (system output omitted) with the required commands:

```
$ cd /cygdrive/c/smodels-2.31
$ cat $XSB/packages/xasp/patch-smodels-2.31 | patch -p1
$ make lib
```

---

[5]A special makefile is needed for OS X since the GNU libtool is called glibtool on this platform.
[6]This section was written by Goncalo Lopes.

After that, you should have a file called `smodels.dll` in the current directory, as well as a file called `smodels.a`. You should make the former "visible" to Windows. Two alternatives are either (a) change the `PATH` environment variable to contain `c:\smodels-2.31`, or (b) copy `smodels.dll` to some other directory in your PATH (such as `c:\windows`, for instance). One simple way to do this is to copy `smodels.dll` to `$XSB/config/i686-pc-cygwin/bin`, *after* the configure XSB step (step 2), since that directory has to be in your path in order to make XSB fully functional.

2. *Configure XSB.* In order to properly configure XSB, you must tell it where the Smodels sources and library (the `smodels.a` file) are. In addition, you must compile XSB such that it doesn't use the Cygwin DLL (using the `-mno-cygwin` option for gcc). The following is a sample command:

```
$ cd $XSB/build
$ ./configure --enable-no-cygwin -with-smodels="/cygdrive/c/smodels-2.31''
```

You can optionally include the extended Cygwin w32 API using the configuration option `--with-includes=<PATH_TO_API>`, (this allows XSB's build procedure to find `makedepend` for instance), but you'll probably do fine with just the standard Cygwin apps.

There are some compiler variables which may not be automatically set by the configure script in `xsb_config.h`, namely the configuration names and some activation flags. To correct this, do the following:

(a) cd to `$XSB/config/i686-pc-cygwin`

(b) open the file `xsb_config.h` and add the following lines:

```
#define CONFIGURATION "i686-pc-cygwin"
#define FULL_CONFIG_NAME "i686-pc-cygwin"
#define SLG_GC
```

(Still more flags may be needed depending on Cygwin configuration)

After applying these changes, cd back to the `$XSB/build` directory and compile XSB:

```
$ ./makexsb
```

Now you should have in `$XSB/config/i686-pc-cygwin/bin` directory both a `xsb.exe` and a `xsb.dll`.

3. *Compiling XASP.* First, go to the XASP directory and execute the `makelinks.sh` script in order to make the headers and libraries in Smodels be accessible to XSB, i.e.:

```
$ cd $XSB/packages/xasp
$ sh makelinks.sh /cygdrive/c/smodels-2.31
```

Now you must copy the `smoMakefile` from the `config` directory to the `xasp` directory and run both its directives:

```
$ cp $XSB/config/i686-pc-cygwin/smoMakefile .
$ make -f smoMakefile module
$ make -f smoMakefile all
```

At this point, you can consult `xnmr` as you can with any other package, or xsb with the `xnmr` command line parameter, like this: (don't forget to add XSB bin directory to the `$PATH` environment variable)

```
$ xsb xnmr
```

Lots of error messages will probably appear because of some runtime load compiler, but if everything goes well you can ignore all of them since your xasppkg will be correctly loaded and everything will be functioning smoothly from there on out.

13.2   The Smodels Interface

The Smodels interface contains two levels: the *cooked* level and the *raw* level. The cooked level interns rules in an XSB *clause store*, and translates general weight constraint rules [21] into a *normal form* that the Smodels engine can evaluate. When the programmer has determined that enough clauses have been added to the store to form a semantically complete sub-program, the program is *committed*. This means that information in the clauses is copied to Smodels and interned using Smodels data structures so that stable models of the clauses can be computed and examined. By convention, the cooked interface ensures that the atom `true` is present in all stable models, and the atom `false` is false in all stable models. The raw level models closely the Smodels API, and demands, among other things, that each atom in a stable sub-program has been translated into a unique integer. The raw level also does not provide translation of arbitrary weight constraint rules into the normal form required by the Smodels engine. As a result, the raw level is significantly more difficult to directly use than the cooked level. While we make public the APIs for both the raw and cooked level, we provide support only for users of the cooked interface.

As mentioned above Smodels extends normal programs to allow weight constraints, which can be useful for combinatorial problems. However, the syntax used by Smodels for weight constraints does not follow ISO Prolog syntax so that the XSB syntax for weight constraints differs in some respects from that of Smodels. Our syntax is defined as follows, where A is a Prolog atom, N a non-negative integer, and I an arbitrary integer.

- *GeneralLiteral ::= WeightConstraint | Literal*

- *WeightConstraint ::= weightConst(Bound,WeightList,Bound)*

- *WeightList ::= List of WeightLiterals*

- *WeightLiteral ::= Literal | weight(Literal,N)*

- *Literal ::= A | not(A)*

- *Bound ::== I* | `undef`

Thus an example of a weight constraint might be:

- `weightConst(1,[weight(a,1),weight(not(b),1)],2)`

We note that if a user does not wish to put an upper or lower bound on a weight constraint, she may simply set the bound to `undef` or to an integer less than `0`.

The intuitive semantics of a weight constraint `weightConst(Lower,WeightList,Upper)`, in which `List` is is list of *WeightLiterals* that it is true in a model M whenever the sum of the weights of the literals in the constraint that are true in M is between the lower `Lower` and `Upper`. Any literal in a *WeightList* that does not have a weight explicitly attached to it is taken to have a weight of *1*.

In a typical session, a user will initialize the Smodels interface, add rules to the clause store until it contains a semantically meaningful sub-problem. He can then specify a compute statement if needed, commit the rules, and compute and examine stable models via backtracking. If desired, the user can then re-initialize the interface, and add rules to or retract rules from the clause store until another semantically meaningful sub-program is defined; and then commit, compute and examine another stable model [7].

The process of adding information to a store and periodically evaluating it is vaguely reminiscent of the Constraint Logic Programming (CLP) paradigm, but there are important differences. In CLP, constraints are part of the object language of a Prolog program: constraints are added to or projected out of a constraint store upon forward execution, removed upon backwards execution, and iteratively checked. When using this interface, on the other hand, an XSB program essentially acts as a compiler for the clause store, which is treated as a target language. Clauses must be explicitly added or removed from the store, and stable model computation cannot occur incrementally – it must wait until all clauses have been added to the store. We note in passing that the `xnmr` module provides an elegant but specialized alternative. `xnmr` integrates stable models into the object language of XSB, by computing ""relevant"" stable models from the the residual answers produced by query evaluation. It does not however, support the weighted constraint rules, compute statements and so on that this module supports.

Neither the raw nor the cooked interface currently supports explicit negation.

Examples of use of the various interfaces can be found in the subdirectory `intf_examples`

smcInit

  Initializes the XSB clause store and the Smodels API. This predicate must be executed before building up a clause store for the first time. The corresponding raw predicate, `smrInit(Num)`, initializes the Smodels API assuming that it will require at most `Num` atoms.

smcReInit

  Reinitializes the Smodels API, but does *not* affect the XSB clause store. This predicate is

---

[7]Currently, only normal rules can be retracted.

provided so that a user can reuse rules in a clause store in the context of more than one sub-program.

smcAddRule(+Head,+Body)

Interns a ground rule into the XSB clause store. `Head` must be a *GeneralLiteral* as defined at the beginning of this section, and `Body` must be a list of *GeneralLiterals*. Upon interning, the rule is translated into a normal form, if necessary, and atoms are translated to unique integers. The corresponding raw predicates, `smrAddBasicRule/3`, `smrAddChoiceRule/3`, `smrAddConstraintRule/4`, and `smrAddWeightRule/3` can be used to add raw predicates immediately into the SModels API.

smcRetractRule(+Head,+Body)

Retracts a ground (basic) rule from the XSB clause store. Currently, this predicate cannot retract rules with weight constraints: `Head` must be a *Literal* as defined at the beginning of this section, and `Body` must be a list of *GeneralLiterals*.

smcSetCompute(+List)

Requires that `List` be a list of literals – i.e. atoms or the default negation of atoms). This predicate ensures that each literal in `List` is present in the stable models returned by Smodels. By convention the cooked interface ensures that `true` is present and `false` absent in all stable models. After translating a literal it calls the raw interface predicates `smrSetPosCompute/1` and `smrSetNegCompute/1`

smcCommitProgram

This predicate translates all of the clauses from the XSB clause store into the data structures of the Smodels API. It then signals to the API that all clauses have been added, and initializes the Smodels computation. The corresponding raw predicate, `smrCommitProgram`, performs only the last two of these features.

smComputeModel

This predicate calls Smodels to compute a stable model, and succeeds if a stable model can be computed. Upon backtracking, the predicate will continue to succeed until all stable models for a given program cache have been computed. `smComputeModel/0` is used by both the raw and the cooked levels.

smcExamineModel(+List,-Atoms)

`smcExamineModel/(+List,-Atoms)` filters the literals in `List` to determine which are true in the most recently computed stable model. These true literals are returned in the list `Atoms`. `smrExamineModel(+N,-Atoms)` provides the corresponding raw interface in which integers from `0` to `N`, true in the most recently computed stable model, are input and output.

smEnd

Reclaims all resources consumed by Smodels and the various APIs. This predicate is used by both the cooked and the raw interfaces.

print_cache

This predicate can be used to examine the XSB clause store, and may be useful for debugging.

13.2.1 Using the Smodels Interface with Multiple Threads

If XASP has been compiled under the multi-threaded engine, the Smodels interface will be fully thread-safe: this means that Smodels and all interface predicates described in this section can be used concurrently by different threads. In multi-threaded XASP, each XSB thread can initialize and query its own instance of Smodels, and build up its own private clause store at both the cooked and raw levels (shared clause stores are not yet available). Figure 13.1 provides a simple example of how this can be done. For each thread that will generate stable models, a message queue is created that will be used to communicate back results. Two threads are then created and these threads concurrently add rules to their private clause stores, call Smodels, and send the results back to the calling thread using the appropriate message queue. Of course the example here is just one of many possible: answers could be returned using different configurations of message queues, through shared tables, through shared asserted code, and so on.

13.3 The xnmr_int Interface

. This module provides the interface from the `xnmr` module to Smodels. It does not use the `sm_int` interface, but rather directly calls the Smodels C interface, and can be thought of as a special-purpose alternative to `sm_int`.

init_smodels(+Query)
> Initializes smodels with the residual program produced by evaluating `Query`. `Query` must be a call to a tabled predicate that is currently completely evaluated (and should have a delay list)

atom_handle(?Atom,?AtomHandle)
> The *handle* of an atom is set by `init_smodels/1` to be an integer uniquely identifying each atoms in the residual program (and thus each atom in the Herbrand base of the program for which the stable models are to be derived. The initial query given to `init_smodels` has the atom-handle of 1.

in_all_stable_models(+AtomHandle,+Neg)
> `in_all_stable_models/2` returns true if `Neg` is 0 and the atom numbered `AtomHandle` returns true in all stable models (of the residual program set by the previous call to `init_smodels/1`). If `Neg` is nonzero, then it is true if the atom is in NO stable model.

pstable_model(+Query,-Model,+Flag)
> returns nondeterministically a list of atoms true in the partial stable model total on the atoms relevant to instances of `Query`, if `Flag` is 0. If `Flag` is 1, it only returns models in which the instance of `Query` is true.

a_stable_model
> This predicate invokes Smodels to find a (new) stable model (of the program set by the previous invocation of `init_smodels/1`.) It will compute all stable models through backtracking. If there are no (more) stable models, it fails. Atoms true in a stable model can be examined by `in_current_stable_model/1`.

```
:- ensure_loaded(xasp).
:- import smcInit/0, smcAddRule/2, smcCommitProgram/0 smcSetCompute/1,
         smComputeModel/0, smcExamineModel/1, smEnd/0 from sm_int.
:- import thread_create/1 from thread.
:- import thread_get_message/2,  thread_send_message/2, message_queue_create/1 from mutex_xsb.

test:-
     message_queue_create(Queue1),
     message_queue_create(Queue2),
     thread_create(test1(Queue1)),
     thread_create(test2(Queue2)),
     read_models(Queue1),
     read_models(Queue2).

test1(Queue) :-
     smcInit,
     smcAddRule(a1,[]),
     smcAddRule(b1,[]),
     smcAddRule(d1,[a1,not(c1)]),
     smcAddRule(c1,[b1,not(d1)]),
     smcCommitProgram,
     write('All Solutions: '),nl,
     (    smComputeModel,
          smcExamineModel(Model),
          thread_send_message(Queue,solution(program1,Model)),
          fail
     ;
          thread_send_message(Queue,no_more_solutions),
          smEnd  ).

test2(Queue) :-
     smcInit,
     smcAddRule(a2,[]),
     smcAddRule(b2,[]),
     smcAddRule(d2,[a2,not(c2)]),
     smcAddRule(c2,[b2,not(d2)]),
     smcCommitProgram,
     write('All Solutions: '),nl,
     (    smComputeModel,
          smcExamineModel(Model),
          thread_send_message(Queue,solution(program2,Model)),
          fail
     ;
          thread_send_message(Queue,no_more_solutions),
          smEnd  ).

read_models(Queue):-
     repeat,
     thread_get_message(Queue,Message),
     (Message = no_more_solutions ->
          true
       ; writeln(Message),
          fail ).
```

Figure 13.1: Using the Smodels Interface with Multi-Threading

`in_current_stable_model(?AtomHandle)`

> This predicate is true of handles of atoms true in the current stable model (set by an invocation of `a_stable_model/0`.)

`current_stable_model(-AtomList)`

> returns the list of atoms true in the current stable model.

`print_current_stable_model`

> prints the current stable model to the stream to which answers are sent (i.e `stdfbk`)

Chapter 14

Other XSB Packages

Many of the XSB packages are maintained somewhat independently of XSB and have their own manuals. For these packages: *Flora2*, *XMC*, *xsbdoc* and *Cold Dead Fish* we provide summaries; full information can be obtained in the packages themselves. In addition, we provide full documentation here for two of the smaller packages, `slx` and `GAP`.

14.1   Programming with FLORA-2

ℱLORA-*2* is a sophisticated object-oriented knowledge base language and application development platform. It is implemented as a set of run-time libraries and a compiler that translates a unified language of F-logic [14], HiLog [7], and Transaction Logic [4, 3] into tabled Prolog code.

Applications of ℱLORA-*2* include intelligent agents, Semantic Web, ontology management, integration of information, and others.

The programming language supported by ℱLORA-*2* is a dialect of F-logic with numerous extensions, which include a natural way to do meta-programming in the style of HiLog and logical updates in the style of Transaction Logic. ℱLORA-*2* was designed with extensibility and flexibility in mind, and it provides strong support for modular software design through its unique feature of dynamic modules. Other extensions, such as the versatile syntax of FLORID path expressions, are borrowed from FLORID, a C++-based F-logic system developed at Freiburg University.[1] Extensions aside, the syntax of ℱLORA-*2* differs in many important ways from FLORID, from the original version of F-logic, as described in [14], and from an earlier implementation of ℱLORA. These syntactic changes were needed in order to bring the syntax of ℱLORA-*2* closer to that of Prolog and make it possible to include simple Prolog programs into ℱLORA-*2* programs without choking the compiler. Other syntactic deviations from the original F-logic syntax are a direct consequence of the added support for HiLog, which obviates the need for the "@" sign in method invocations (this sign is now used to denote calls to ℱLORA-*2* modules).

ℱLORA-*2* is distributed in two ways. First, it is part of the official distribution of XSB and thus is installed together with XSB. Second, a more up-to-date version of the system is available

---

[1] See `http://www.informatik.uni-freiburg.de/∼dbis/florid/` for more details.

on $\mathcal{F}$LORA-2's Web site at

```
http://flora.sourceforge.net
```

These two versions can be installed at the same time and used independently (*e.g.*, if you want to keep abreast with the development of $\mathcal{F}$LORA-2 or if a newer version was released in-between the releases of XSB). The installation instructions are somewhat different in these two cases. Here we only describe the process of configuring the version $\mathcal{F}$LORA-2 included with XSB.

Installing $\mathcal{F}$LORA-2 under UNIX.   To configure a version of $\mathcal{F}$LORA-2 that was downloaded as part of the distribution of XSB, simply configure XSB as usual:

```
cd XSB/build
configure
makexsb
```

and then run

```
makexsb packages
```

If you downloaded XSB from its CVS repository earlier and are updating your copy using the `cvs update` command, then it might be a good idea to also do the following:

```
cd packages/flora2
makeflora clean
makeflora
```

Installing $\mathcal{F}$LORA-2 in Windows.   First, you need Microsoft's `nmake`. Then use the following commands to configure $\mathcal{F}$LORA-2 (assuming that XSB is already installed and configured):

```
cd flora2
makeflora clean
makeflora path-to-prolog-executable
```

Also make sure that the `packages` directory contains a shortcut called `flora2.P` to the file `packages\flora2\flora2.P`.

Running $\mathcal{F}$LORA-2. $\mathcal{F}$LORA-2 is fully integrated into the underlying XSB engine, including its module system. In particular, $\mathcal{F}$LORA-2 modules can invoke predicates defined in other Prolog modules, and Prolog modules can query the objects defined in $\mathcal{F}$LORA-2 modules.

Due to certain problems with XSB, $\mathcal{F}$LORA-2 runs best when XSB is configured with *local* scheduling, which is the default XSB configuration. However, with this type of scheduling, many Prolog intuitions that relate to the operational semantics do not work. Thus, the programmer

must think "more declaratively" and, in particular, to not rely on the order in which answers are returned.

The easiest way to get a feel of the system is to start $\mathcal{F}$LORA-*2* shell and begin to enter queries interactively. The simplest way to do this is to use the shell script

```
.../flora2/runflora
```

where "..." is the directory where $\mathcal{F}$LORA-*2* is downloaded. For instance, to invoke the version supplied with XSB, you would type something like

```
~/XSB/packages/flora2/runflora
```

At this point, $\mathcal{F}$LORA-*2* takes over and F-logic syntax becomes the norm. To get back to the Prolog command loop, type `Control-D` (Unix) or Control-Z (Windows), or

```
| ?- flEnd.
```

If you are using $\mathcal{F}$LORA-*2* shell frequently, it pays to define an alias, say (in Bash):

```
alias runflora='~/XSB/packages/flora2/runflora'
```

$\mathcal{F}$LORA-*2* can then be invoked directly from the shell prompt by typing `runflora`. It is even possible to tell $\mathcal{F}$LORA-*2* to execute commands on start-up. For instance,

```
foo>  runflora -e "flHelp."
```

will cause the system to execute the help command right after after the initialization. Then the usual $\mathcal{F}$LORA-*2* shell prompt is displayed.

$\mathcal{F}$LORA-*2* comes with a number of demo programs that live in

```
.../flora2/demos/
```

The demos can be run issuing the command "`flDemo(demo-filename).`" at the $\mathcal{F}$LORA-*2* prompt, *e.g.*,

```
flora2 ?- flDemo(flogic_basics).
```

There is no need to change to the demo directory, as `flDemo` knows where to find these programs.

14.2   Summary of xmc: Model-checking with XSB

No documentation yet available.

the Ciao [6] system's *lpdoc* which has been adapted to generate a reference manual automatically from one or more XSB source files. The target format of the documentation can be Postscript, HTML, PDF, or nicely formatted ASCII text. `xsbdoc` can be used to automatically generate a description of full applications, library modules, README files, etc. A fundamental advantage of using `xsbdoc` to document programs is that it is much easier to maintain a true correspondence between the program and its documentation, and to identify precisely to what version of the program a given printed manual corresponds. Naturally, the `xsbdoc` manual is generated by `xsbdoc` itself.

The quality of the documentation generated can be greatly enhanced by including within the program text:

- *assertions* (indicating types, modes, etc. ...) for the predicates in the program, via the directive `pred/1`; and

- *machine-readable comments* (in the "literate programming" style).

The assertions and comments included in the source file need to be written using the forthcoming XSB *assertion language*, which supports most of the features of Ciao's assertion language within a simple and (hopefully) intuitive syntax.

`xsbdoc` is distributed under the *GNU general public license*.

Unlike `lpdoc`, `xsbdoc` does not use Makefiles, and instead maintains information about how to generate a document within Prolog *format files*. As a result, `xsbdoc` can in principle be run in any environment that supports the underlying software, such as XSB, LaTeX, `dvips` and so on. It has been tested on Linux and Windows running with Cygwin.

14.3   slx: Extended Logic Programs under the Well-Founded Semantics

As explained in the section *Using Tabling in XSB*, XSB can compute normal logic programs according to the well-founded semantics. In fact, XSB can also compute *Extended Logic Programs*, which contain an operator for explicit negation (written using the symbol - in addition to the negation-by-failure of the well-founded semantics (`\+` or `not`). Extended logic programs can be extremely useful when reasoning about actions, for model-based diagnosis, and for many other uses [2]. The library, slx provides a means to compile programs so that they can be executed by XSB according to the *well-founded semantics with explicit negation* [1]. Briefly, WFSX is an extension of the well-founded semantics to include explicit negation and which is based on the *coherence principle* in which an atom is taken to be default false if it is proven to be explicitly false, intuitively:

$$-p \Rightarrow not\ p.$$

This section is not intended to be a primer on extended logic programming or on WFSX semantics, but we do provide a few sample programs to indicate the action of WFSX. Consider the program

```
s:- not t.

t:- r.
t.

r:- not r.
```

If the clause -t were not present, the atoms r, t, s would all be undefined in WFSX just as they would be in the well-founded semantics. However, when the clause t is included, t becomes true in the well-founded model, while s becomes false. Next, consider the program

```
s:- not t.

t:- r.
-t.

r:- not r.
```

In this program, the explicitly false truth value for t obtained by the rule -t overrides the undefined truth value for t obtained by the rule t:- r. The WFSX model for this program will assign the truth value of t as false, and that of s as true. If the above program were contained in the file test.P, an XSB session using test.P might look like the following:

```
> xsb

| ?- [slx].
[slx loaded]

yes
| ?- slx_compile('test.P').
[Compiling ./tmptest]
[tmptest compiled, cpu time used: 0.1280 seconds]
[tmptest loaded]

| ?- s.

yes
| ?- t.

no
| ?- naf t.

yes
| ?- r.

no
| ?- naf r.

no
| ?- und r.
```

```
        yes
```

In the above program, the query `?- t.` did not succeed, because `t` is false in WFSX: accordingly the query `naf t` did succeed, because it is true that t is false via negation-as-failure, in addition to `t` being false via explicit negation. Note that after being processed by the SLX preprocessor, `r` is undefined but does not succeed, although `und r` will succeed.

We note in passing that programs under WFSX can be paraconsistent. For instance in the program.

```
        p:- q.

        q:- not q.
        -q.
```

both `p` and `q` will be true *and* false in the WFSX model. Accordingly, under SLX preprocessing, both `p` and `naf p` will succeed.

`slx_compile(+File)`   module: slx
    Preprocesses and loads the extended logic program named `File`. Default negation in `File` must be represented using the operator `not` rather than using `tnot` or `\+`. If L is an objective literal (e.g. of the form A or $-A$ where A is an atom), a query `?- L` will succeed if L is true in the WFSX model, `naf L` will succeed if L is false in the WFSX model, and `und L` will succeed if L is undefined in the WFSX model.

14.4   gapza: Generalized Annotated Programs

Generalized Annotated Programs (GAPs) [15] offer a powerful computational framework for handling paraconsistency and quantitative information within logic programs. The tabling of XSB is well-suited to implementing GAPs, and the gap library provides a meta-interpreter that has proven robust and efficient enough for a commercial application in data mining. The current meta-interpreter is limited to range-restricted programs.

A description of GAPs along with full documentation for this meta-interpreter is provided in [22] (currently also available at `http://www.cs.sunysb.edu/~tswift`). Currently, the interface to the GAP library is through the following call.

`meta(?Annotated_atom)`   module: gap
    If `Annotated_atom` is of the form `Atom:[Lattice_type,Annotation]` the meta-interpreter computes bindings for `Atom` and `Annotation` by evaluating the program according to the definitions provided for `Lattice_type`.

Bibliography

[1] J. Alferes, C. Damasio, and L. Pereira. A logic programming system for non-monotonic reasoning. *Journal of Automated Reasoning*, 1995.

[2] J. Alferes and L. M. Pereira. *Reasoning with Logic Programming*, volume 1111. Springer-Verlag LNAI, 1996.

[3] A. Bonner and M. Kifer. An overview of transaction logic. *Theoretical Computer Science*, 133:205–265, October 1994.

[4] A. Bonner and M. Kifer. A logic for programming database transactions. In J. Chomicki and G. Saake, editors, *Logics for Databases and Information Systems*, chapter 5, pages 117–166. Kluwer Academic Publishers, March 1998.

[5] G. Box and M. Muller. A note on the generation of random normal deviates. *The Annals of Mathematical Statistics*, 29(2):610–611, 1958.

[6] F. Bueno, D. Cabenza, M. Carro, M. Hermenegildo, P. López-García, and G. Puebla. The ciao prolog system, reference manual. Technical report, School of Computer Science, Technical University of Madrid. Available from http://www.clip.dia.fi.upm.es/.

[7] W. Chen, M. Kifer, and D. Warren. HiLog: A foundation for higher-order logic programming. *Journal of Logic Programming*, 15(3):187–230, February 1993.

[8] B. Demoen. Dynamic attributes, their hProlog implementation, and a first evaluation. Report CW 350, Department of Computer Science, K.U.Leuven, Leuven, Belgium, oct 2002. URL = http://www.cs.kuleuven.ac.be/publicaties/rapporten/cw/CW350.abs.html.

[9] C. Draxler. Prolog to SQL compiler, Version 1.0. Technical report, CIS Centre for Information and Speech Processing Ludwig-Maximilians-University, Munich, 1992.

[10] T. Fruehwirth. Thom Fruehwirth's Constraint Handling Rules website. http://www.informatik.uni-ulm.de/pm/mitarbeiter/fruehwirth/chr-intro.html.

[11] T. Frühwirth. Theory and Practice of Constraint Handling Rules. In P. Stuckey and K. Marriot, editors, *Special Issue on Constraint Logic Programming*, volume 37, October 1998.

[12] H. Guo, C. R. Ramakrishnan, and I. V. Ramakrishnan. Speculative beats conservative justification. In *International Conference on Logic Programming*, volume 2237 of *Lecture Notes in Computer Science*, pages 150–165. Springer, 2001.

[13] C. Holzbaur. Ofai clp(q,r) manual, edition 1.3.3. Technical report, Austrian Research Institute for Artificial Intelligence, 1995.

[14] M. Kifer, G. Lausen, and J. Wu. Logical foundations of object-oriented and frame-based languages. *Journal of the ACM*, 42:741–843, July 1995.

[15] M. Kifer and V. S. Subrahmanian. Theory of generalized annotated logic programming and its applications. *J. Logic Programming*, 12(4):335–368, 1992.

[16] T. I. S. Laboratory. *SICStus Prolog User's Manual Version 3.12.5*. Swedish Institute of Computer Science, 2006.

[17] A. McLeod. A remark on algorithm AS 183. *Applied Statistics*, 34:198–200, 1985.

[18] I. Niemelä and P. Simons. Smodels: An implementation of the stable model and well-founded semantics for normal LP. In J. Dix, U. Furbach, and A. Nerode, editors, *Proceedings of the 4th International Conference on Logic Programing and Nonmonotonic Reasoning*, volume 1265 of *LNAI*, pages 420–429, Berlin, July 28–31 1997. Springer.

[19] B. Sanna-Starosta. Chrd: A set-based solver for constraint hanlding rules. available at www.cs.msu.edu/~bss/chr-d, 2006.

[20] B. Sanna-Starosta and C. Ramakrishnan. Compiling constraint handling rules for efficient tabled evaluation. available at www.cs.msu.edu/~bss/chr-d, 2006.

[21] P. Simons, I. Niemelä, and T. Soininen. Extending and implementing the stable model semantics. *Artificial Intelligence*, 138:181–234, 2002.

[22] T. Swift. Tabling for non-monotonic programming. *Annals of Mathematics and Artificial Intelligence*, 25(3-4):201–240, 1999.

[23] B. A. Wichmann and I. D. Hill. Algorithm AS 183: An efficient and portable pseudo-random number generator. *Applied Statistics*, 31:188–190, 1982.

[24] J. Wielemaker. *SWI Prolog version 5.6: Reference Manual*. University of Amsterdam, 2007.

Index